

# OPEN DATA SCIENCE CONFERENCE

London | September 19 - 22 2018



@ODSC

# #ODSC 2018 - London

20.09.2018

## The Path to Deep Learning with Tensorflow + Keras

Juliet Moreiro

Technical Evangelist (Microsoft)

Pablo Doval

Data Pontifex (Plain Concepts)



# Juliet Moreiro

---

TECHNICAL EVANGELIST @ MICROSOFT

Cloud evangelist, you'll never guess my cat's name.

@julietsvq



# Pablo Doval

---

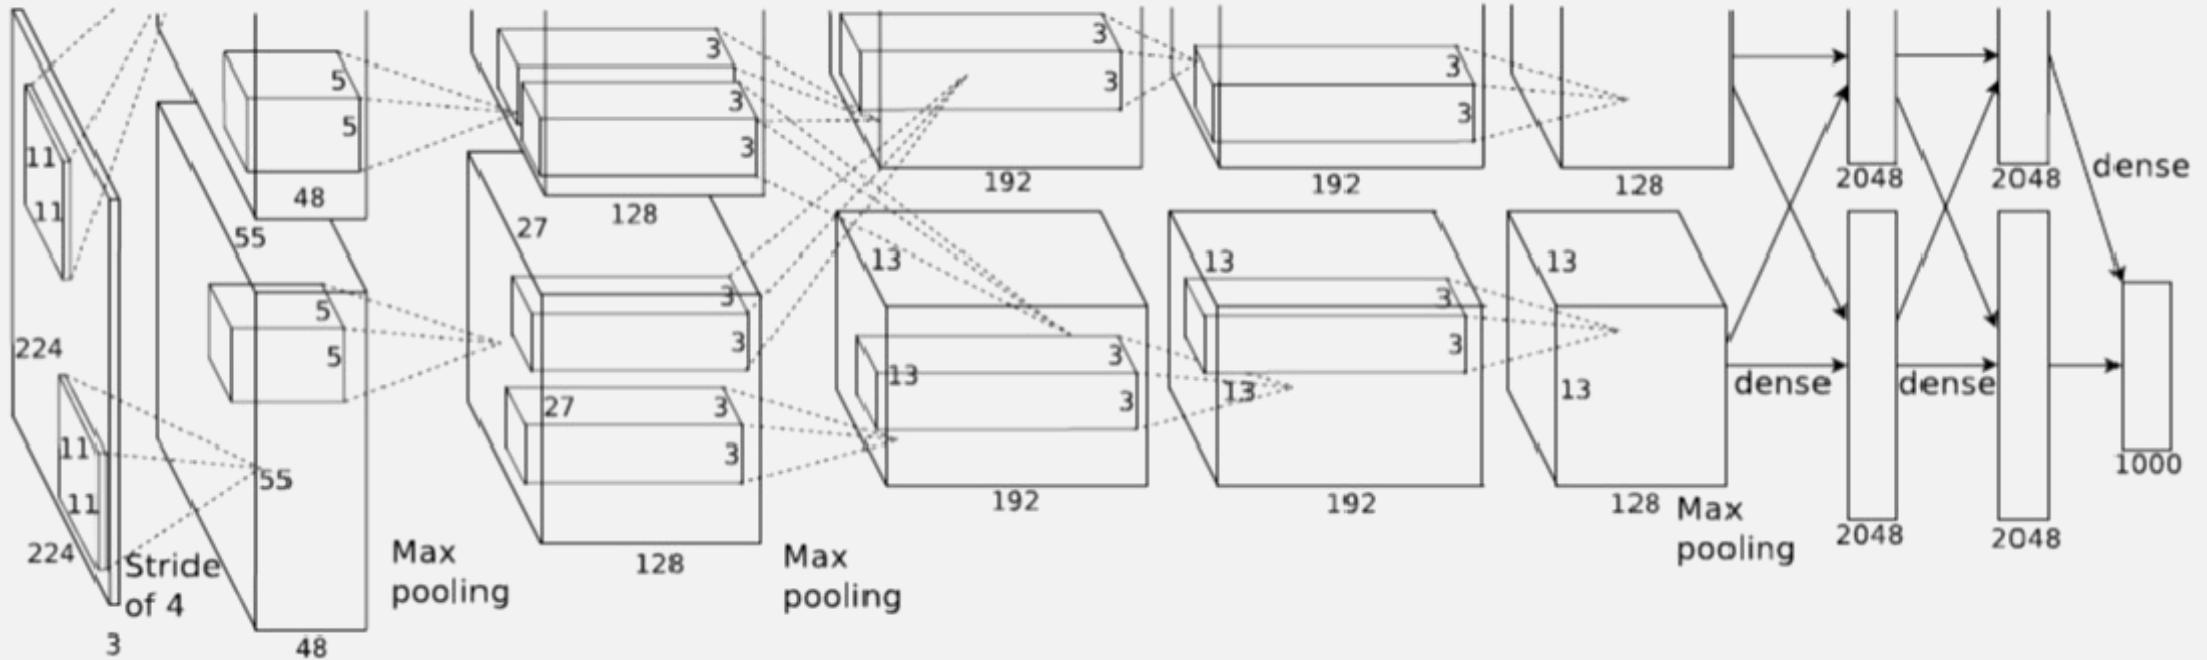
DATA PONTIFEX

I work with code and data, but don't tell my mom; she thinks I'm a piano player in a whorehouse.

**@PabloDoval**

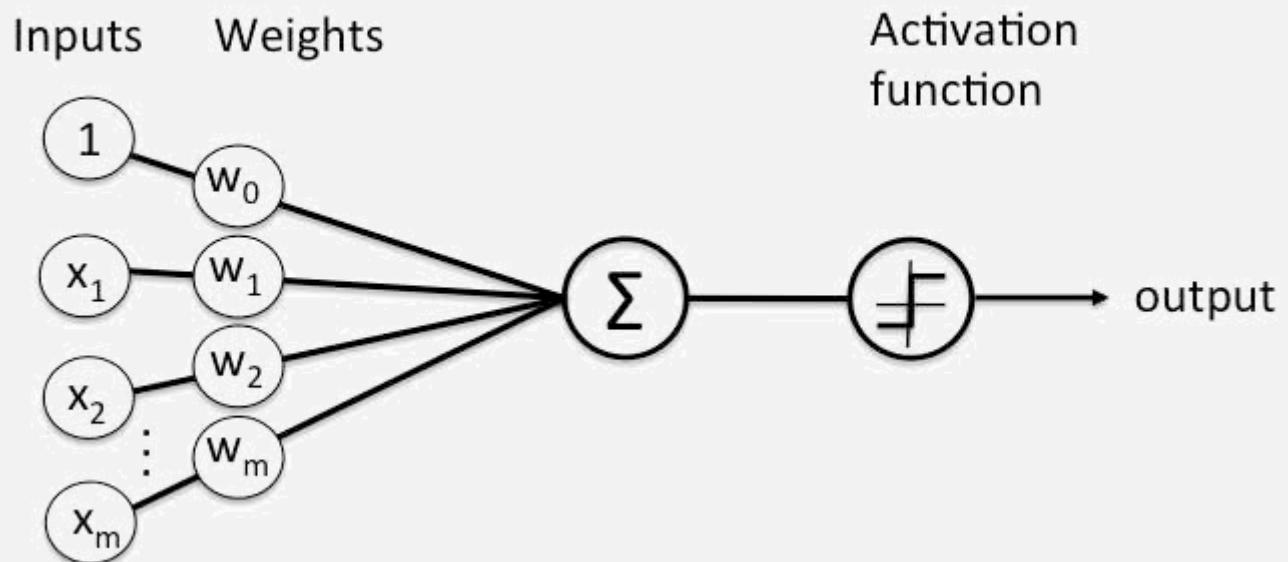
# WHAT IS THIS THING?

---

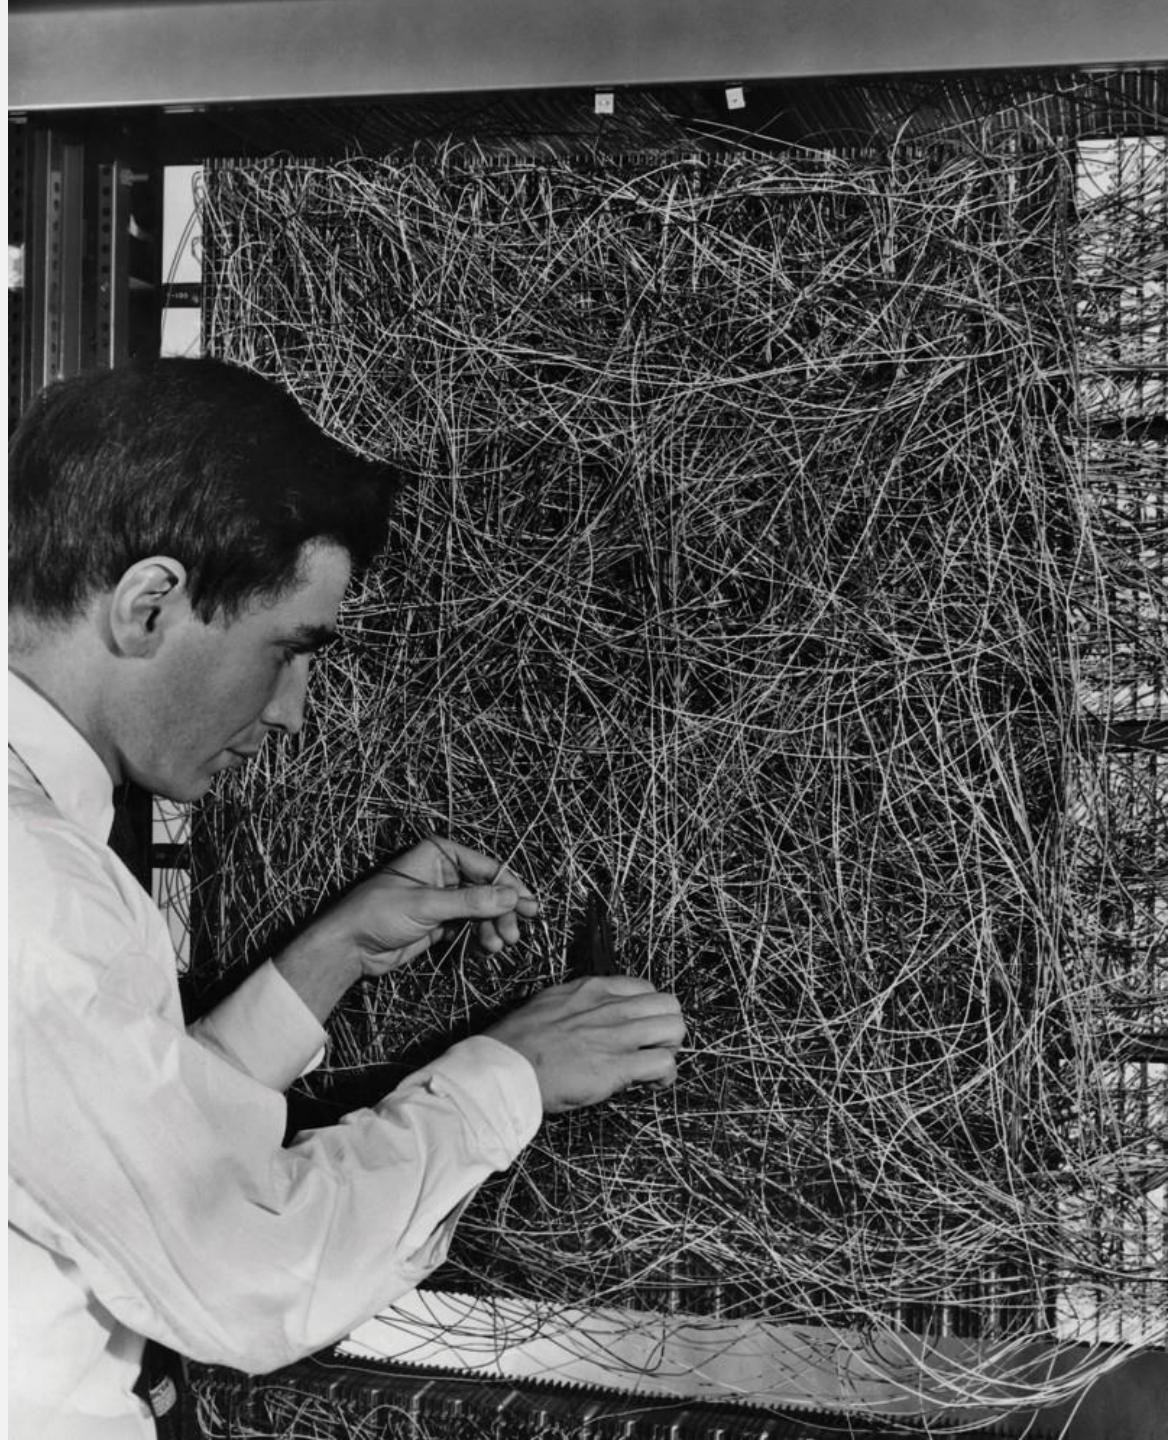
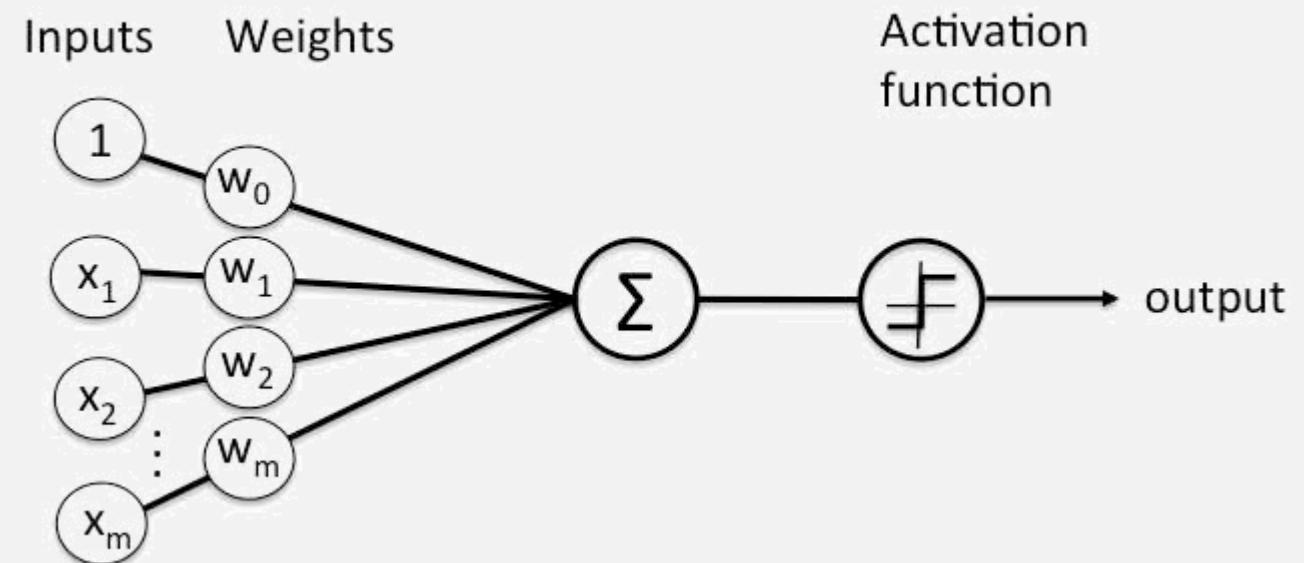




# PERCEPTRON



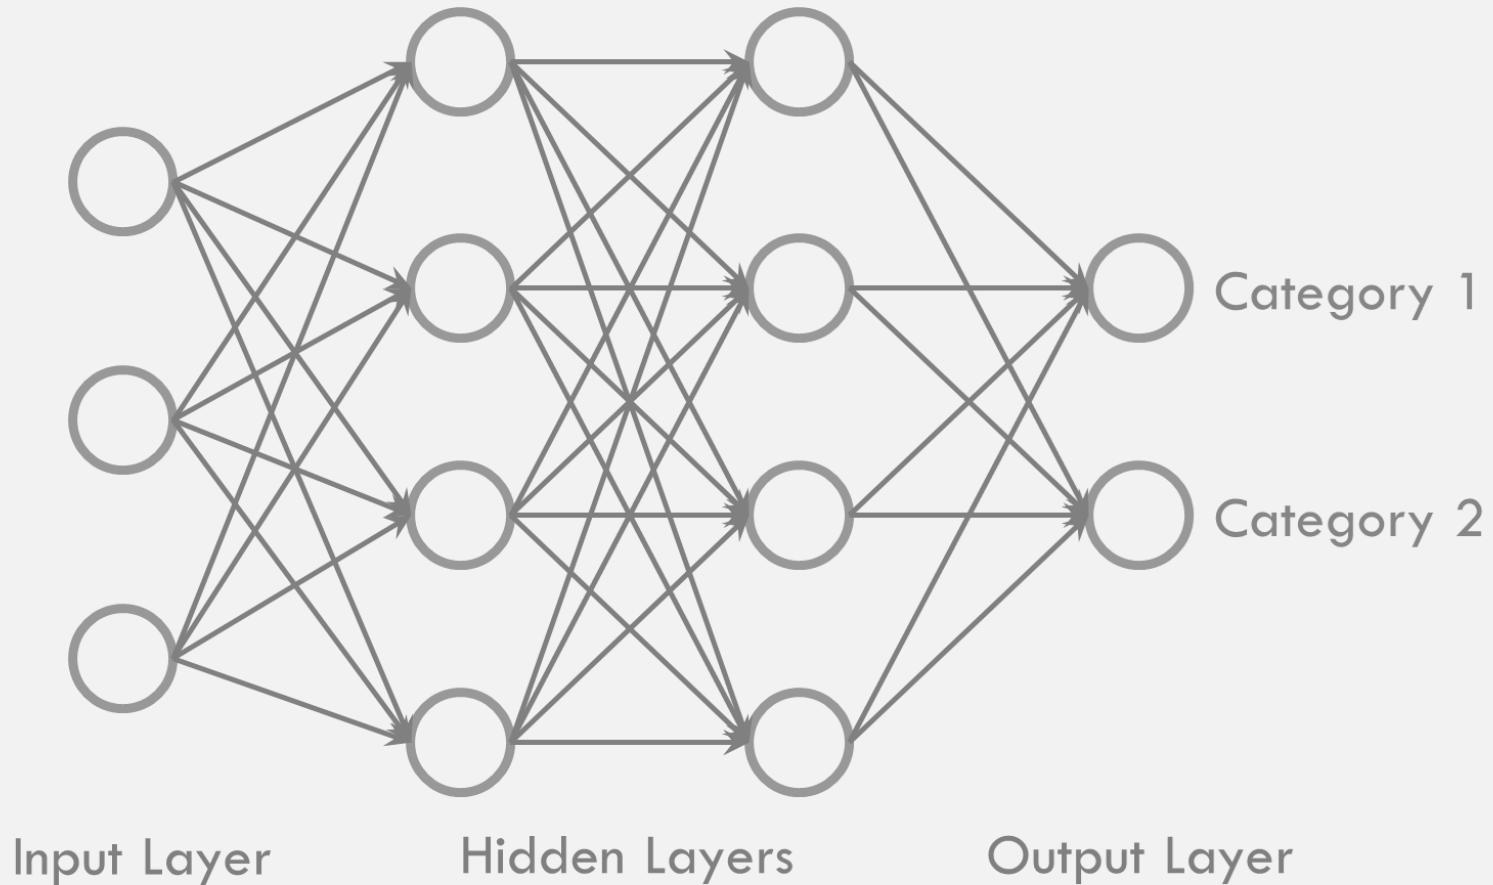
# PERCEPTRON



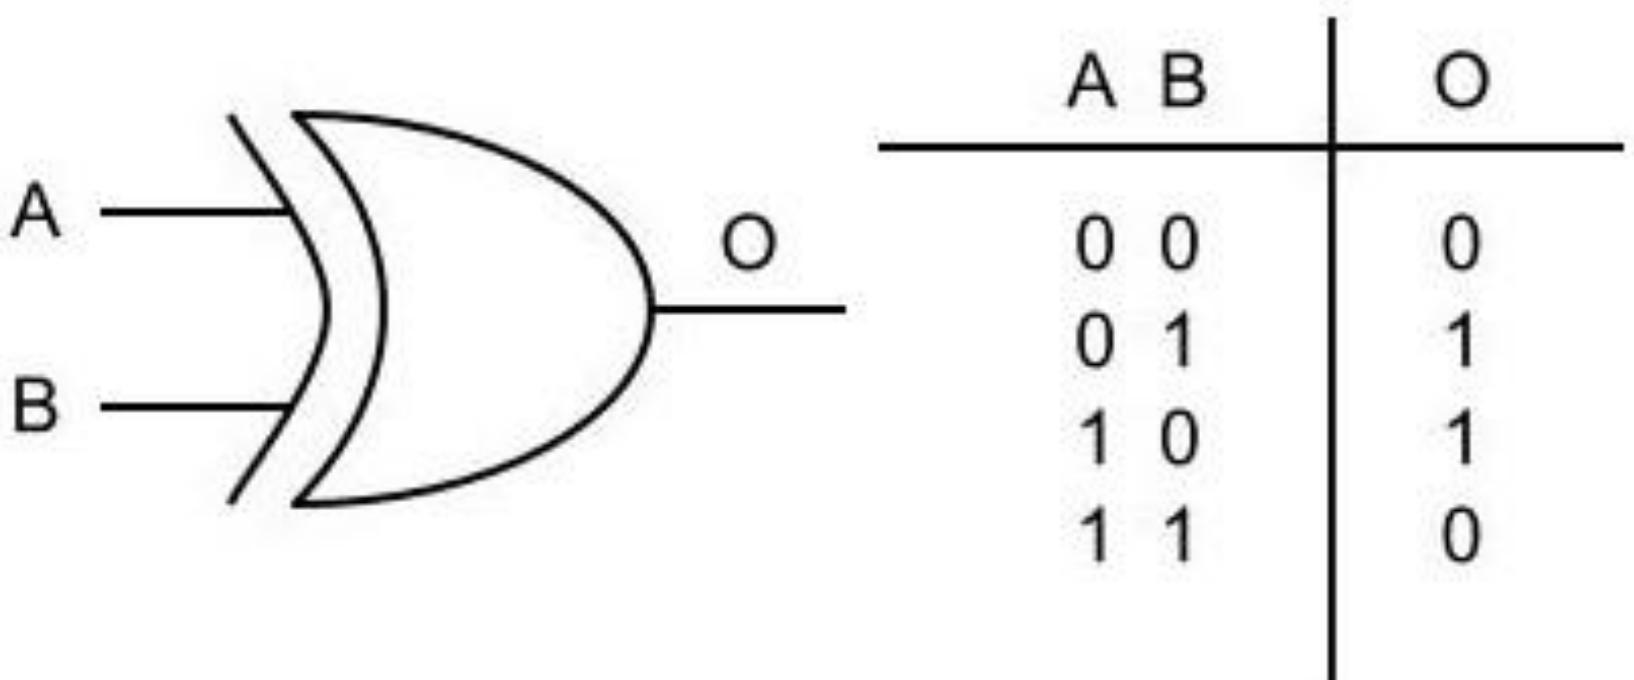
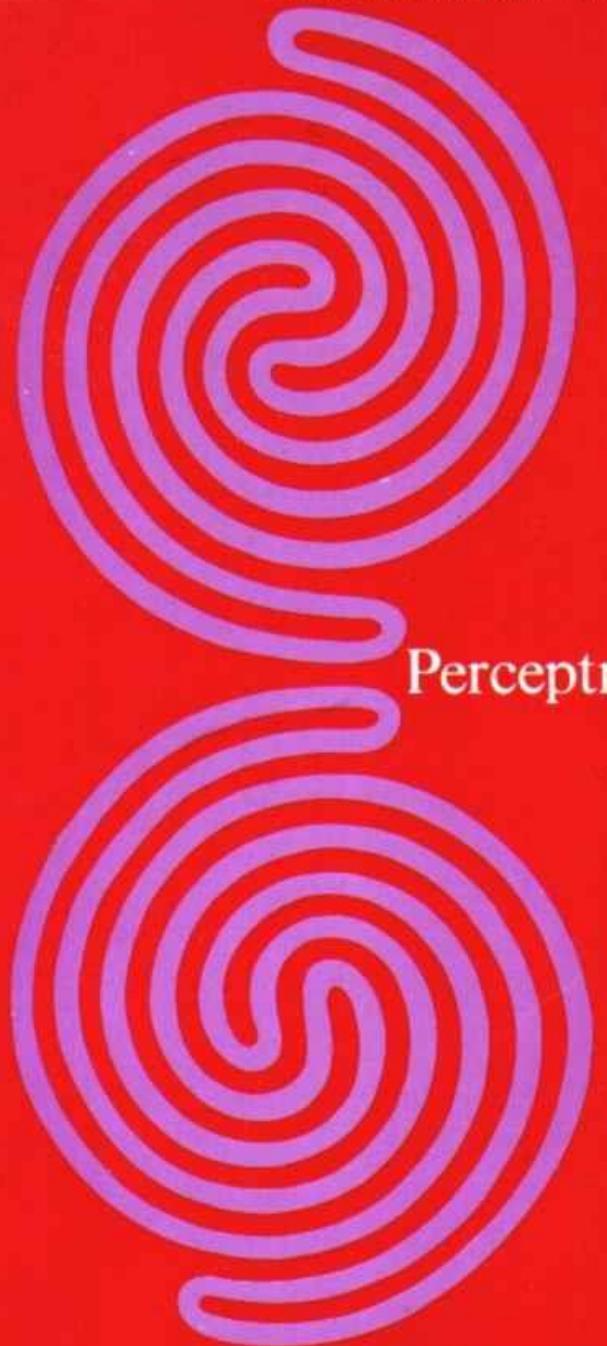
# ARTIFICIAL INTELLIGENCE



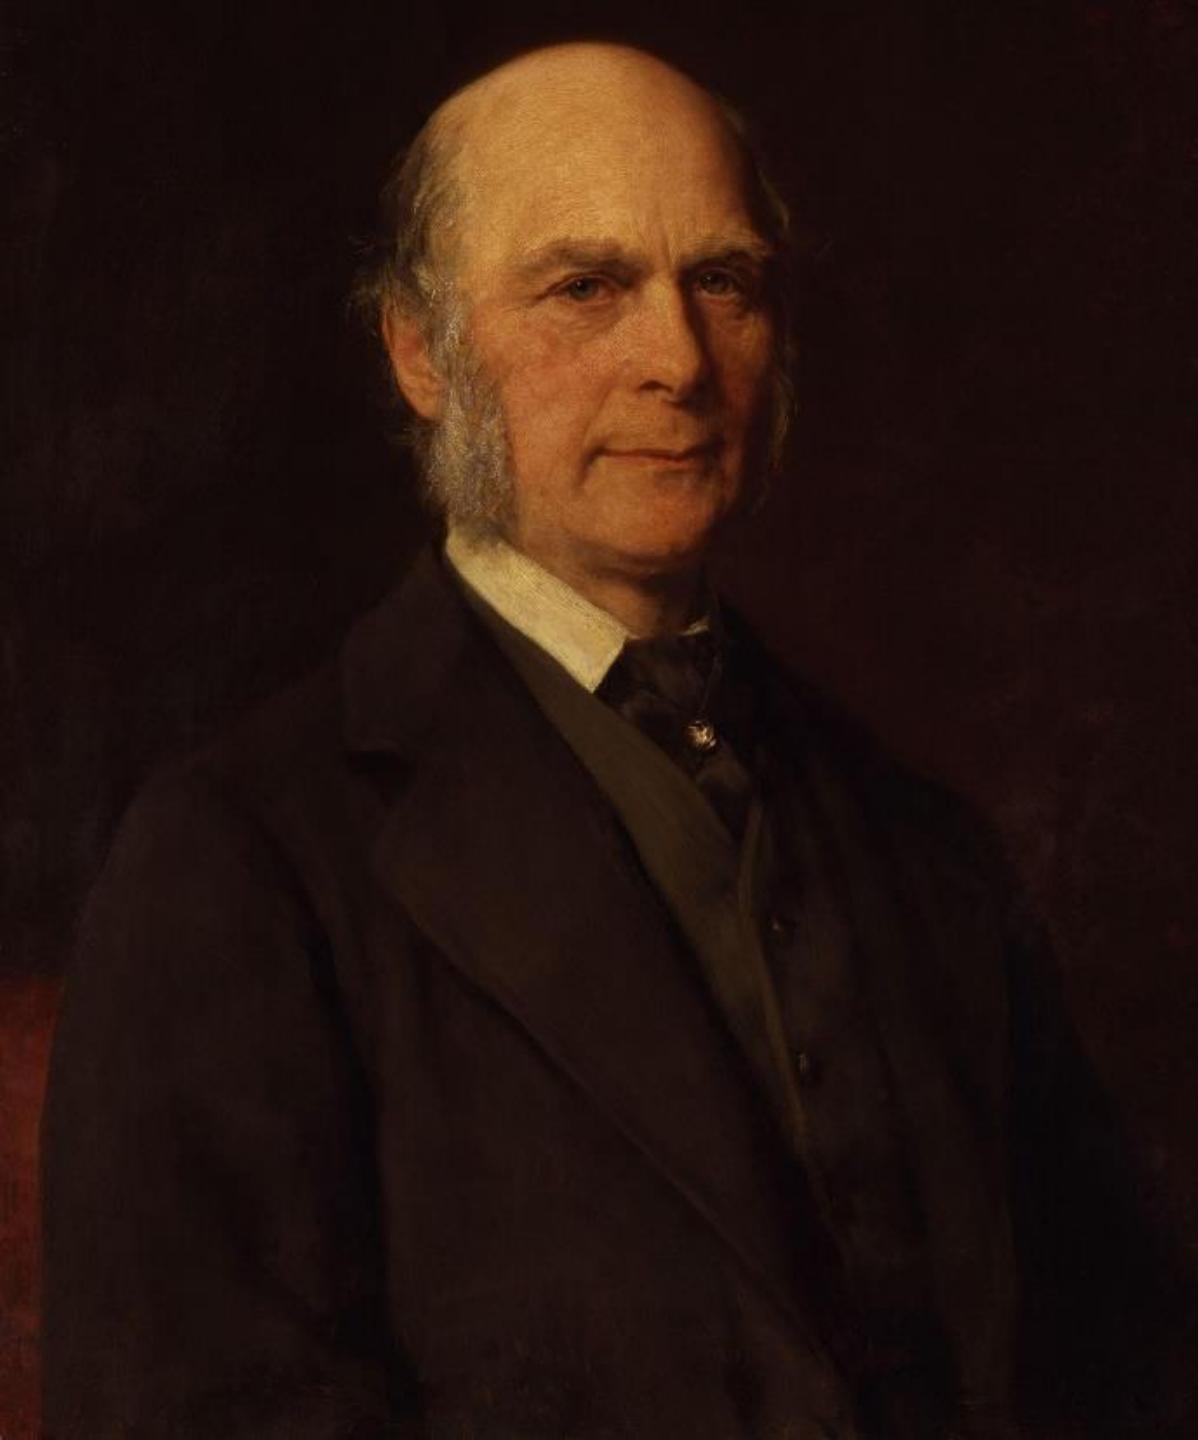
# MULTILAYER PERCEPTRON



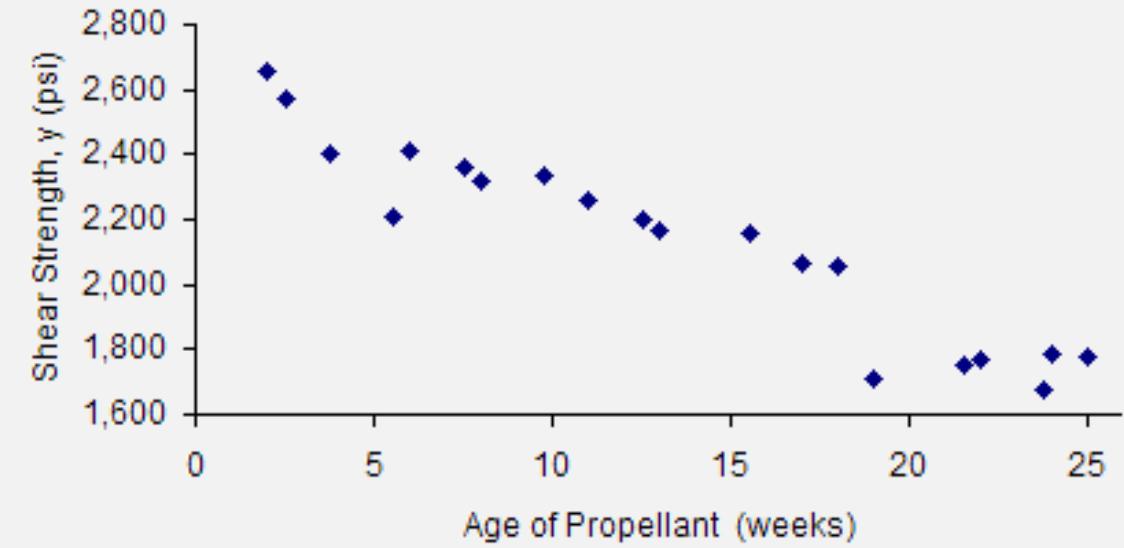
WINTER  
IS COMING



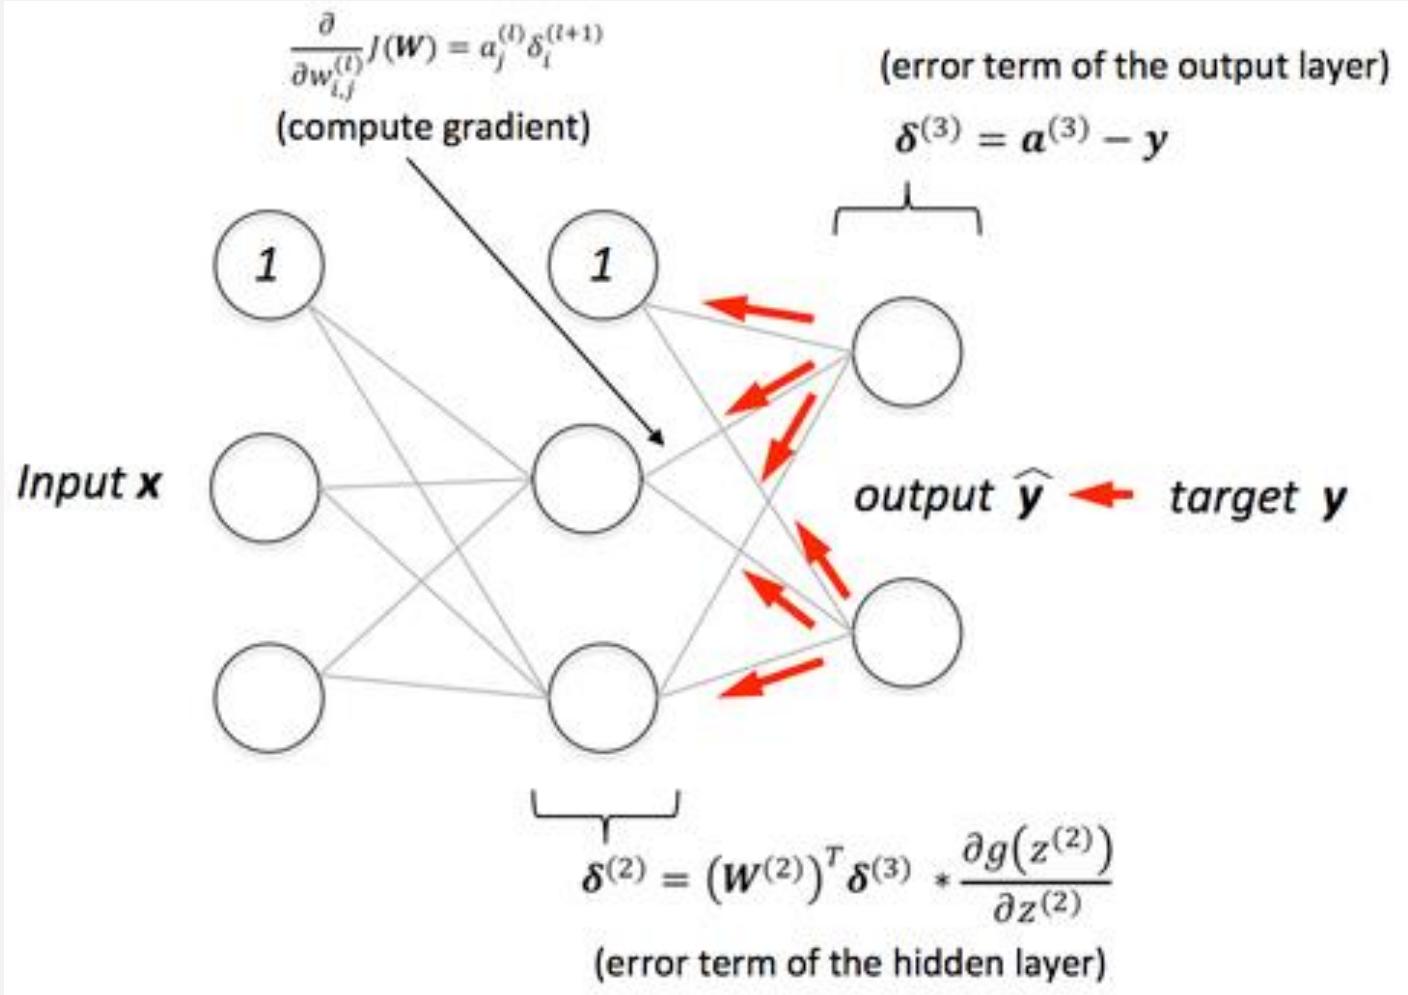
# EXPERT SYSTEMS?



# MACHINE LEARNING



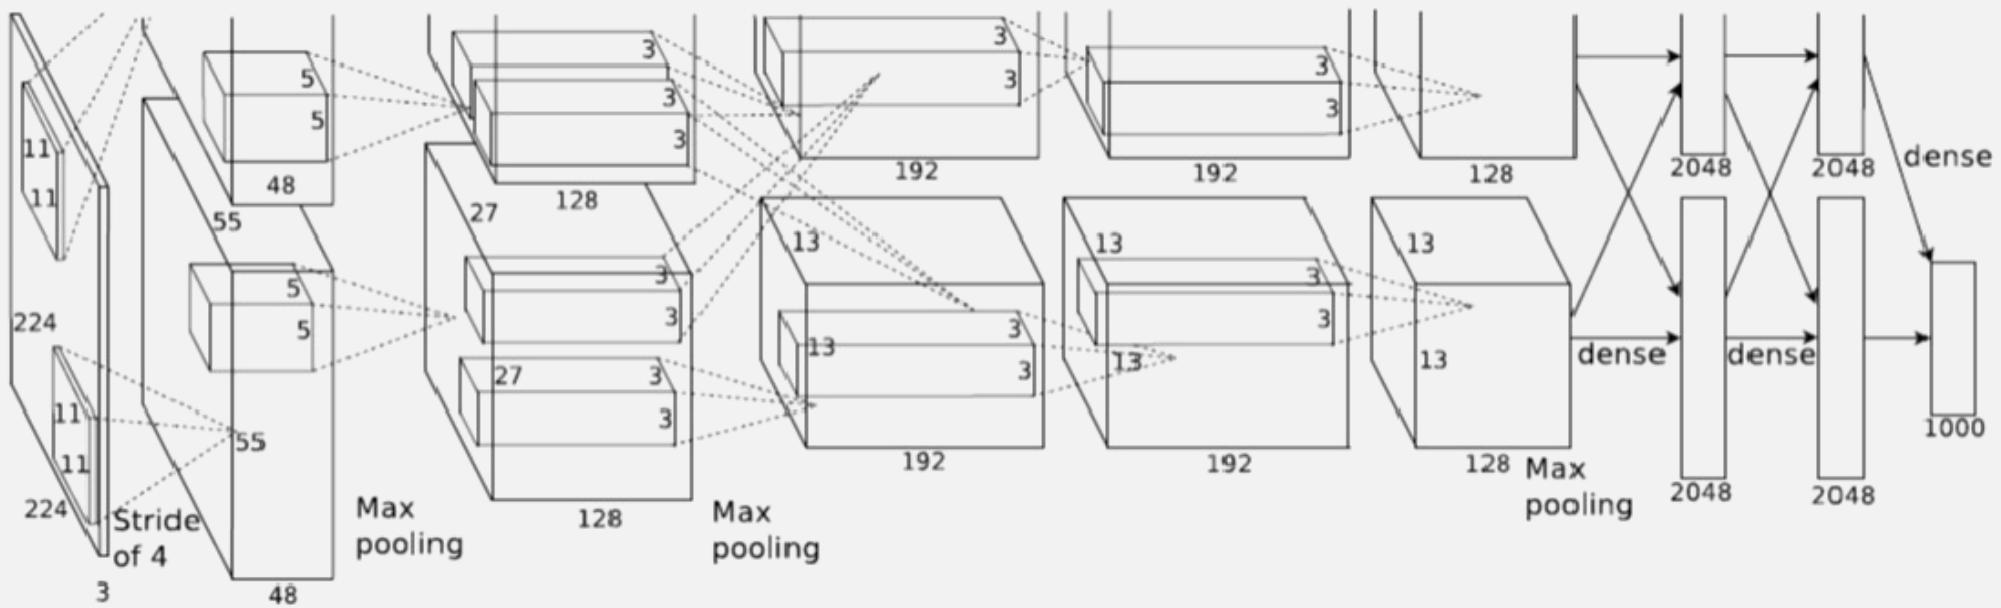
# BACKPROPAGATION

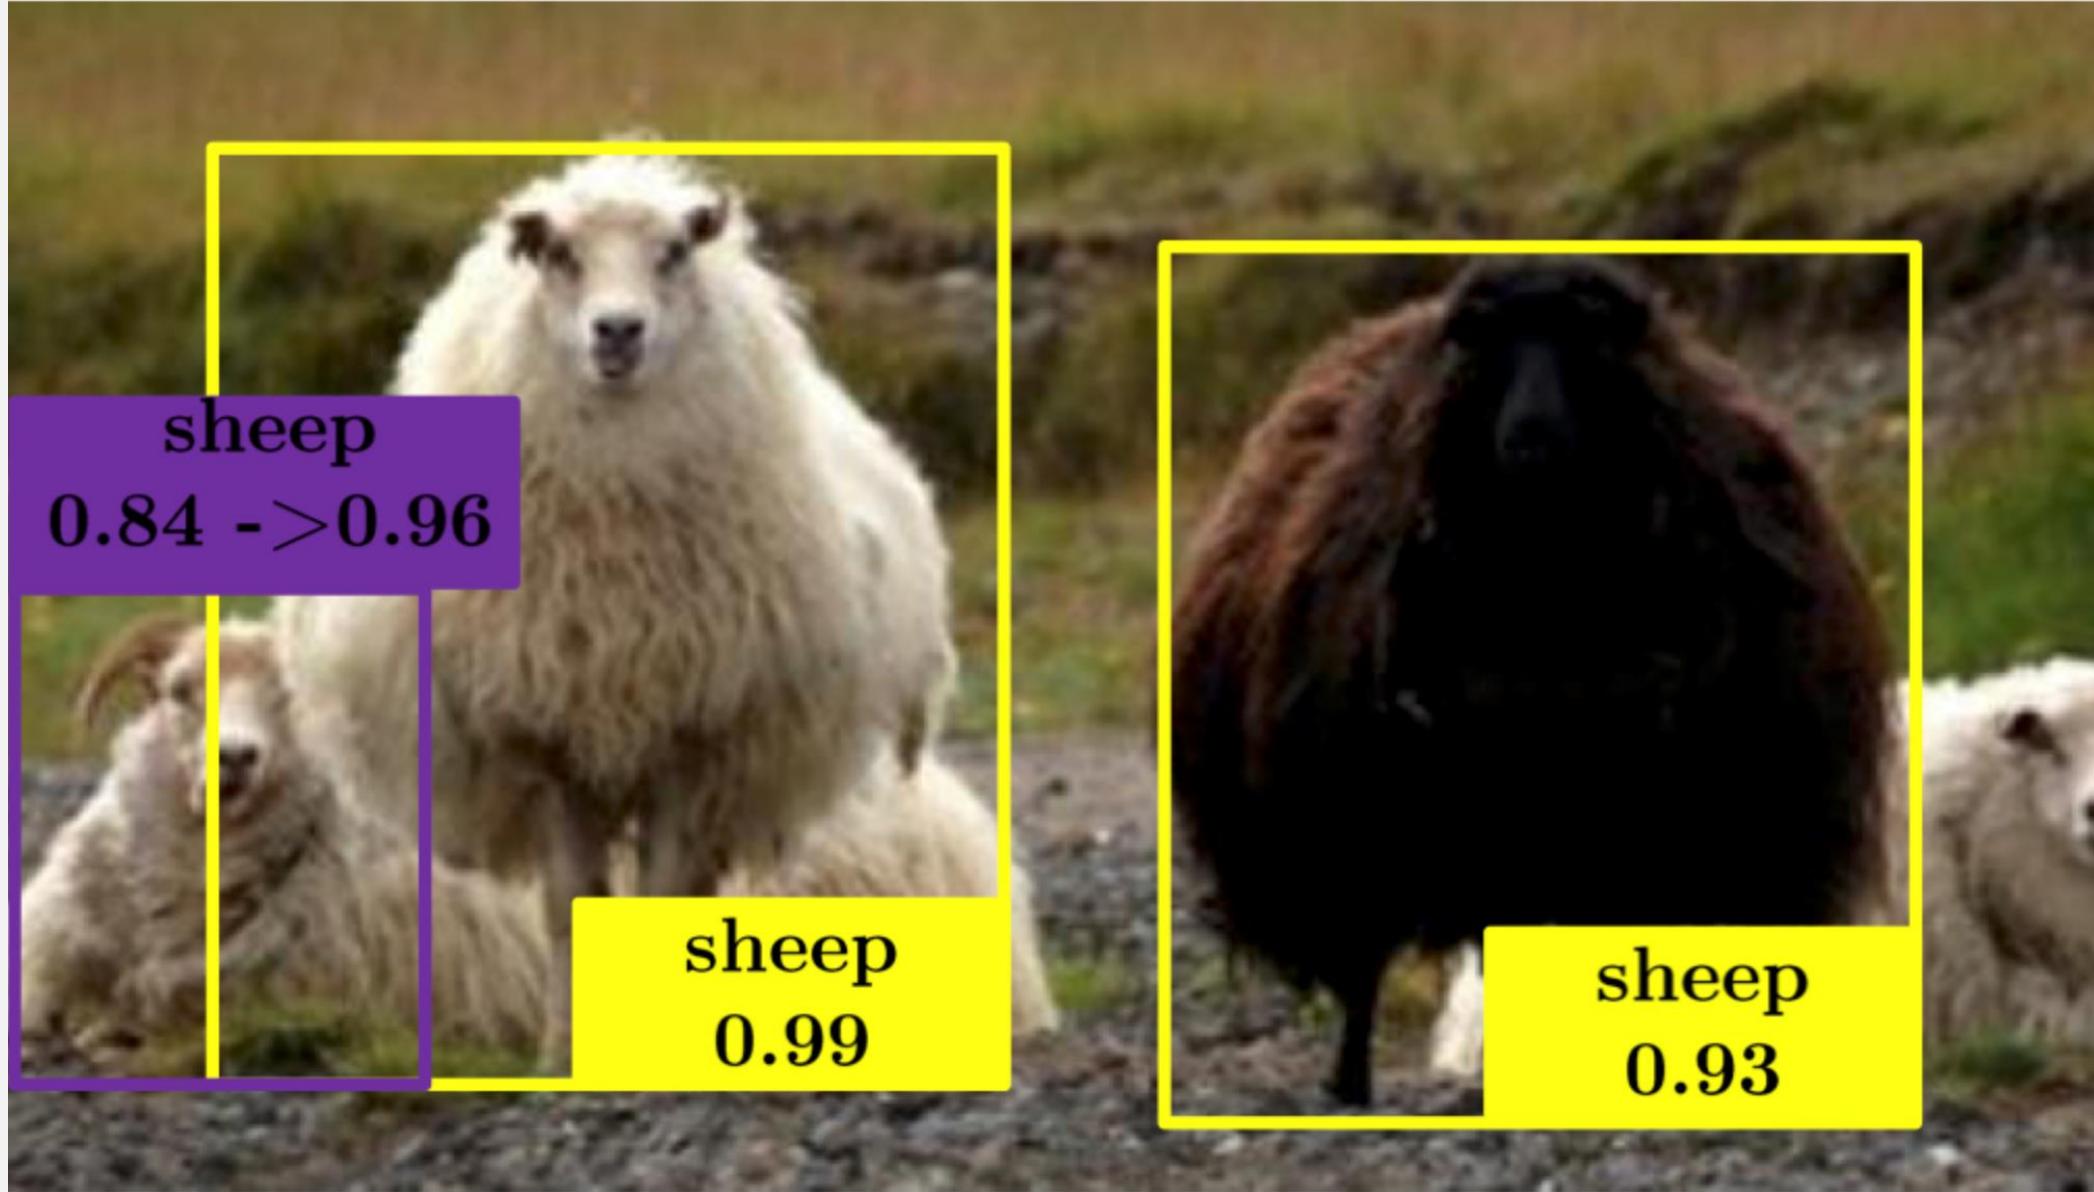


# DEEP LEARNING

A Machine Learning technique

---











# LAB 0: Setup

# CHECKLIST

---

- CUDA 9
  - >= 5.2
  - Python 3.6
  - VS Code (or Atom, ...)
  - Jupyter
- <https://github.com/PabloDoval/odsc18-london>
- 
- The diagram illustrates the checklist items grouped into four categories: CUDA, Anaconda, IDE, and Repository. Each category is preceded by a black text label and followed by a blue curly brace that spans the list of items for that category. A horizontal blue line is positioned above the IDE and Repository sections.
- CUDA
  - Anaconda
  - IDE
  - Repository

# ODSC 2018 Workshop - Deep Learning

---

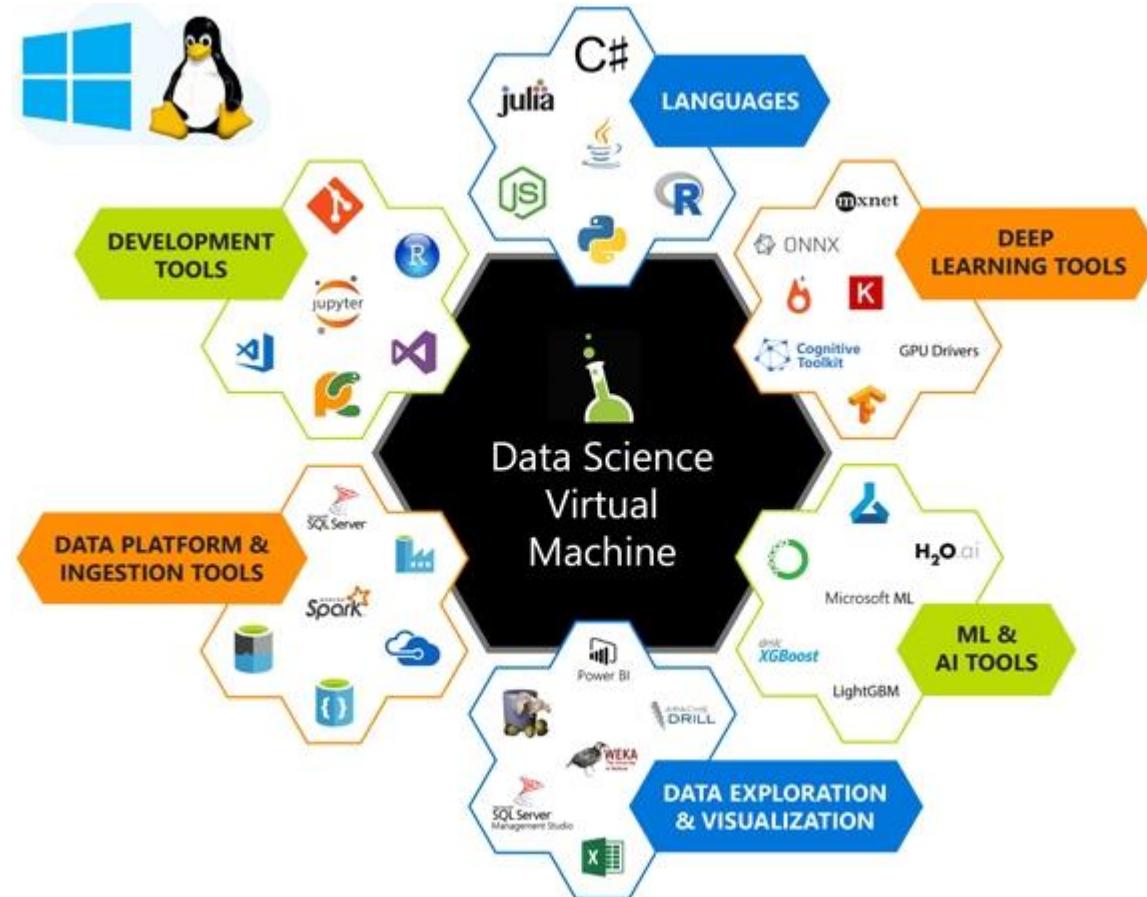
## Initial Setup

---

1. First of all, please do install Anaconda 5.2 in your machine
  - Windows version [here](#).
  - Linux version [here](#).
  - MacOS X version [here](#).
2. Go to the `.\Scripts` folder and run the script that will perform the environment setup:
  - `environment.bat` (if you are running Windows)
  - `./environment.sh` (for you, Linux folks)
3. Once the script finishes running, that environment will be active. The name of the environment is `odsc18-deeplearningworkshop`, in case you want to activate it manually.
4. If you plan on using GPU - and you should! - install the CUDA toolkit from [this link](#) to the NVIDIA web page.
5. Run the `main.py` to check that the environment and dependencies have been properly setup. The output will show the version of the Tensorflow and Keras distributions, as well as the number of CPUs and GPUs available in our machine.

# Azure DataScience Virtual Machines

Pre-installed installed images, configured and tested with popular tools that are commonly used for data analytics, machine learning and AI.



<https://azure.microsoft.com/en-gb/services/virtual-machines/data-science-virtual-machines>

# Deep Learning with Tensorflow

---

# DEEP LEARNING

---

Machine Learning

---

Based on Neural Networks

---

More than one hidden layer (*Deep Models*)

---

# Why?

- 
- Interface for expressing machine learning models
  - Implementation for executing such models
  - Handles multiple CPUs and GPUs
  - Multi-platform

# TensorFlow

---

# TensorFlow

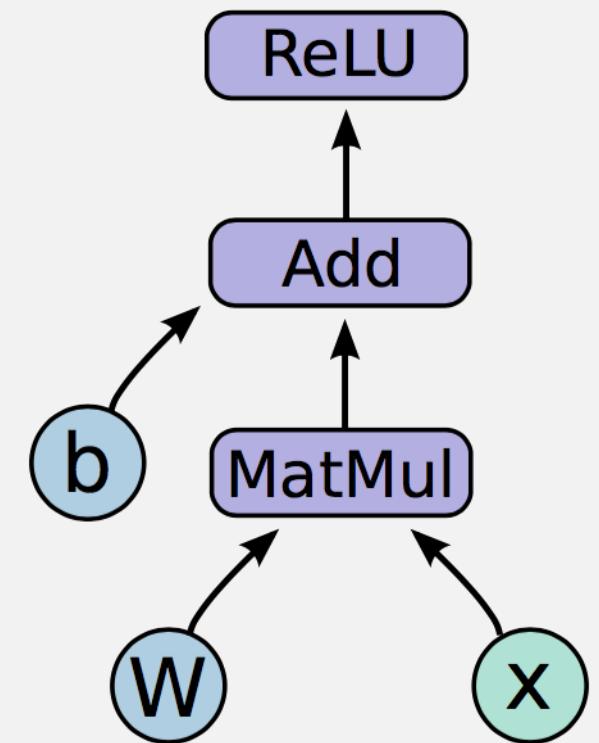
---

- Express a numeric computation as a **graph**
- Graph nodes are **operations** which have any number of inputs and outputs
- Graph edges are **tensors** which flow between nodes
- Development in **Python**

# TensorFlow

---

$$h_i = \text{ReLU}(Wx + b)$$

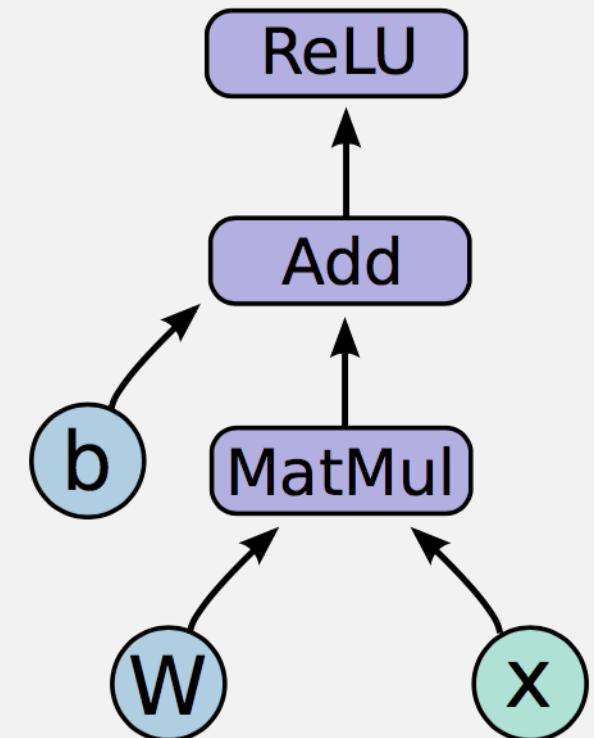


# TensorFlow

---

$$h_i = \text{ReLU}(Wx + b)$$

- **Variables:**
  - Nodes which emit their value as output
  - State is persisted, and retained across multiple executions of the graph
  - Examples: parameters



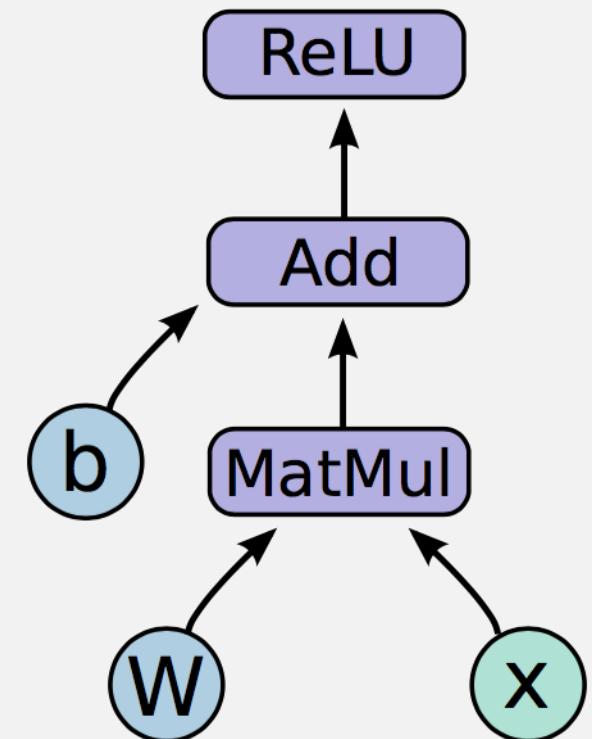
# TensorFlow

---

$$h_i = \text{ReLU}(Wx + b)$$

- **Placeholders:**

- Values are fed at execution time
- Examples: inputs, hyper-parameters

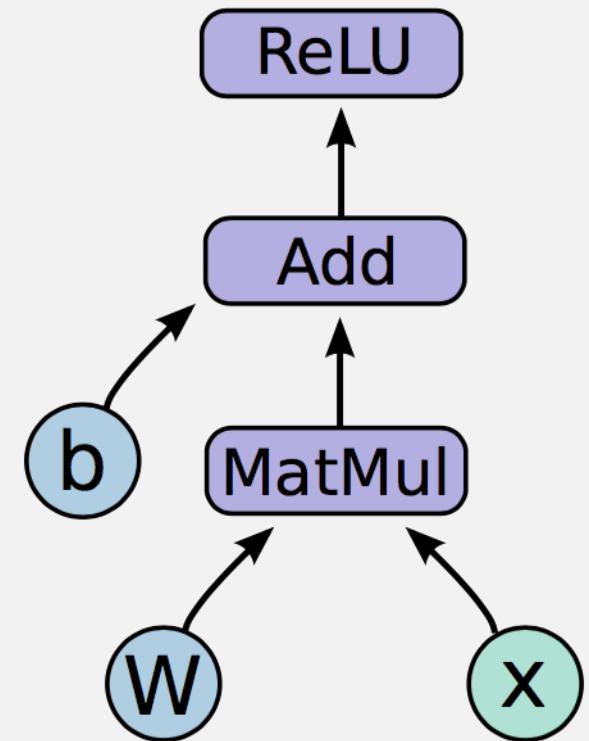


# TensorFlow

---

$$h_i = \text{ReLU}(Wx + b)$$

- **Operations:**
  - Perform the specific operation with the tensor input stream
  - Examples: MatMul, Add, ReLU



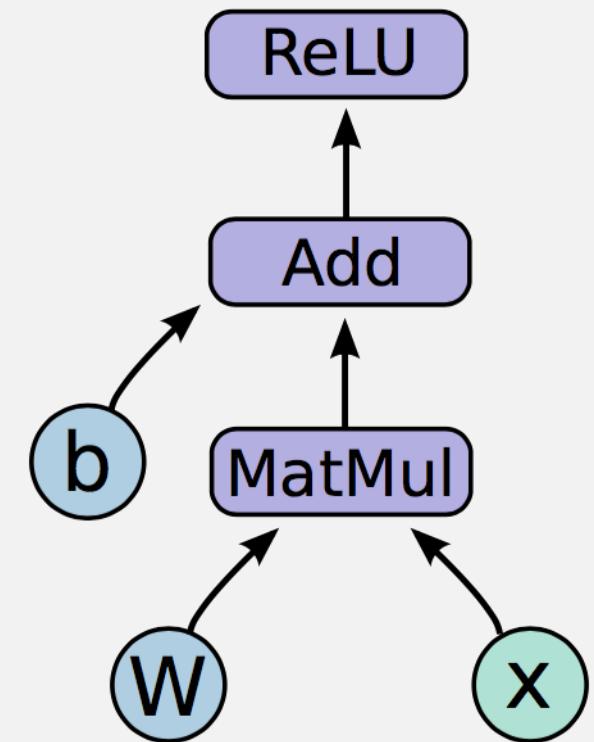
# TensorFlow – Graph Code

---

```
import tensorflow as tf

b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))

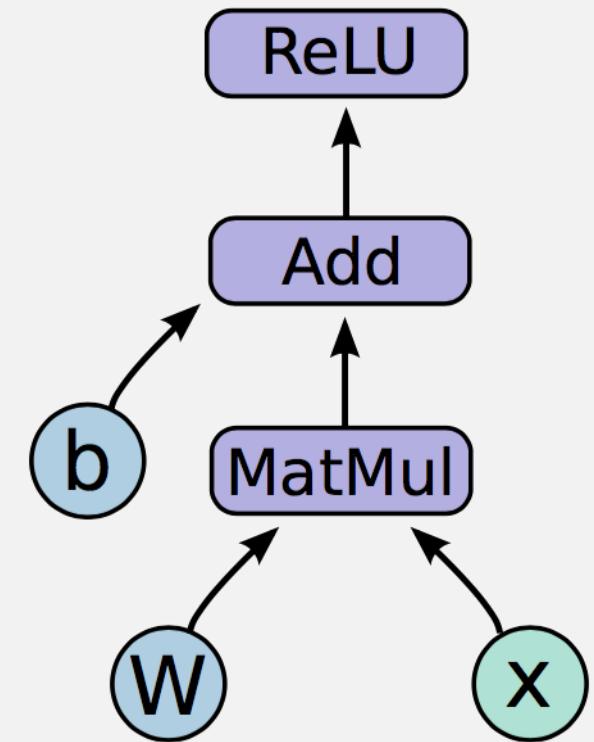
x = tf.placeholder(tf.float32, (None, 784))
h_i = tf.nn.relu(tf.matmul(x, W) + b)
```



# TensorFlow – Graph Execution

---

- **Session:**
  - Once we have defined our **graph**, we need to execute it
  - We can deploy the graph to a **session**
  - These sessions can be binded to an **execution context** (GPU, CPU, ...)



# TensorFlow – Execution Code

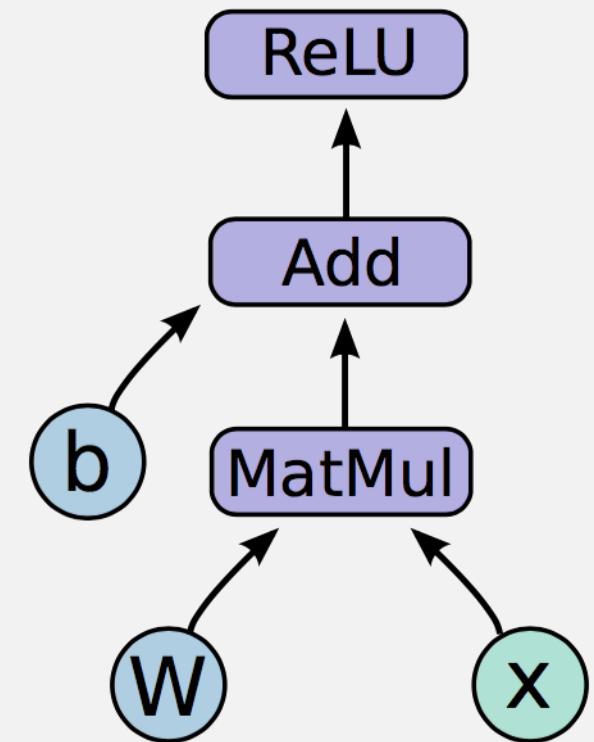
---

```
import tensorflow as tf

b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))

x = tf.placeholder(tf.float32, (None, 784))
h_i = tf.nn.relu(tf.matmul(x, W) + b)

sess = tf.Session()
sess.run(tf.initialize_all_variables())
sess.run(h_i, {x: np.random.random(64, 784)})
```



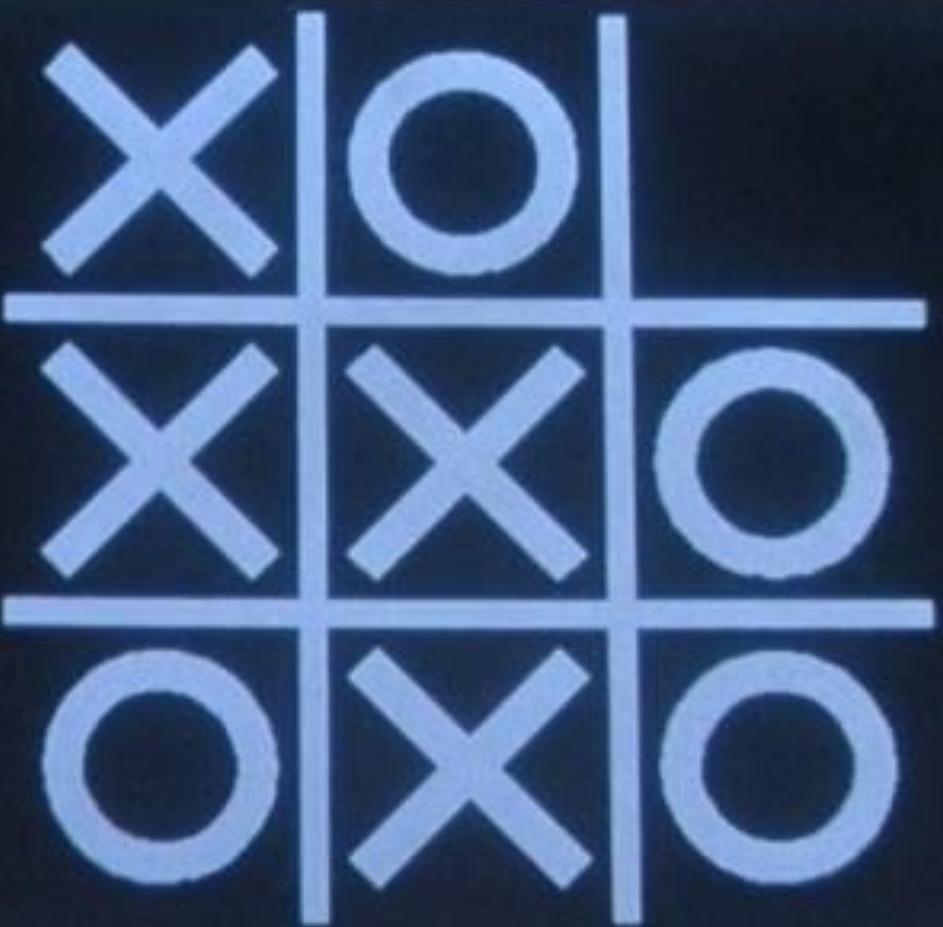
# LAB 1: TensorFlow Operations

# TensorFlow 2.0

---

- Eager Execution as default
  - No need to add `tf.enable_eager_execution()`
  - Blocks or Functions can be defined to be executed in graph mode
- Less conventions, more object-oriented and pythonic design
  - `Variable_scope` removed
- Clean up in libraries and contrib

# SHALL WE PLAY A LITTLE GAME?



YLING  
4  
COM STS  
8365

OMM STS  
3603

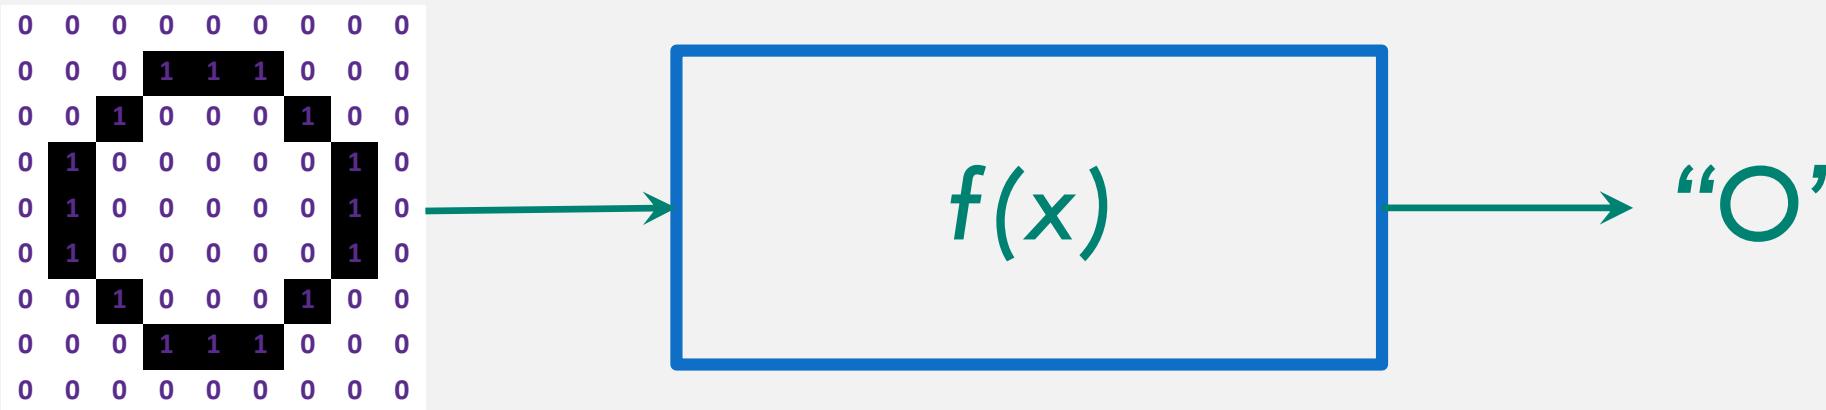
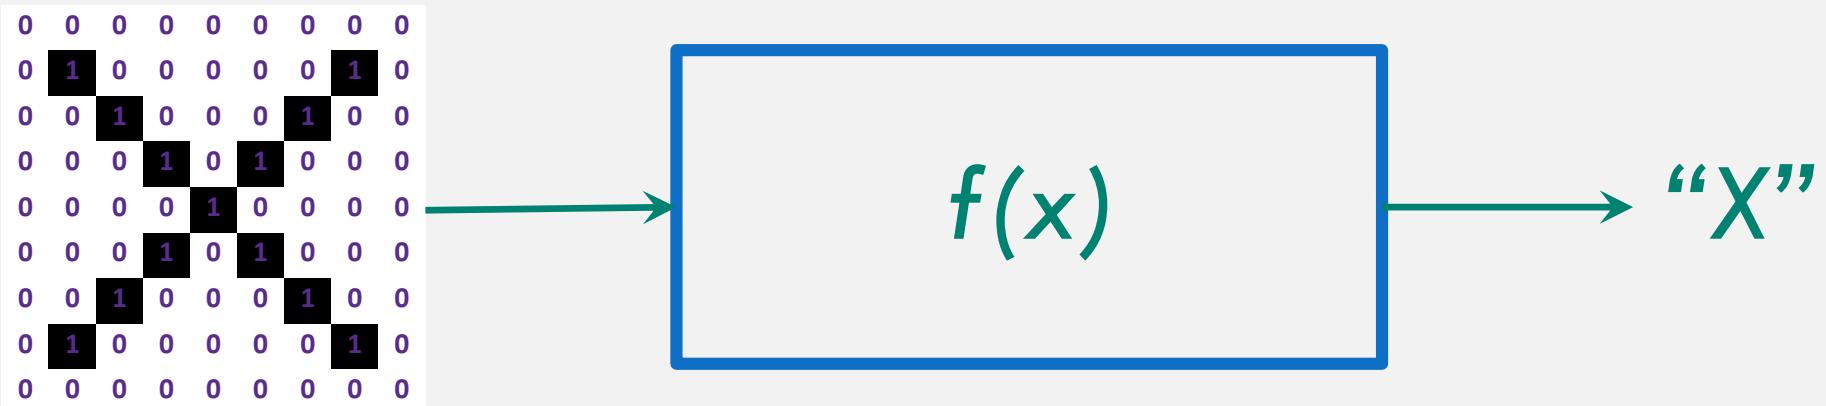
RPL STS  
9071

EOT STS  
4531

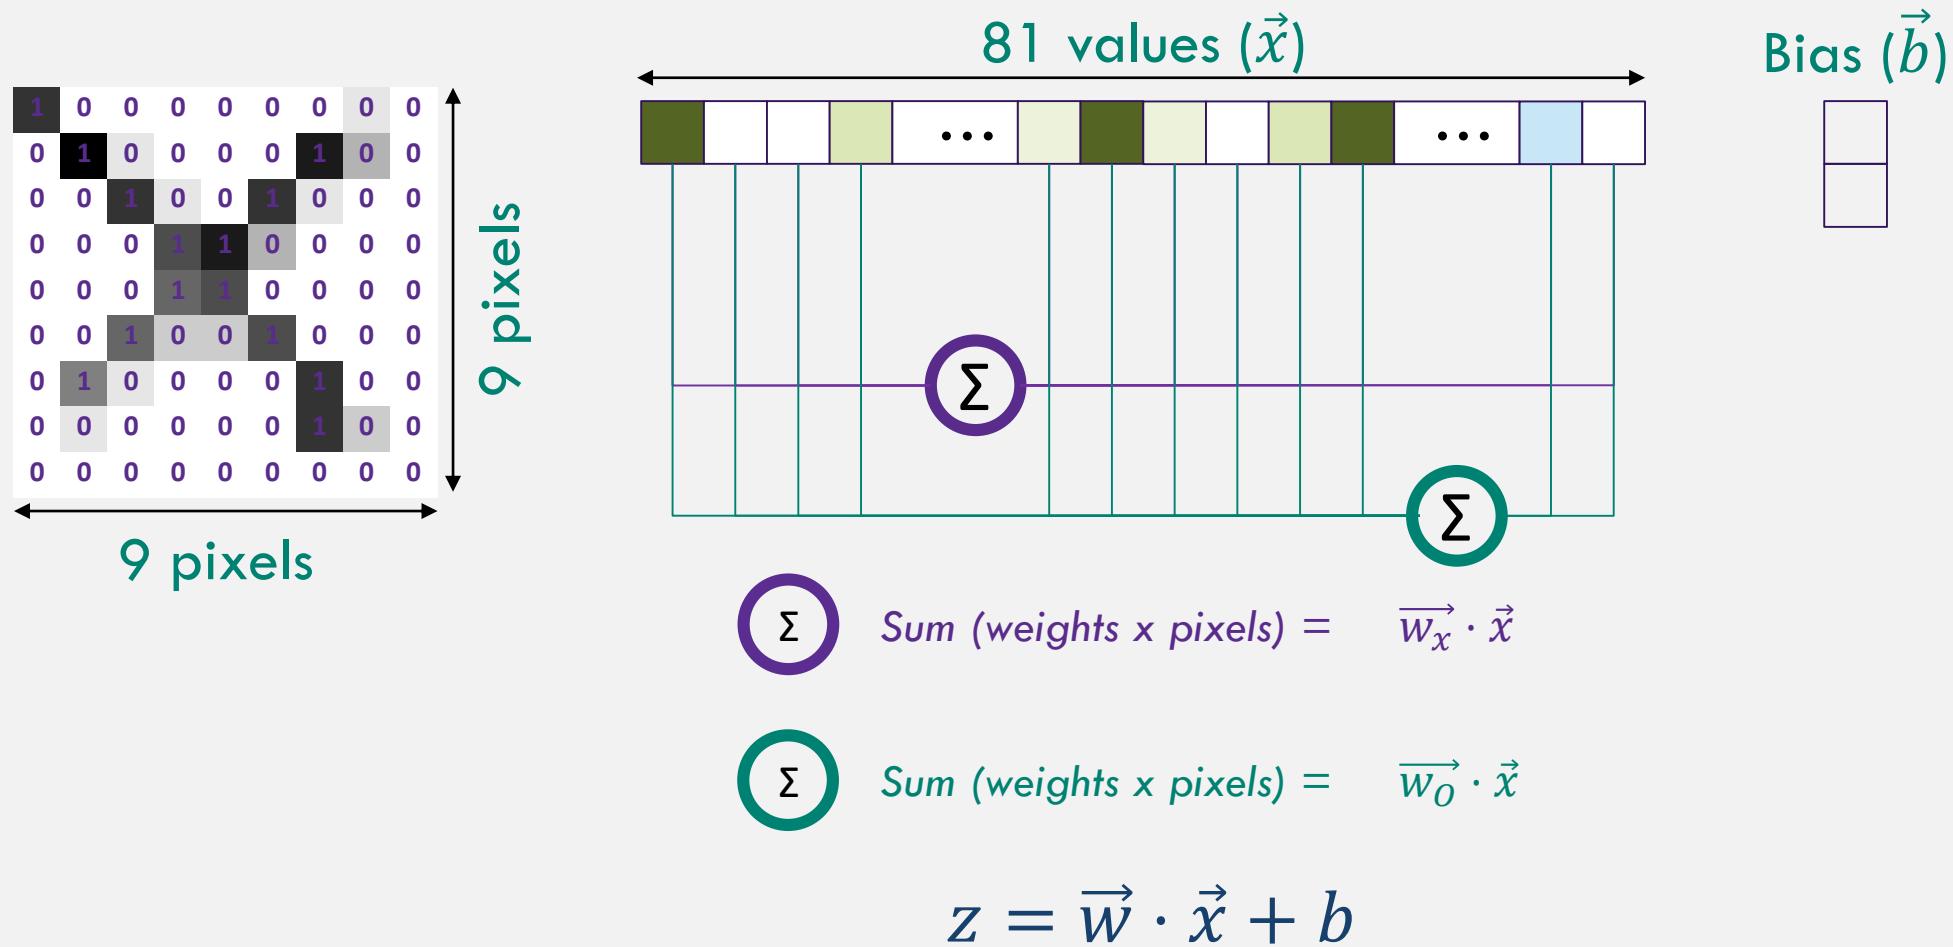
PAC STS  
8838

ADL STS  
6632

# A SIMPLE EXAMPLE

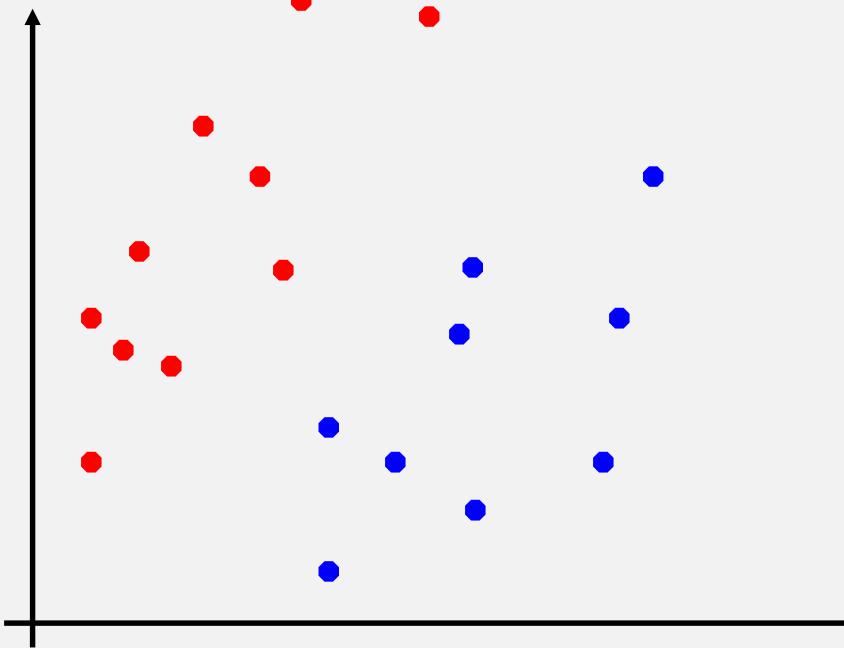


# A SIMPLE EXAMPLE



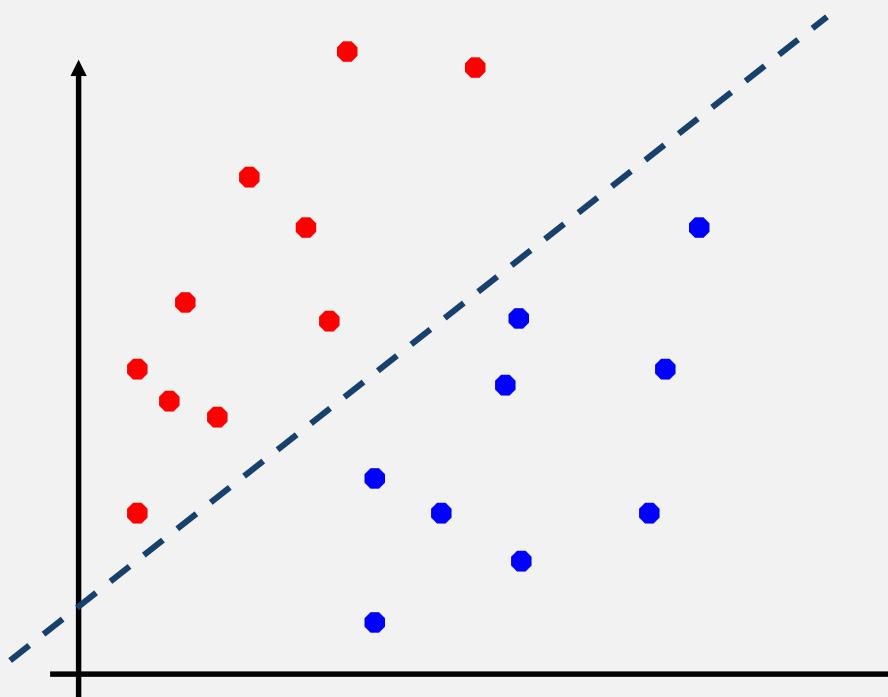
# NOW THAT YOU MENTION IT...

$$z = w \cdot \vec{x} + b$$



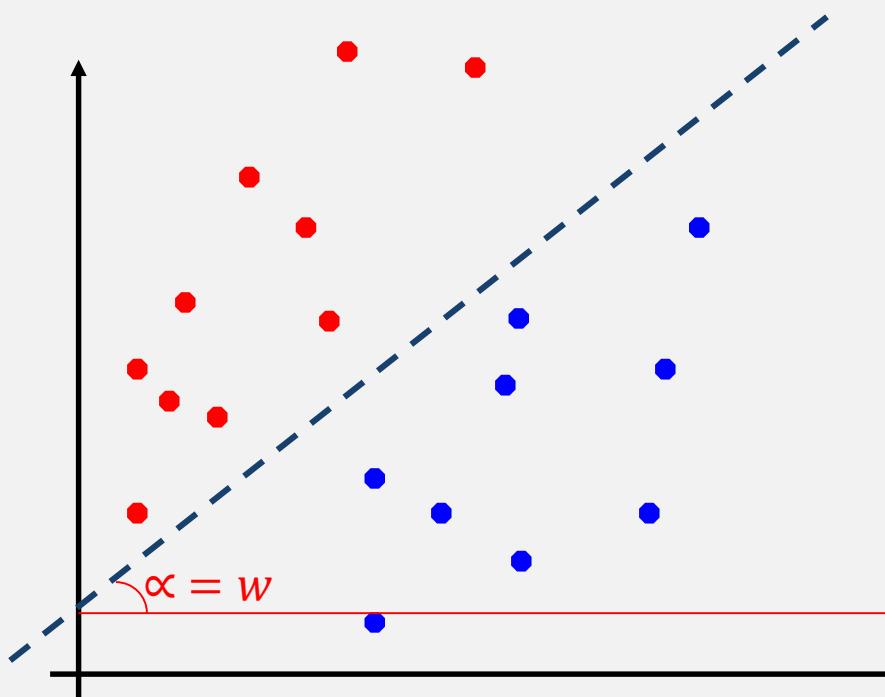
# NOW THAT YOU MENTION IT...

$$z = \mathbf{w} \cdot \vec{x} + b$$

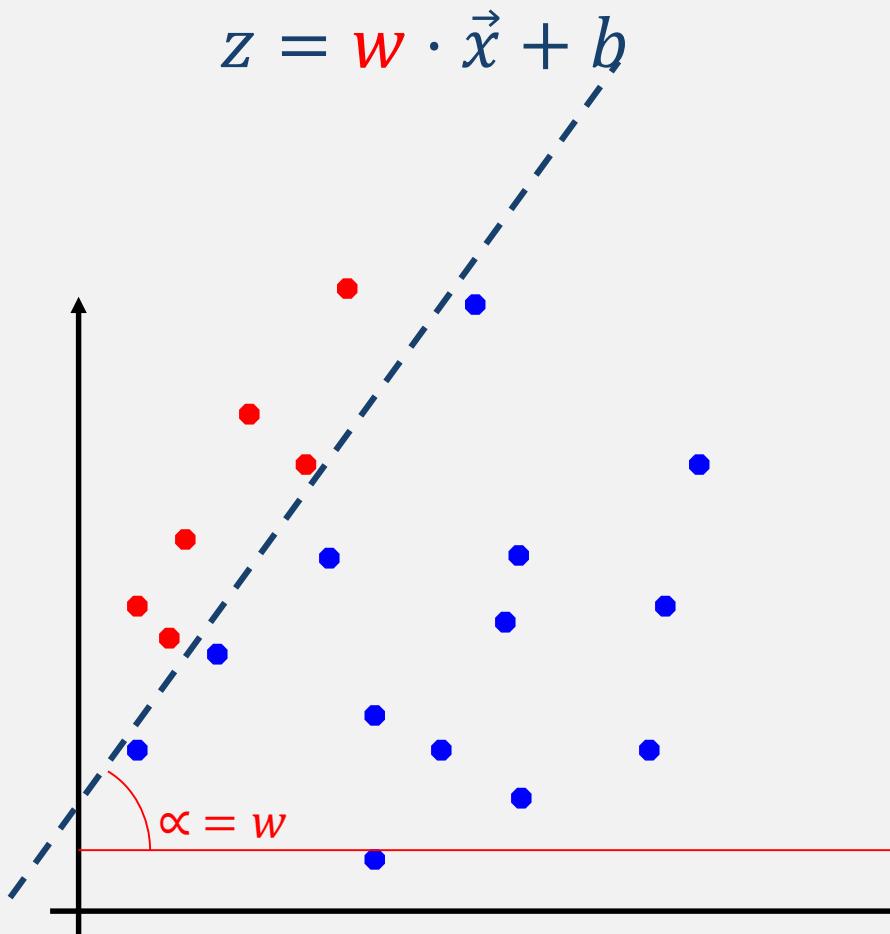


# NOW THAT YOU MENTION IT...

$$z = \mathbf{w} \cdot \vec{x} + b$$

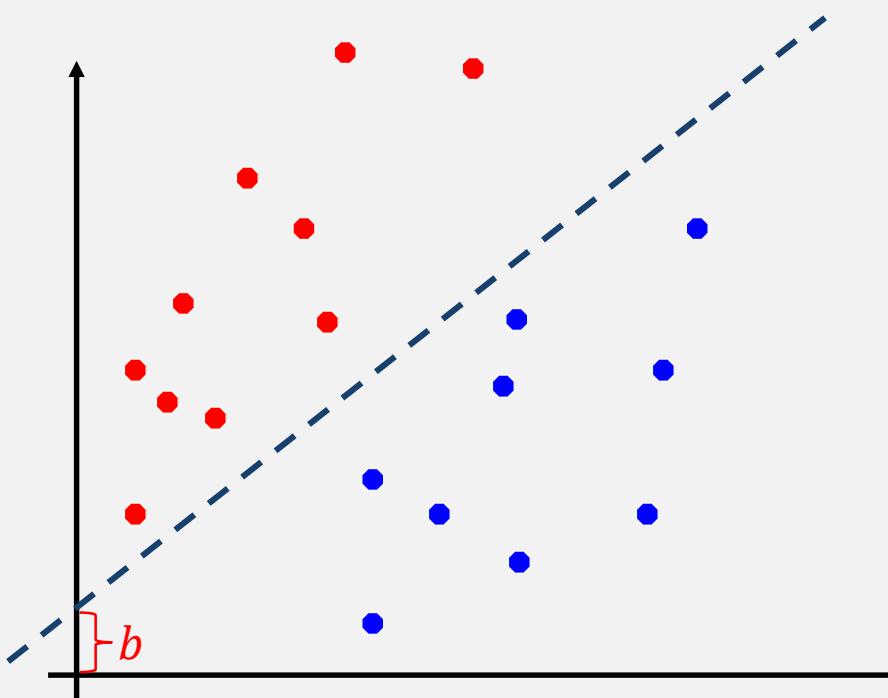


# NOW THAT YOU MENTION IT...



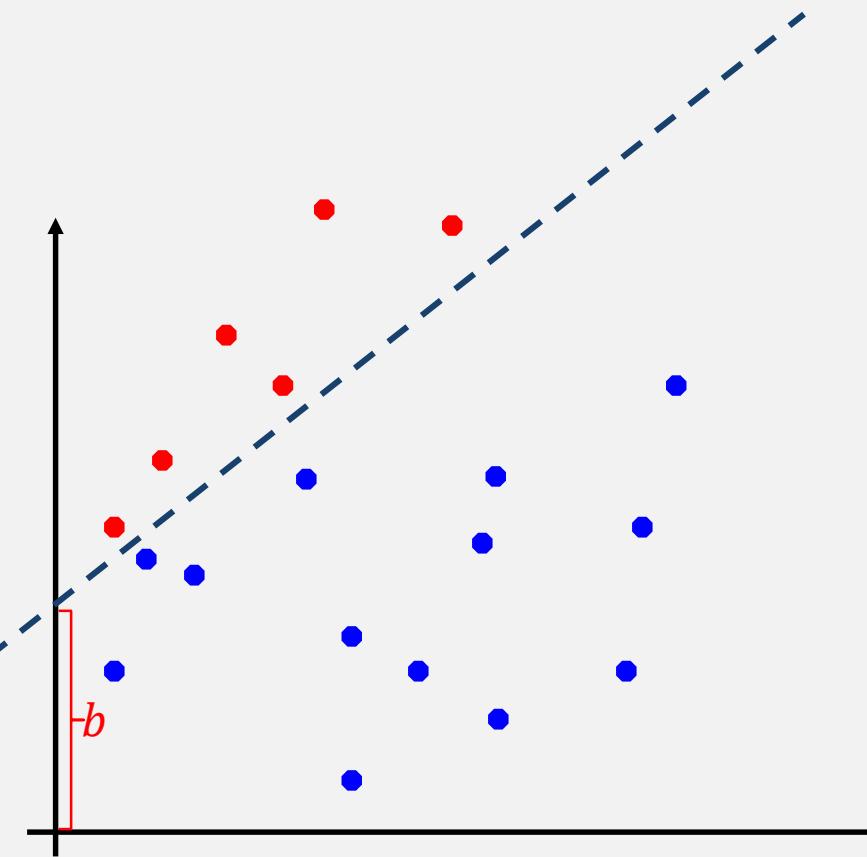
# NOW THAT YOU MENTION IT...

$$z = w \cdot \vec{x} + b$$

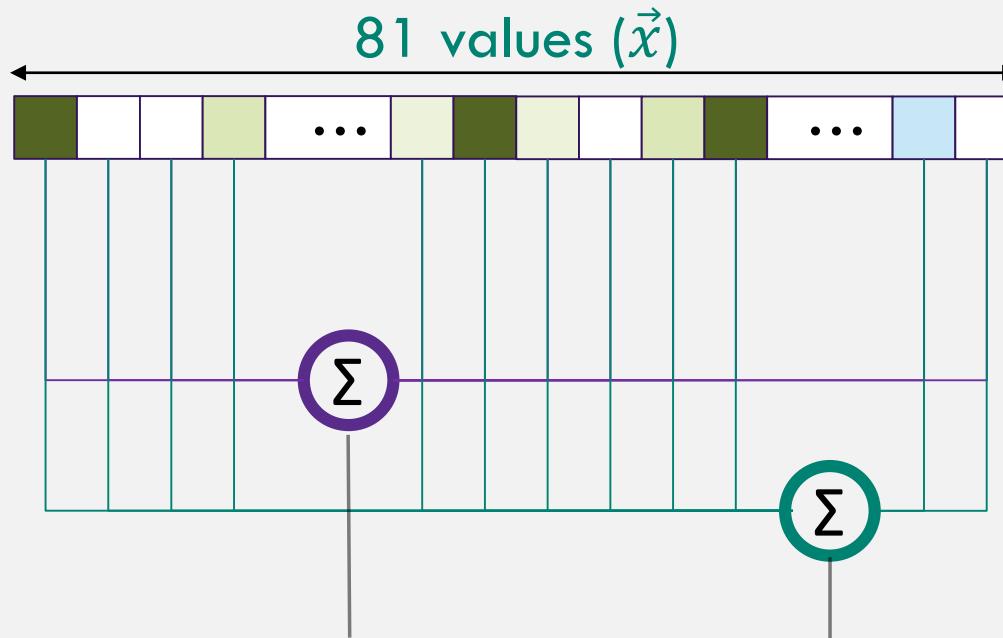
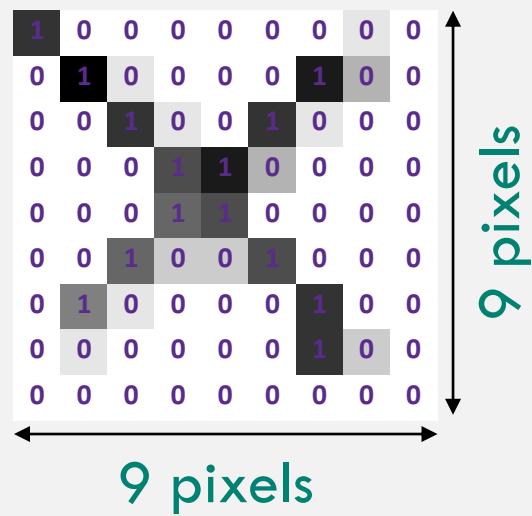


# NOW THAT YOU MENTION IT...

$$z = w \cdot \vec{x} + b$$



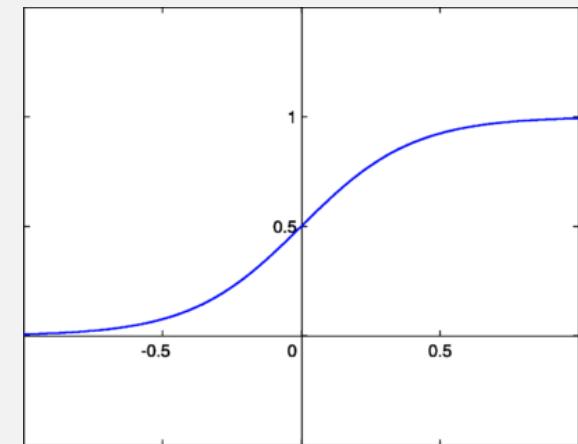
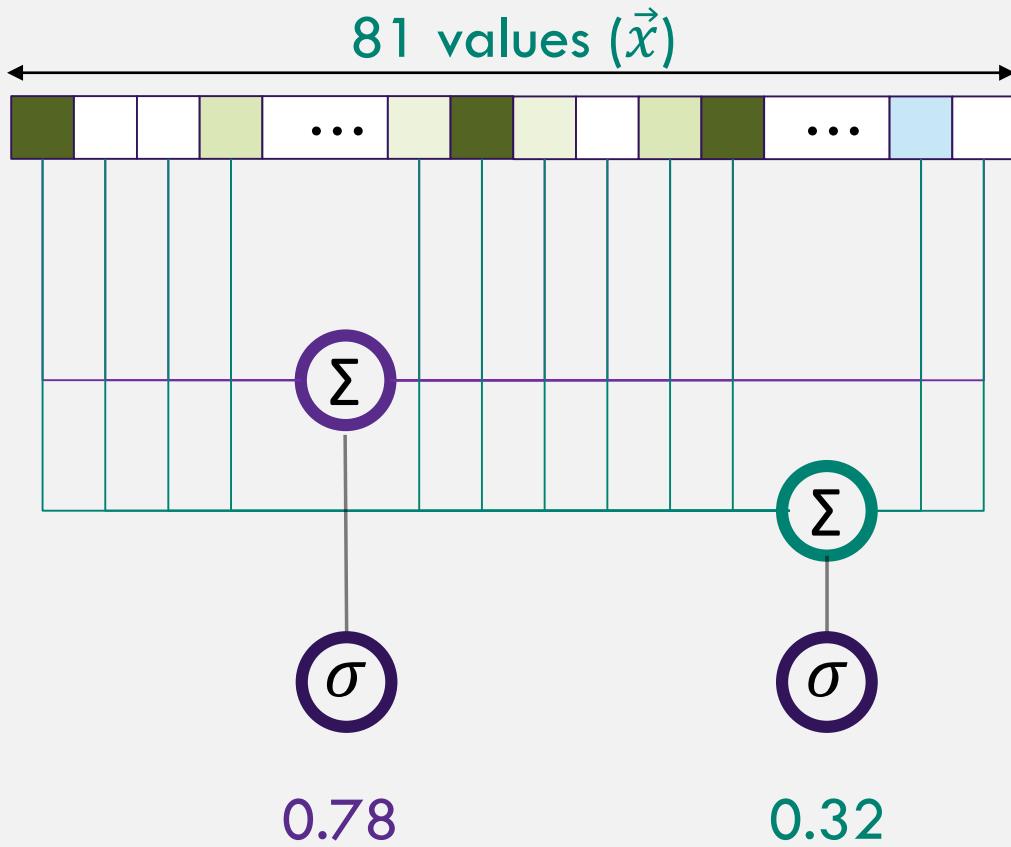
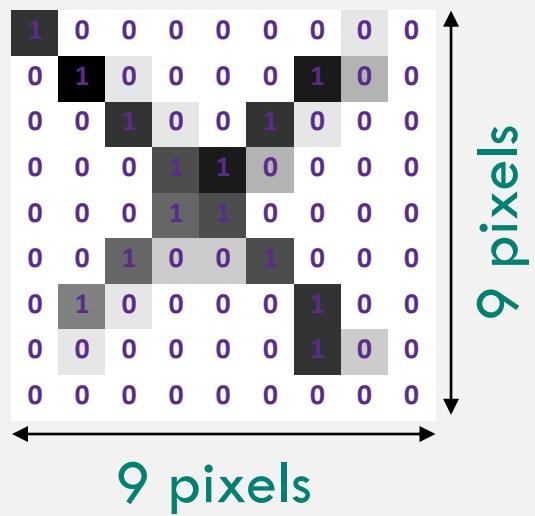
# SIGMOID



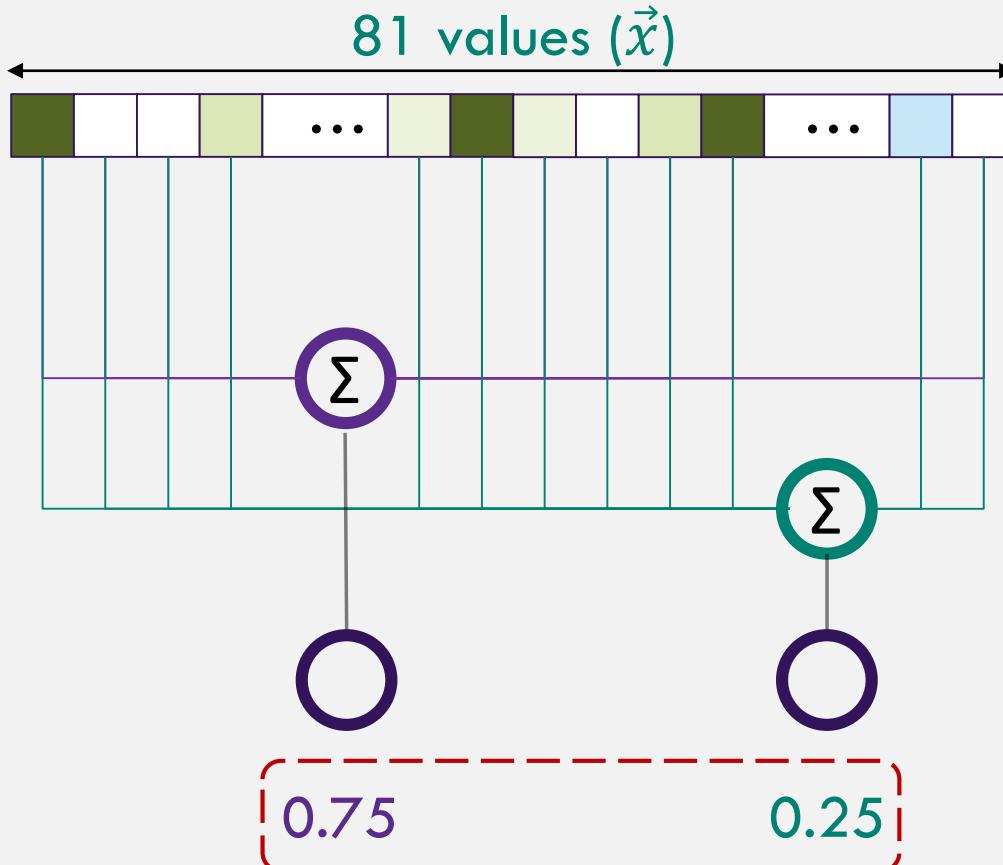
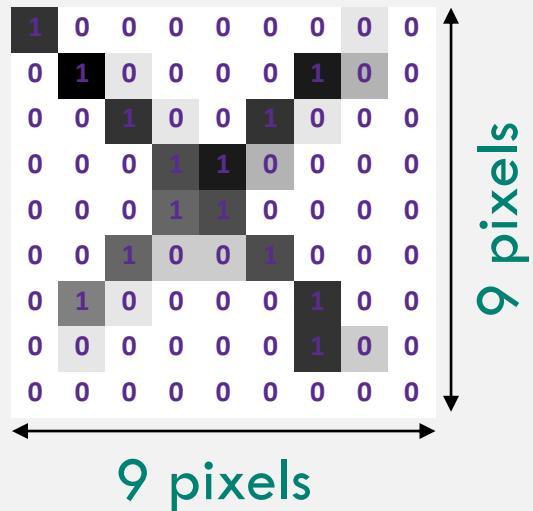
823,731.01

14,78

# SIGMOID



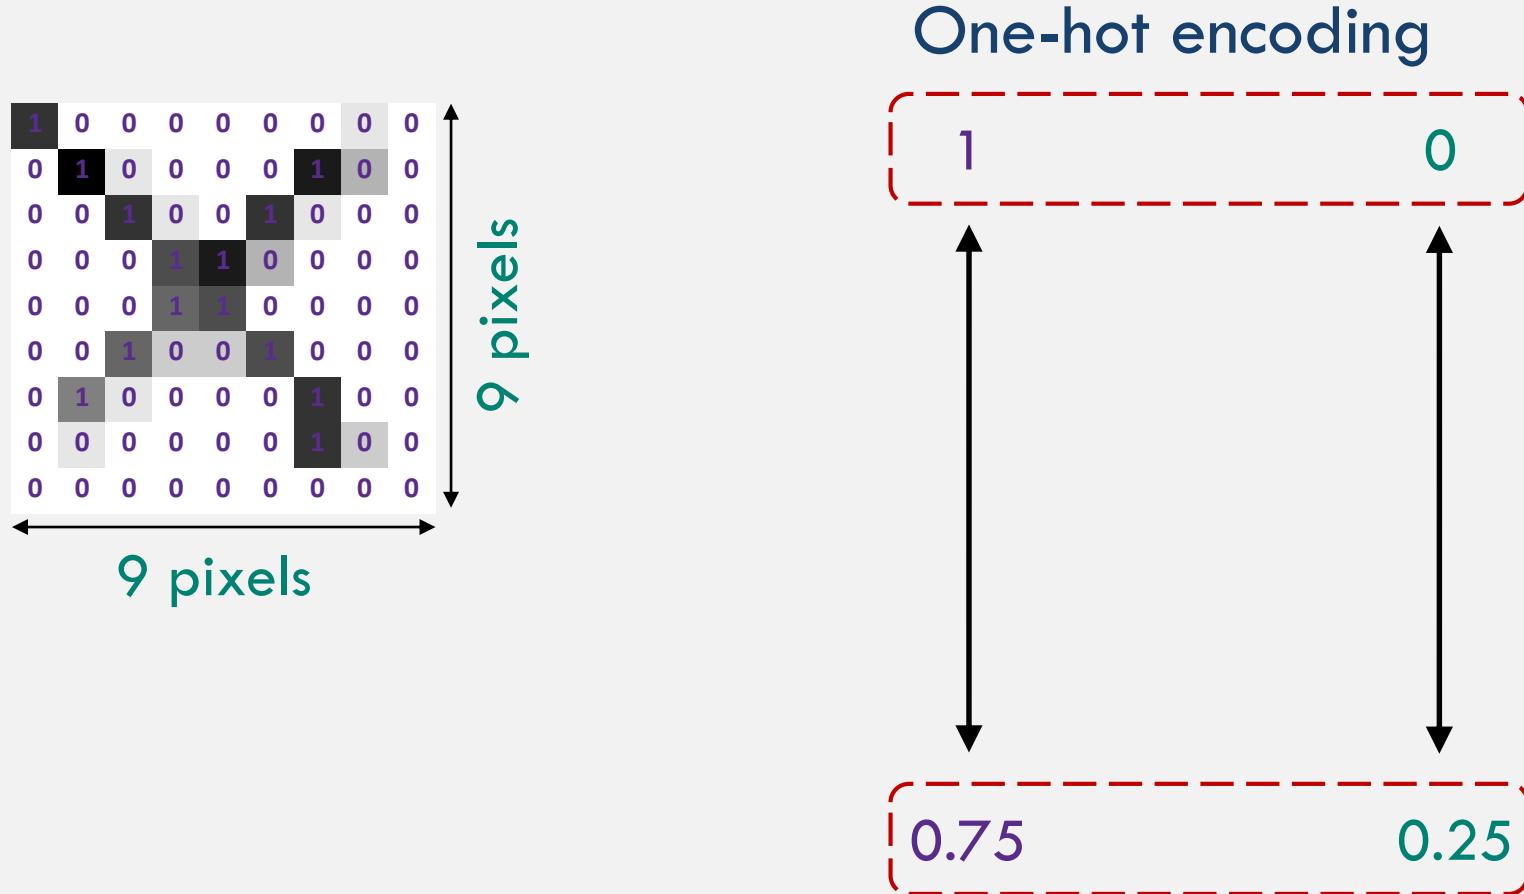
# SOFTMAX



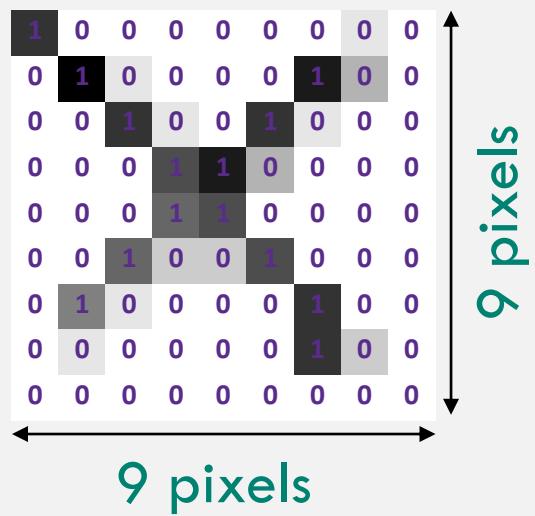
These are actual probabilities

$$s(y_i) = \frac{e^{y_i}}{\sum_j e^{y_i}}$$

# LOSS FUNCTION



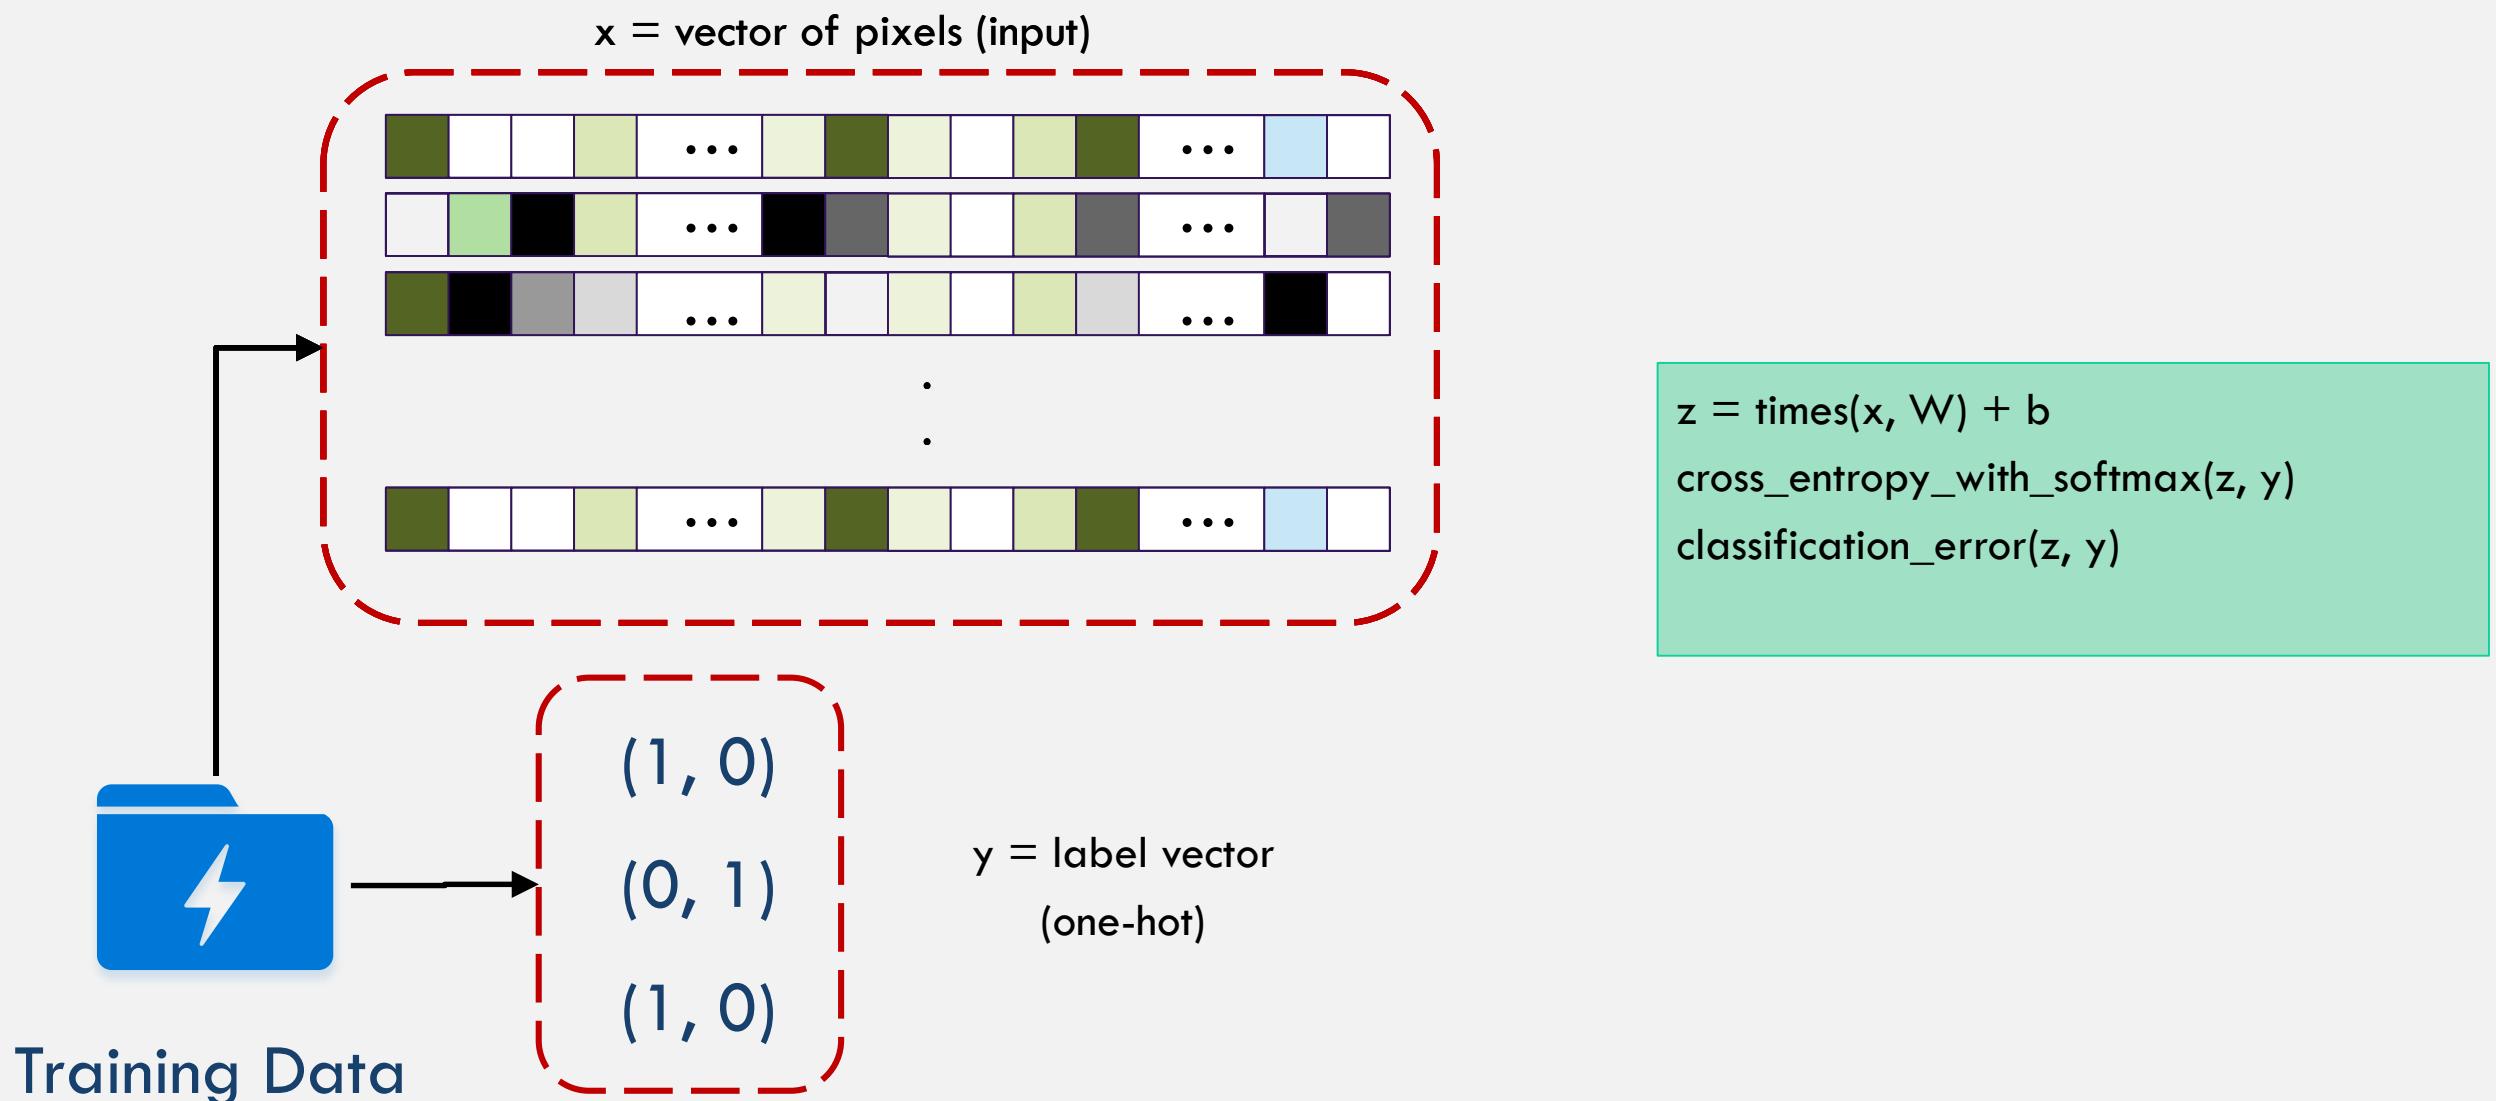
# LOSS FUNCTION



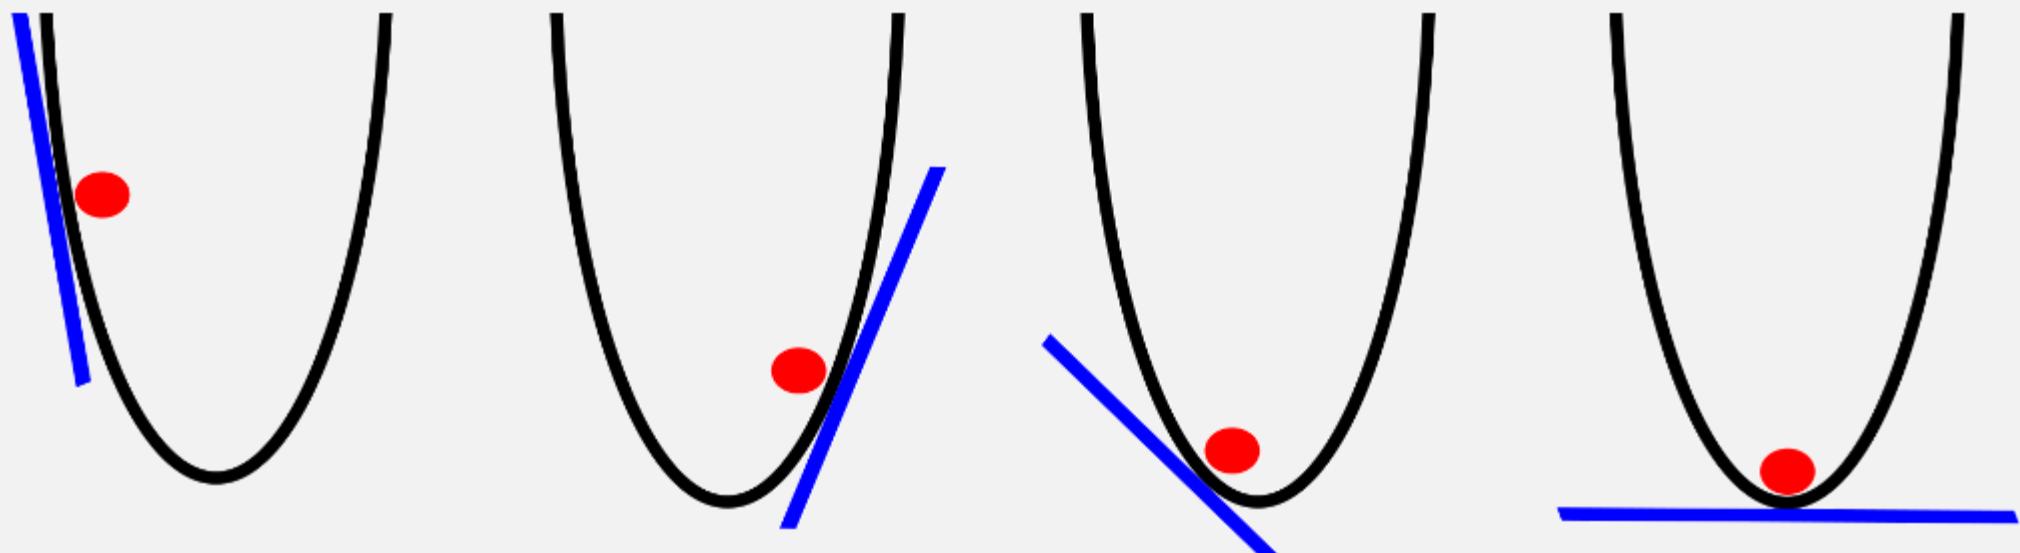
$$\text{error} = \frac{1}{n} \sum_{i=1}^n (y_i - p_i)^2$$



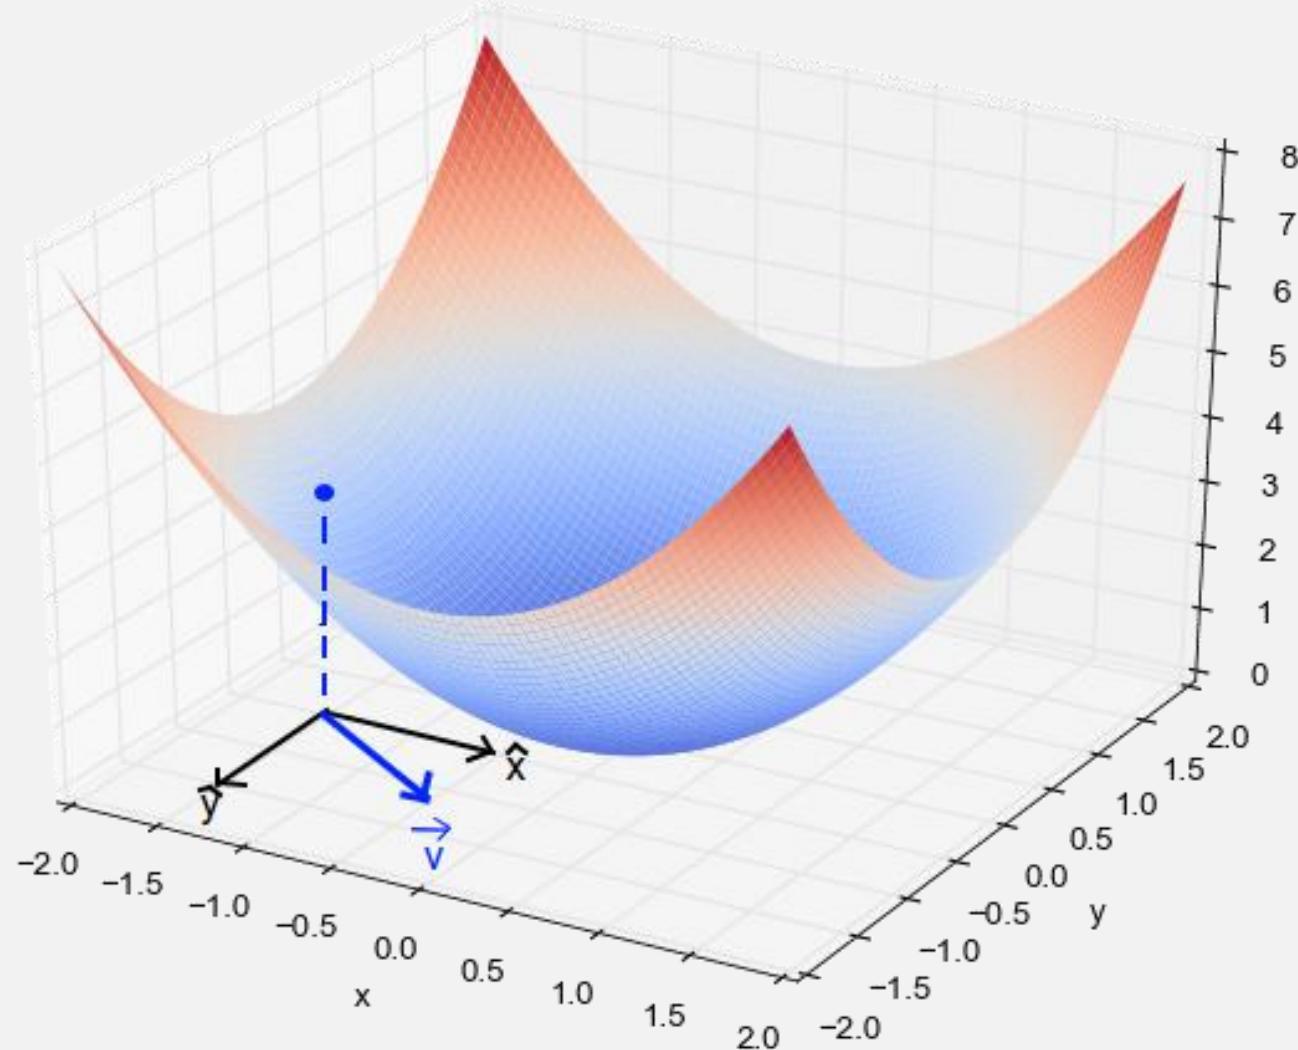
# TRAINING



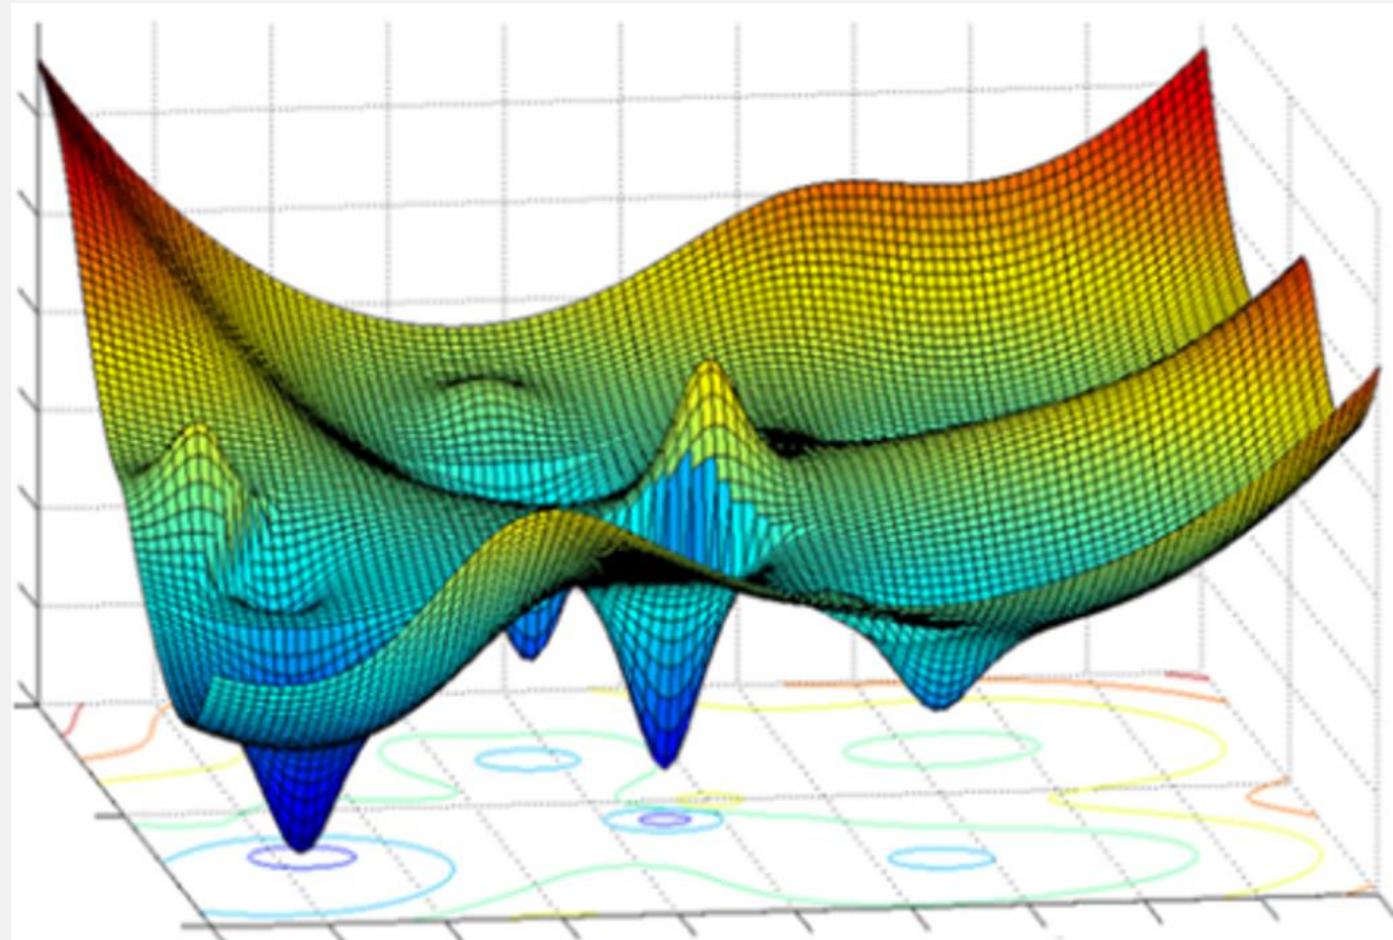
# GRADIENT DESCENT



# GRADIENT DESCENT



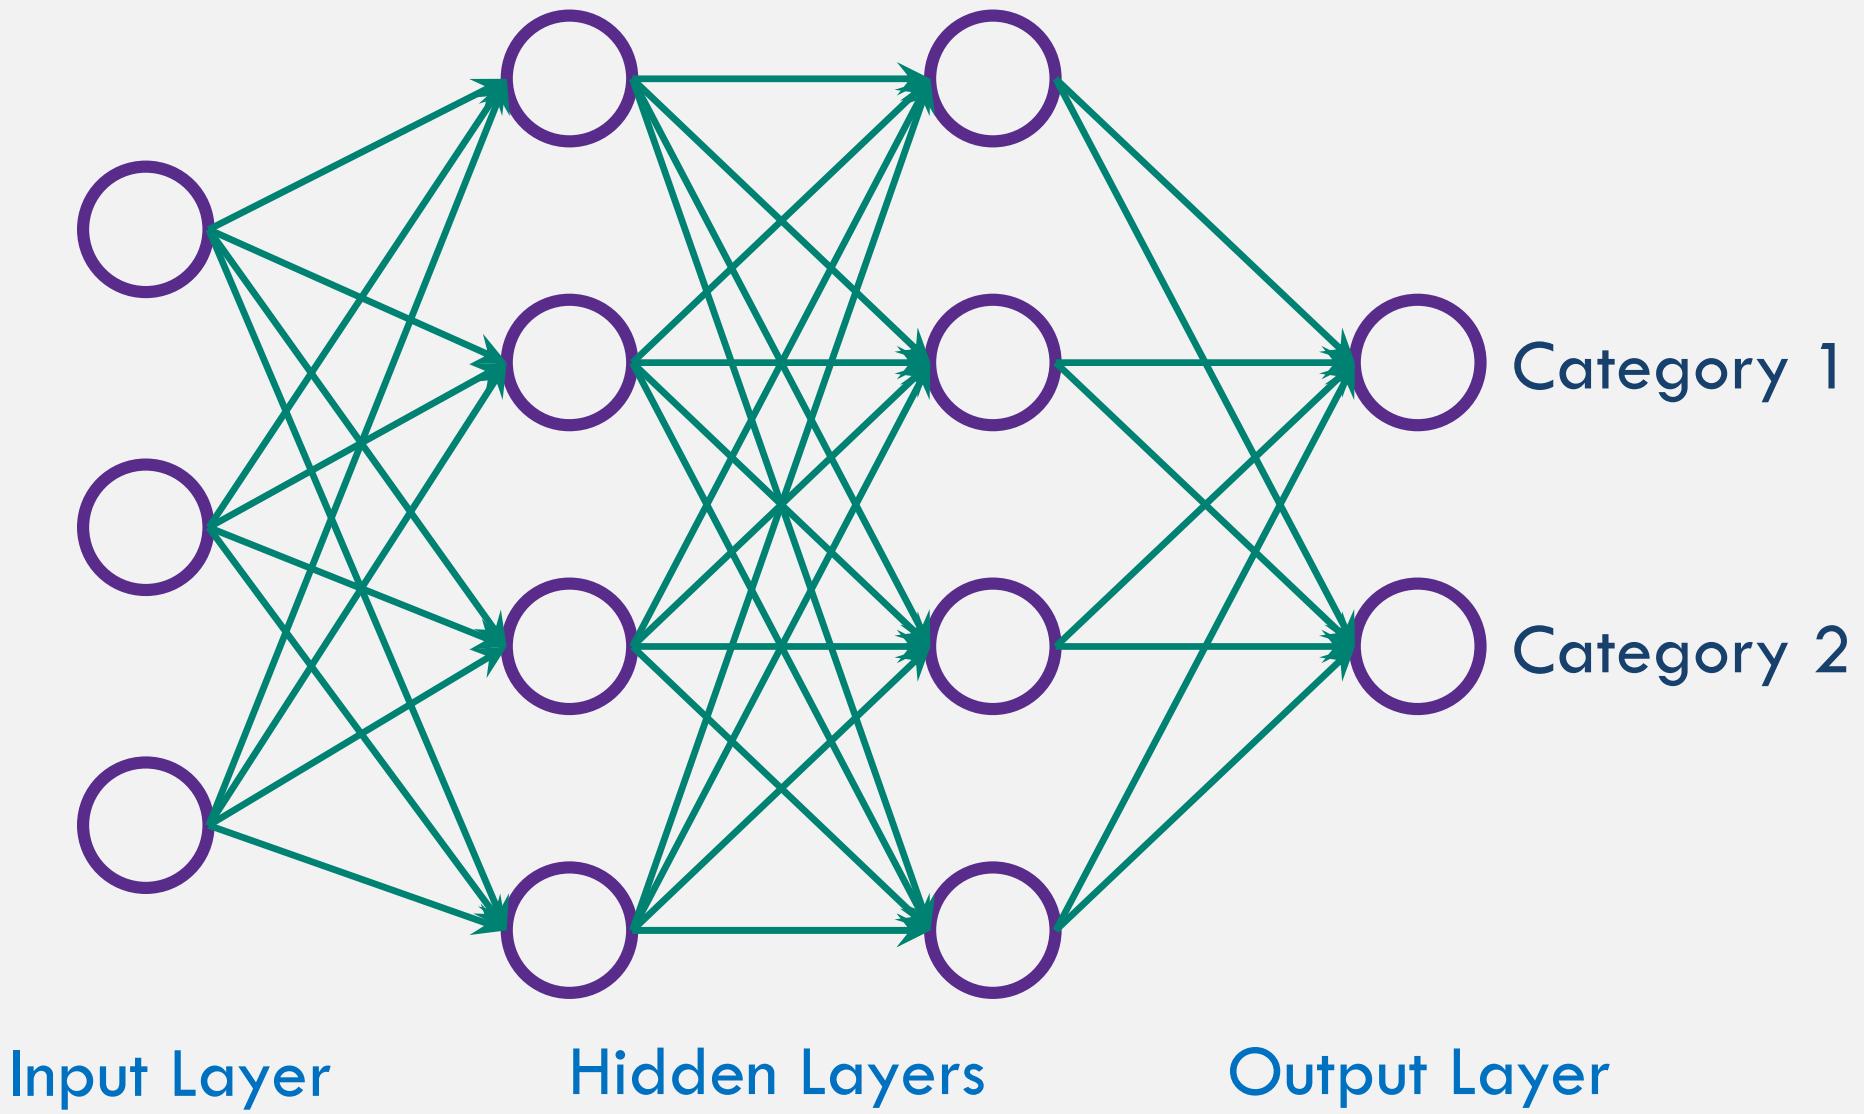
# GRADIENT DESCENT



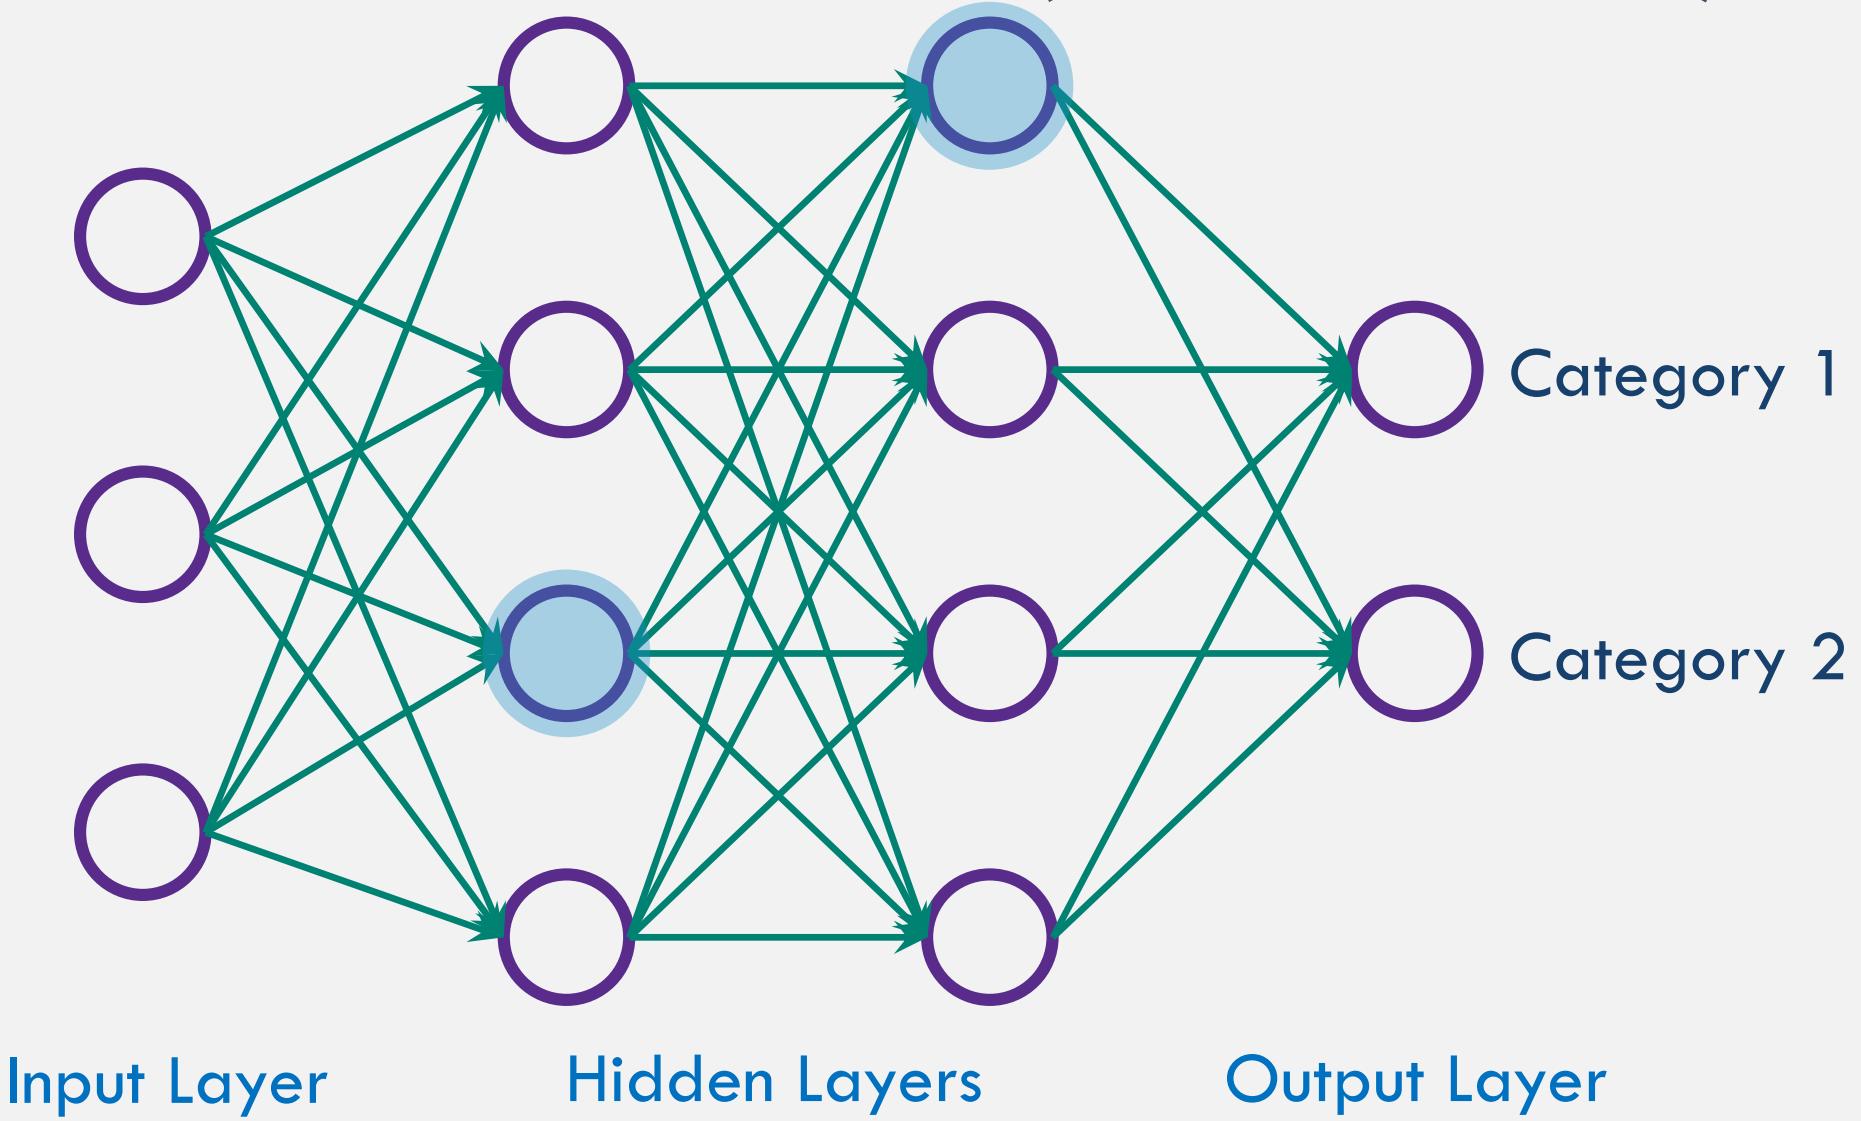
# LAB 2: Logistic Regression

# Neural Networks

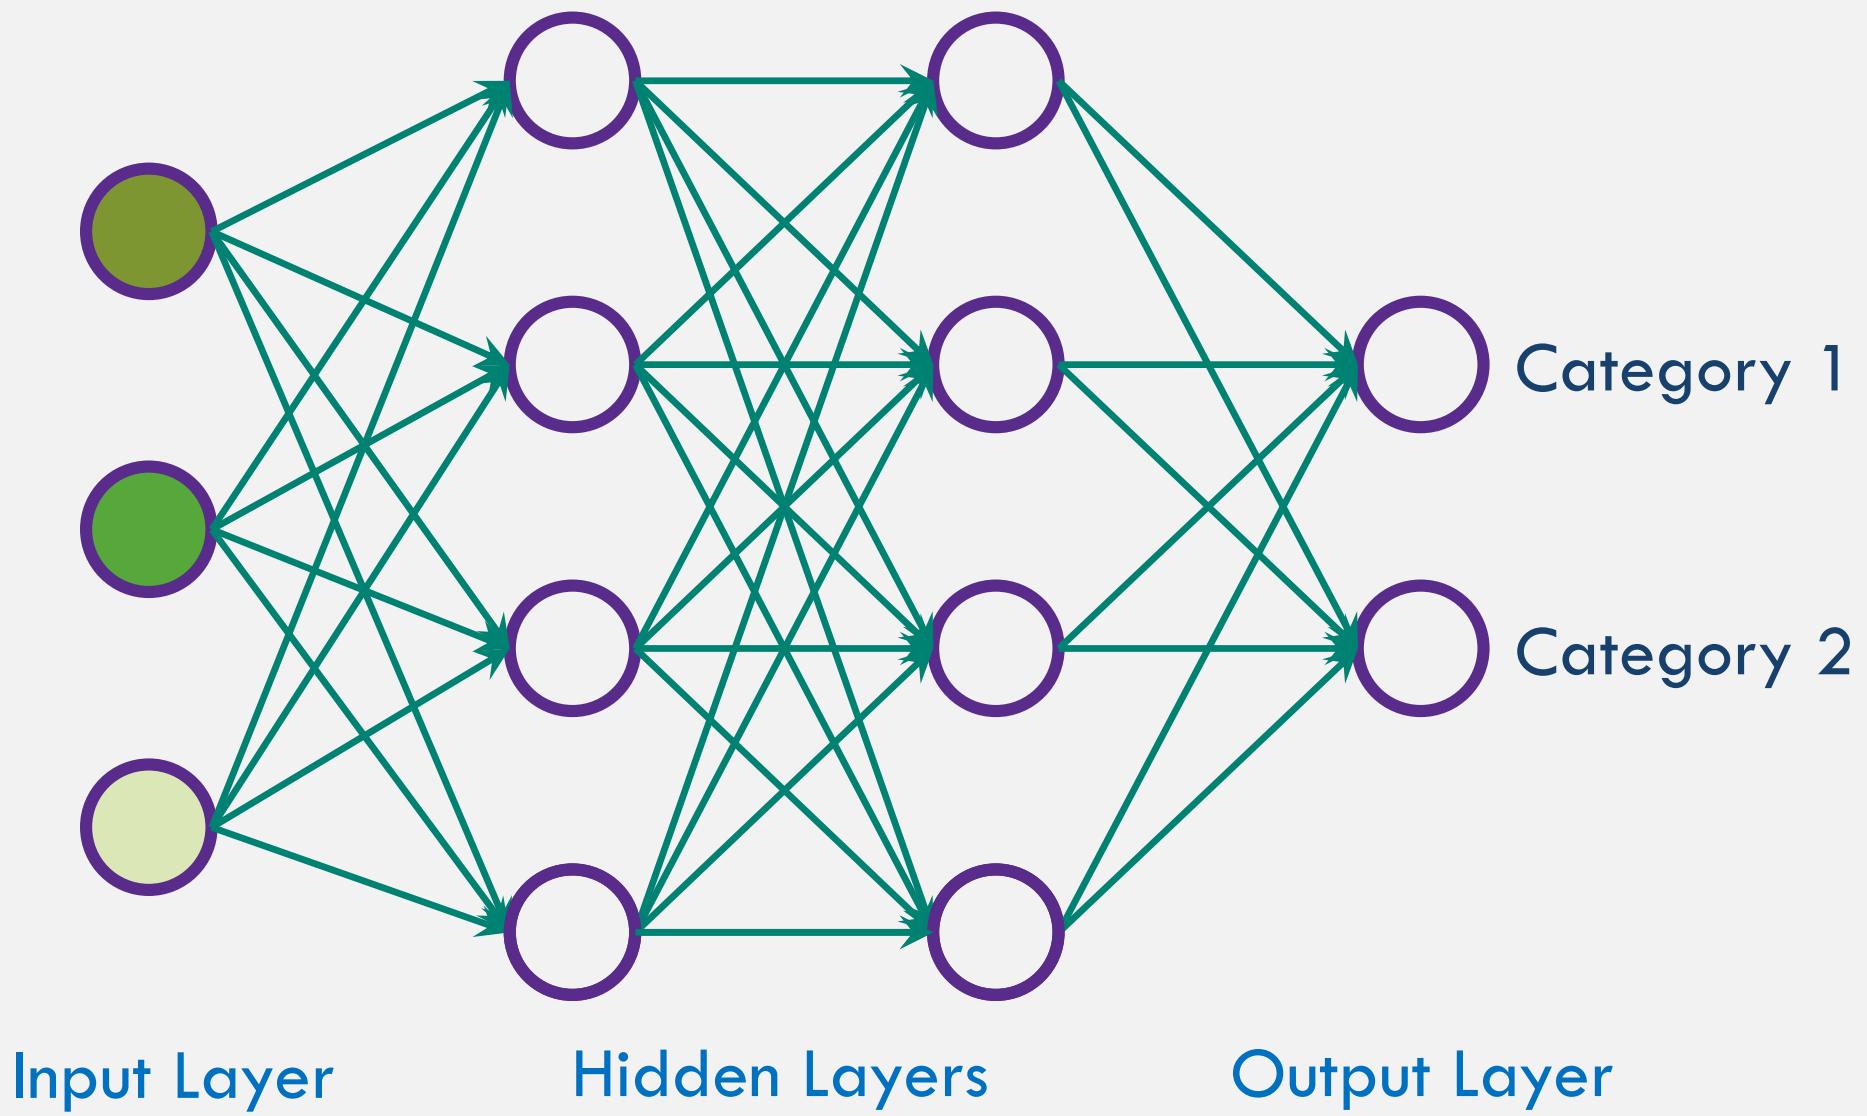
# NEURAL NETWORK (CLASSIFIER)



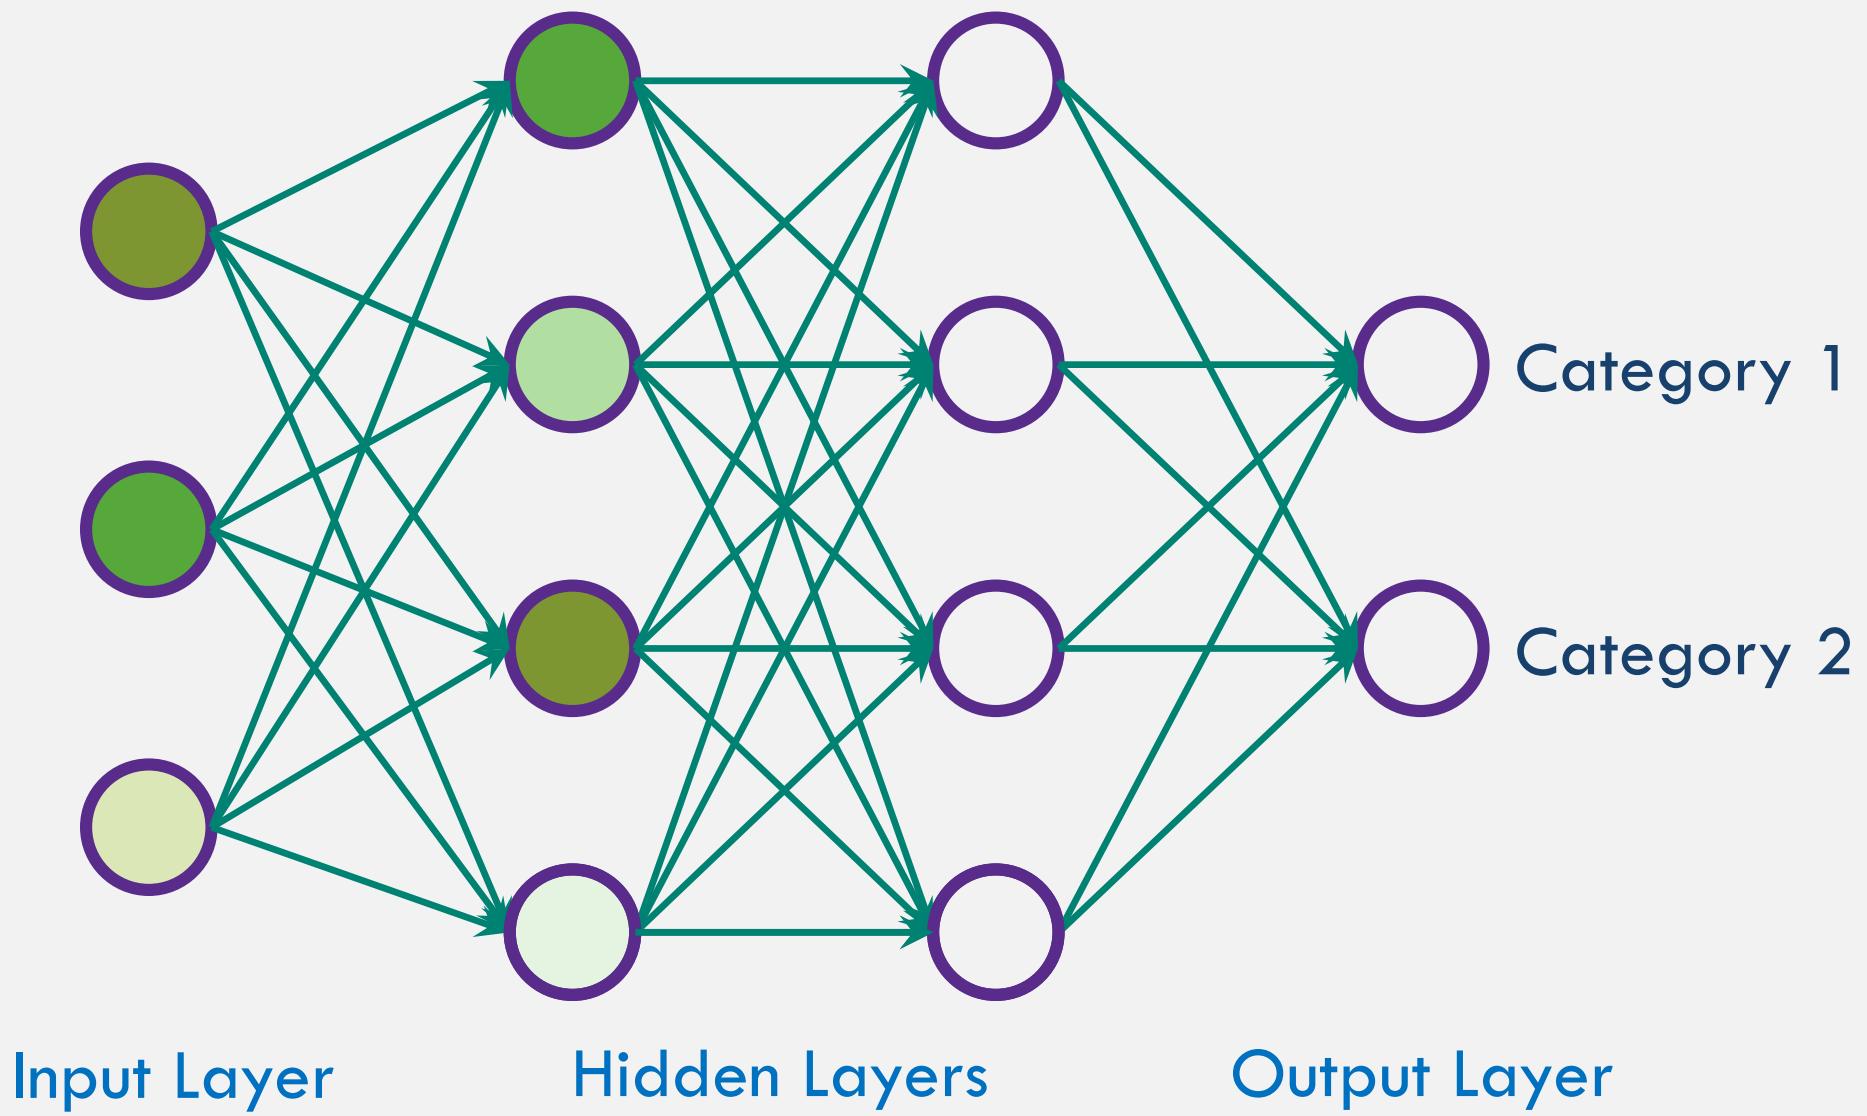
# NEURAL NETWORK (CLASSIFIER)



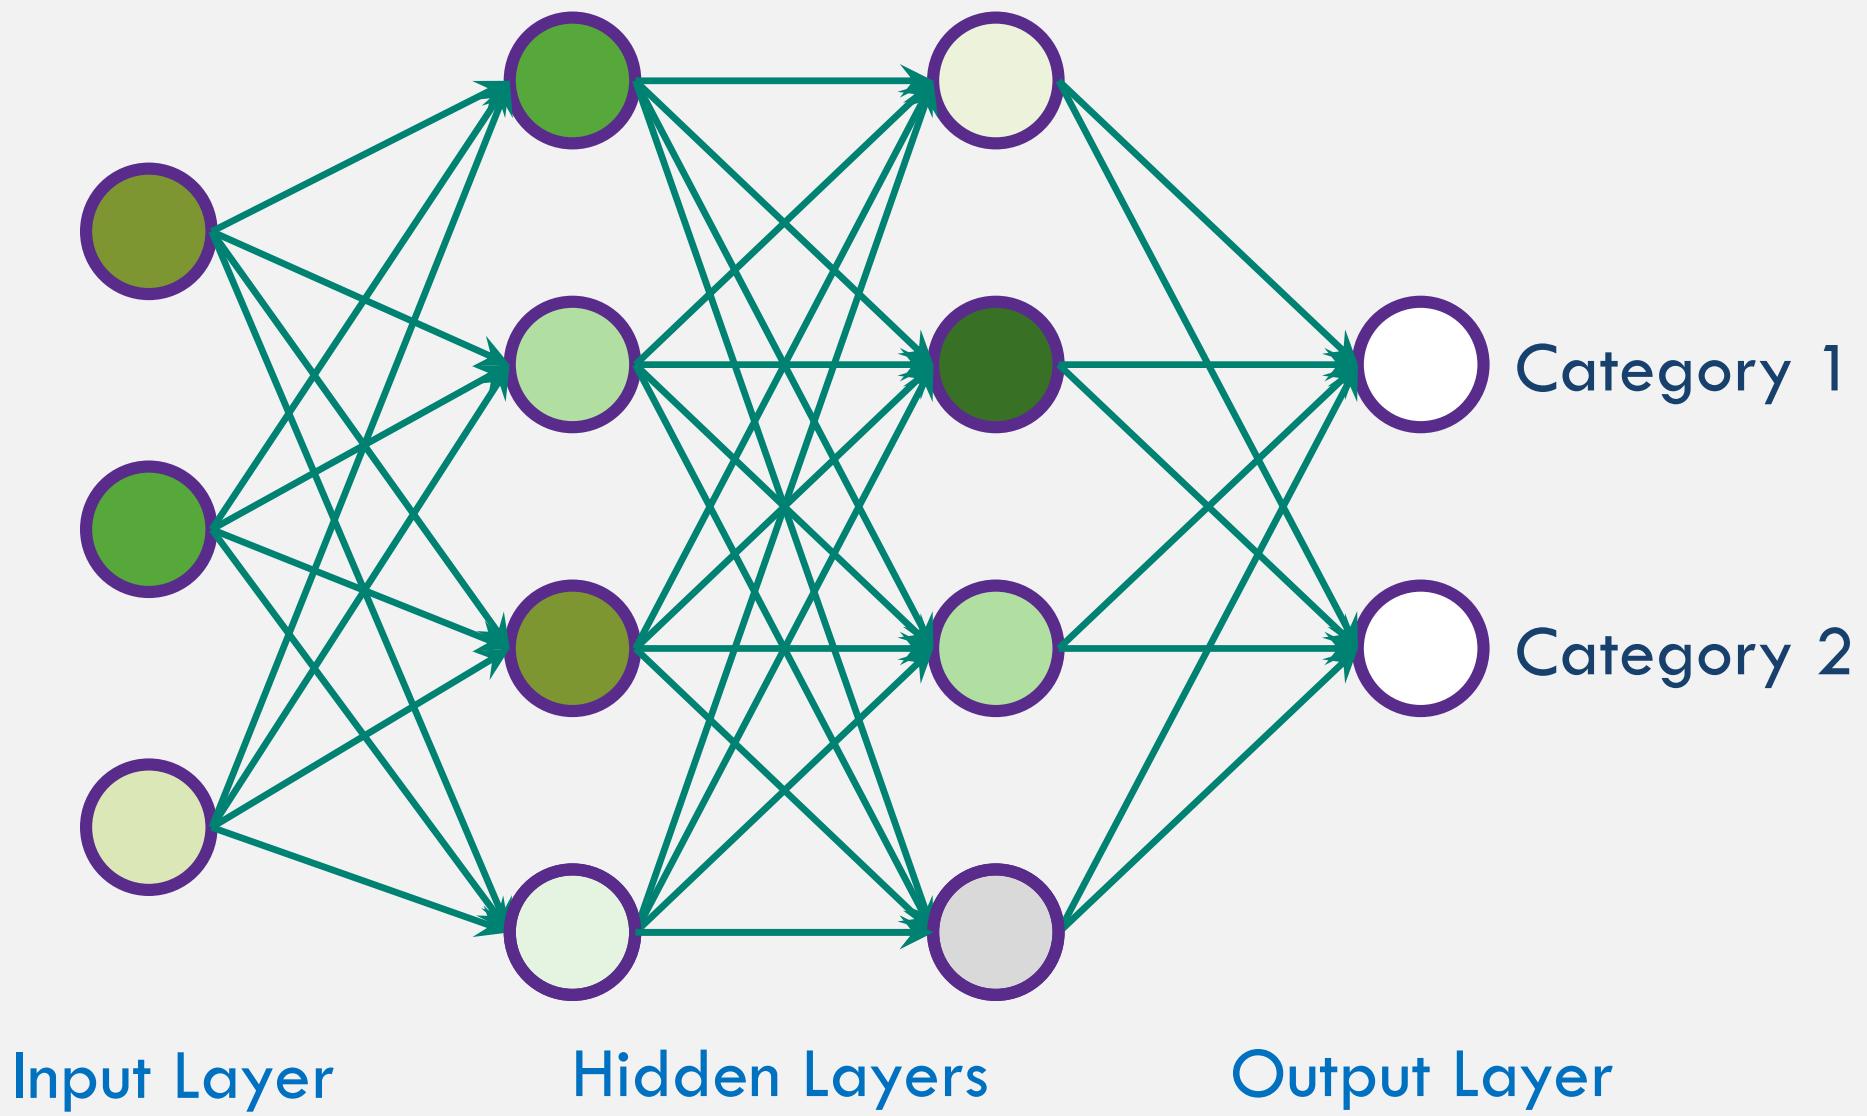
# FORWARD PROPAGATION



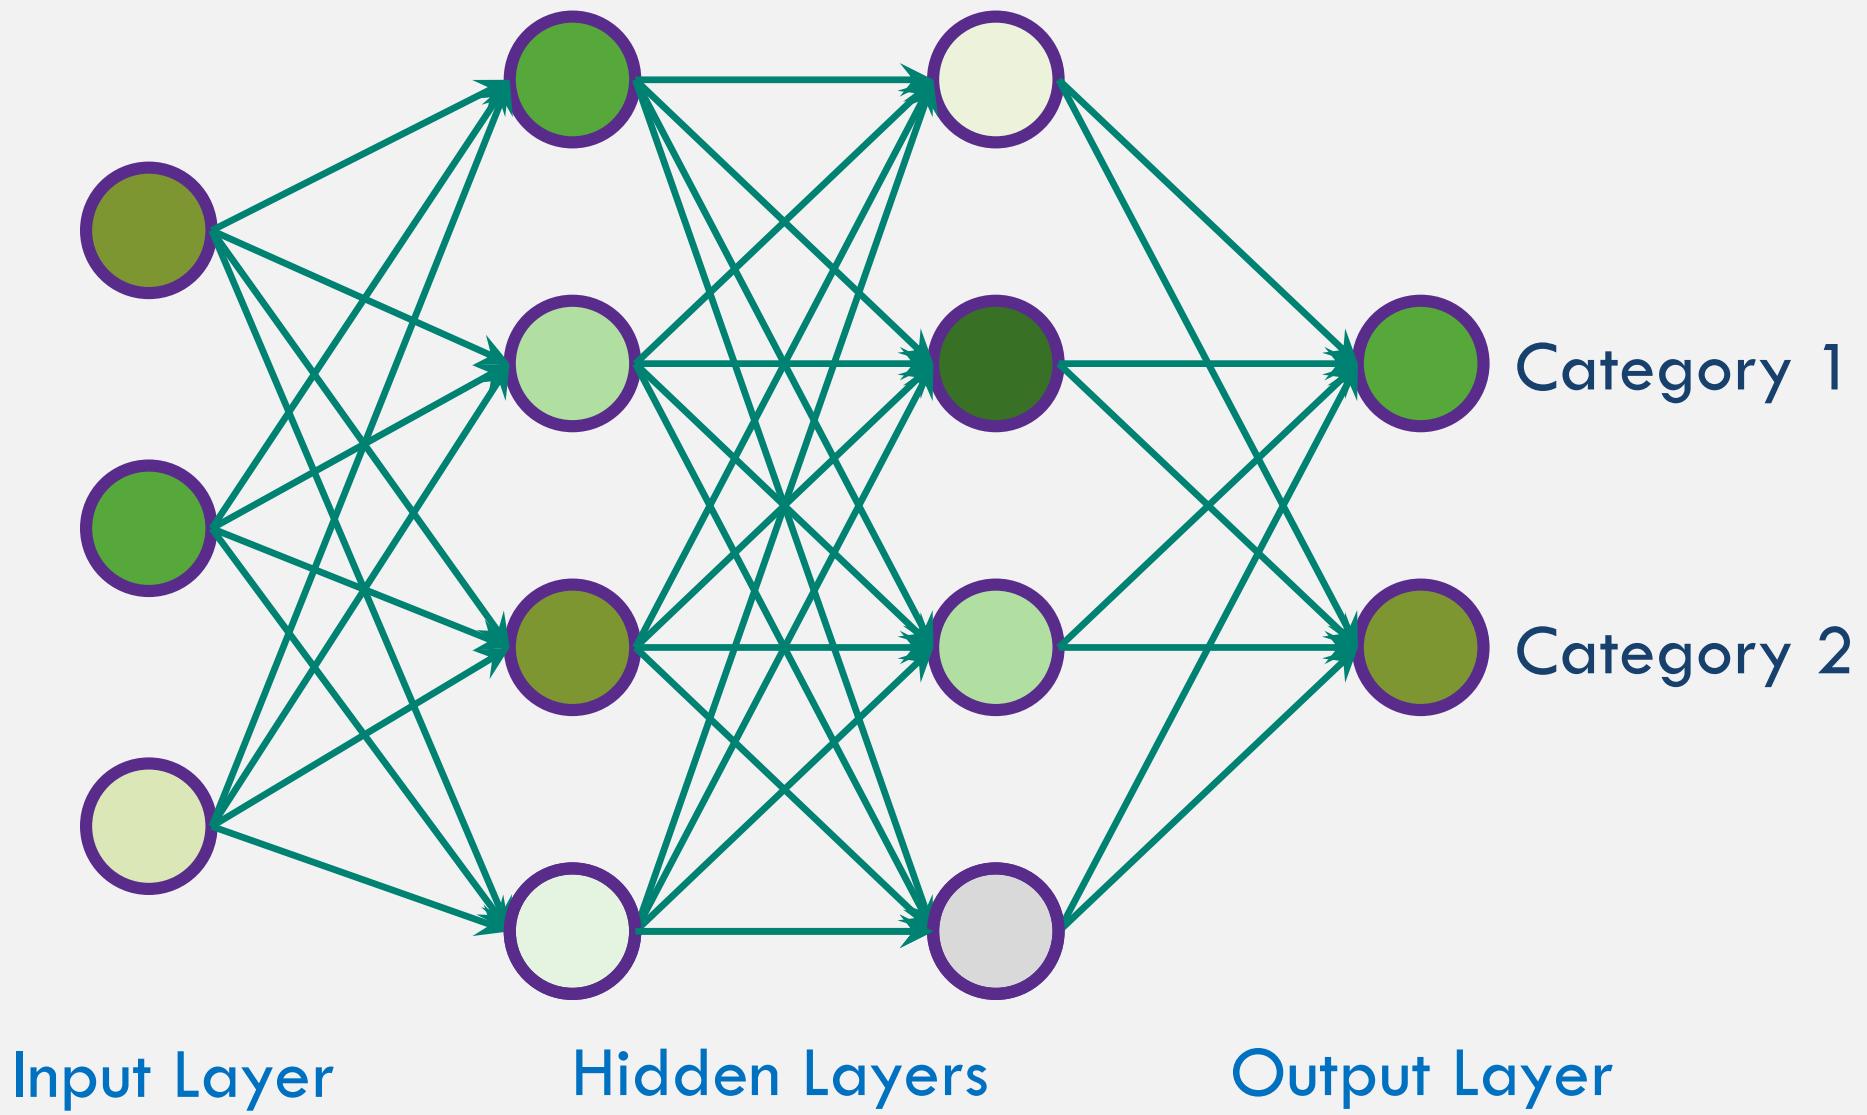
# FORWARD PROPAGATION



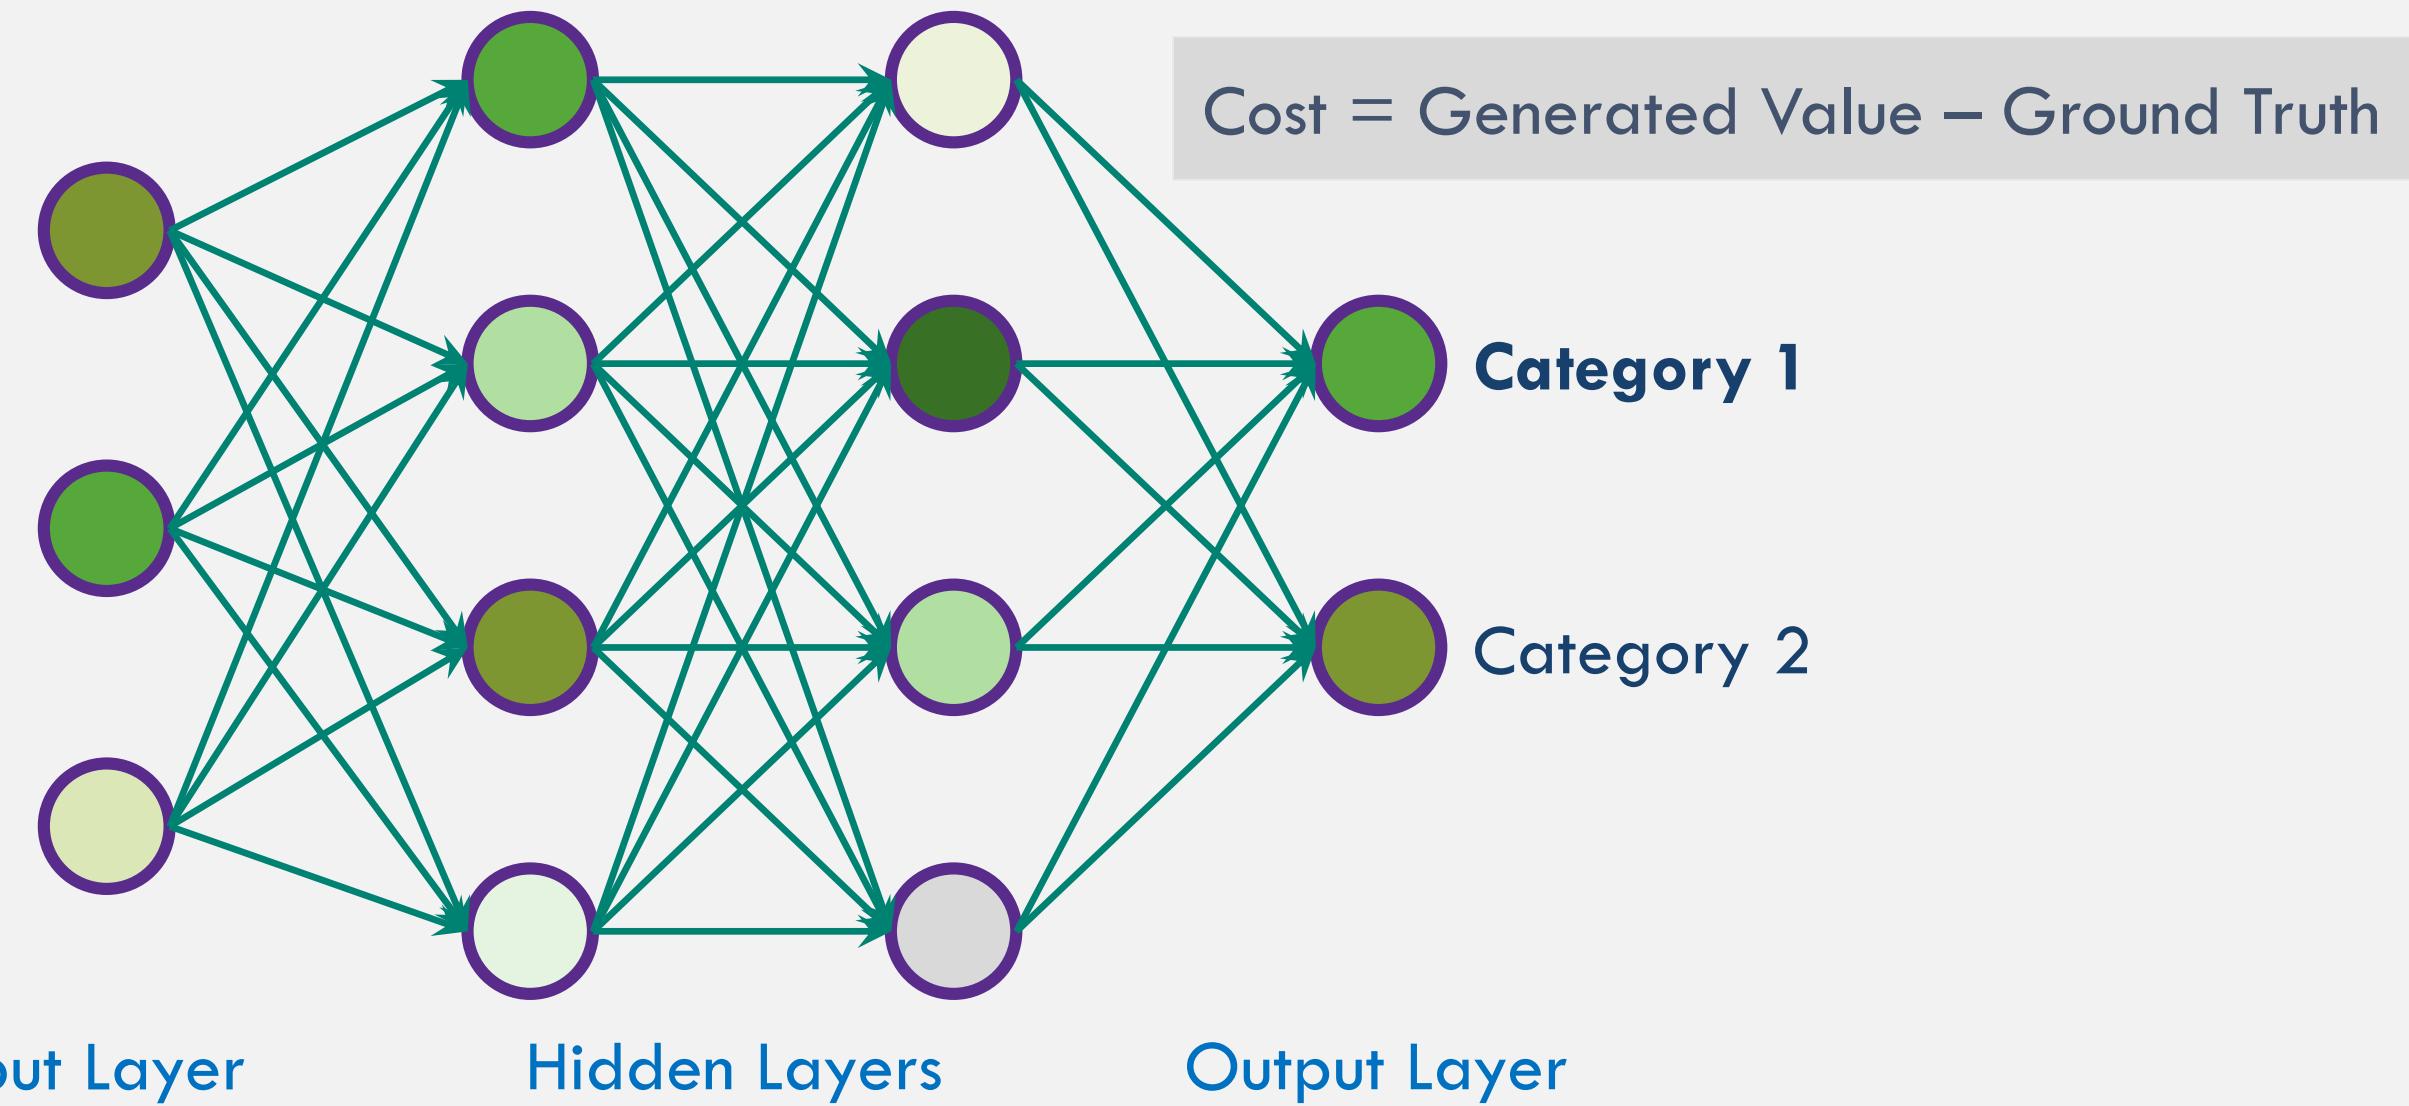
# FORWARD PROPAGATION



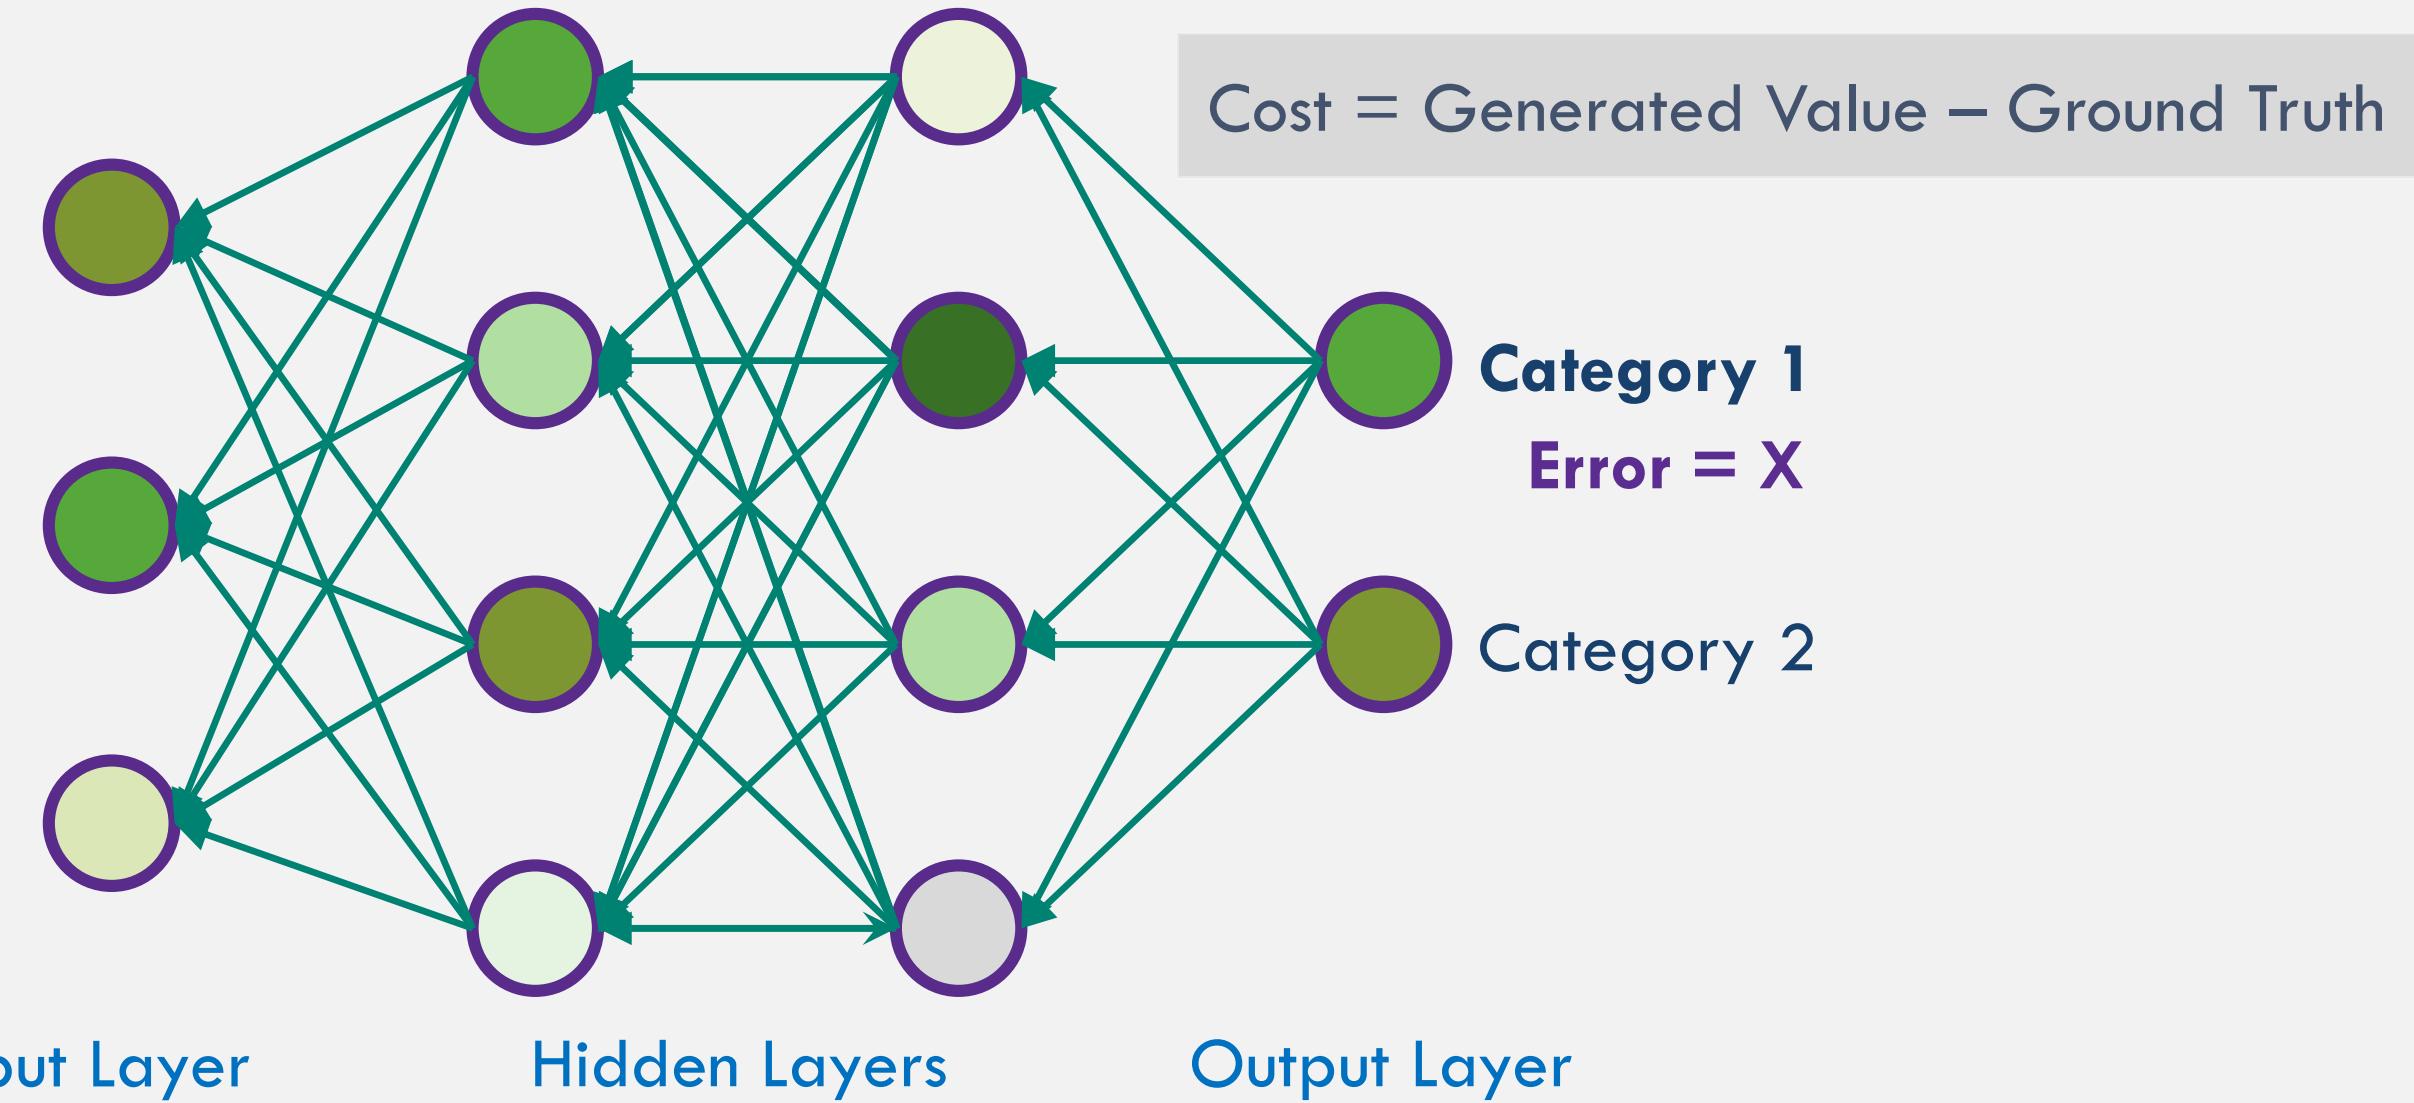
# FORWARD PROPAGATION



# TRAINING



# BACKPROPAGATION

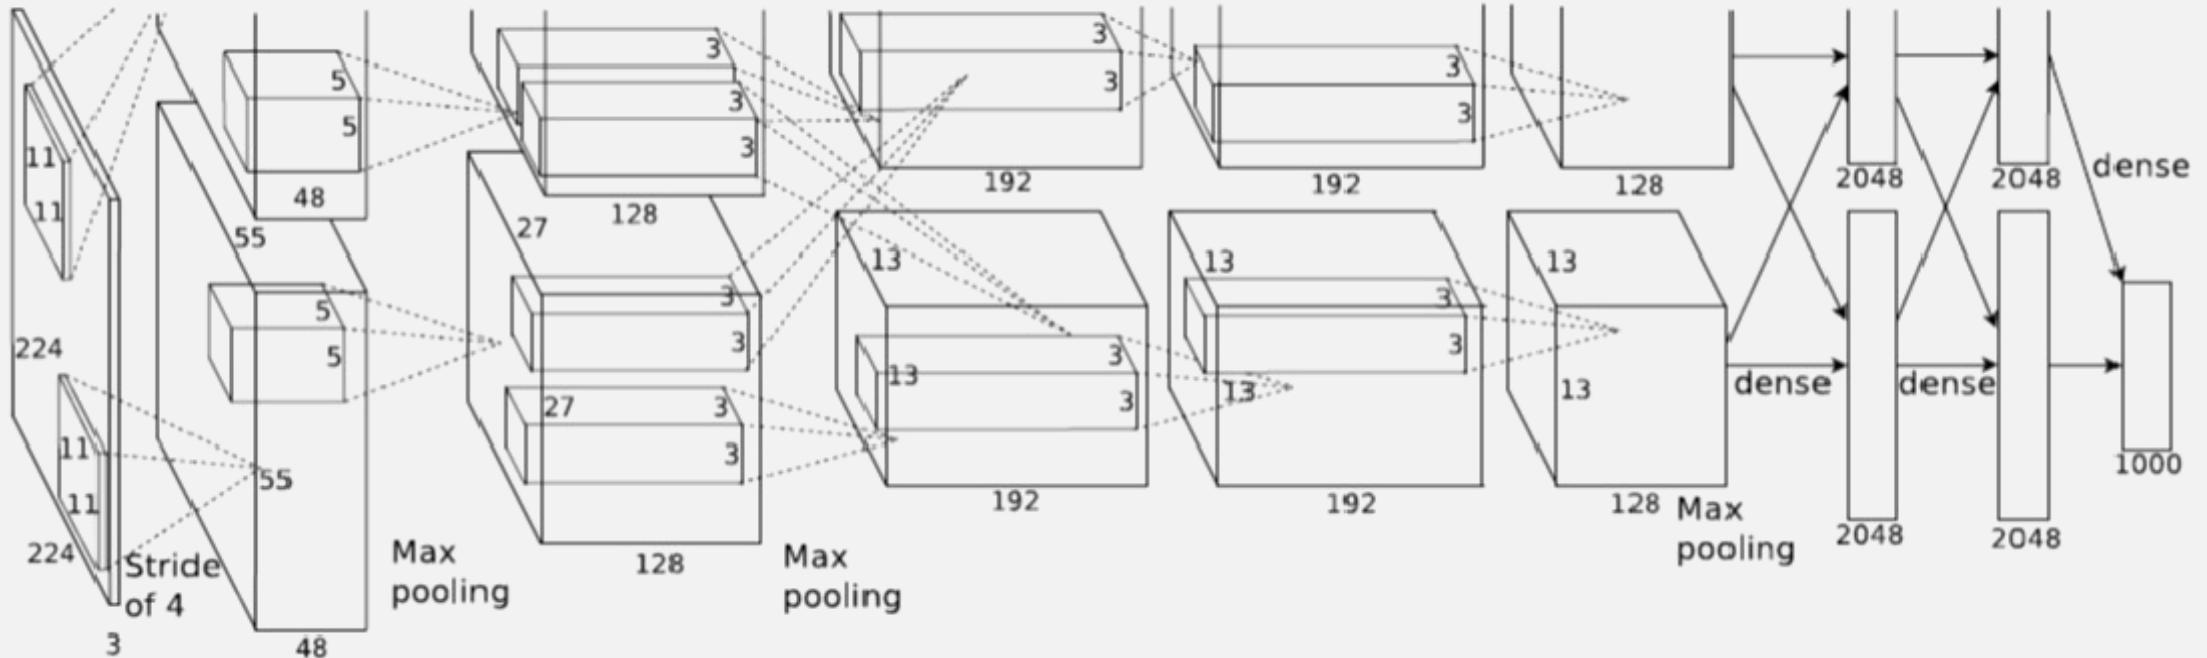


# LAB 3: Multi-Layer Perceptron

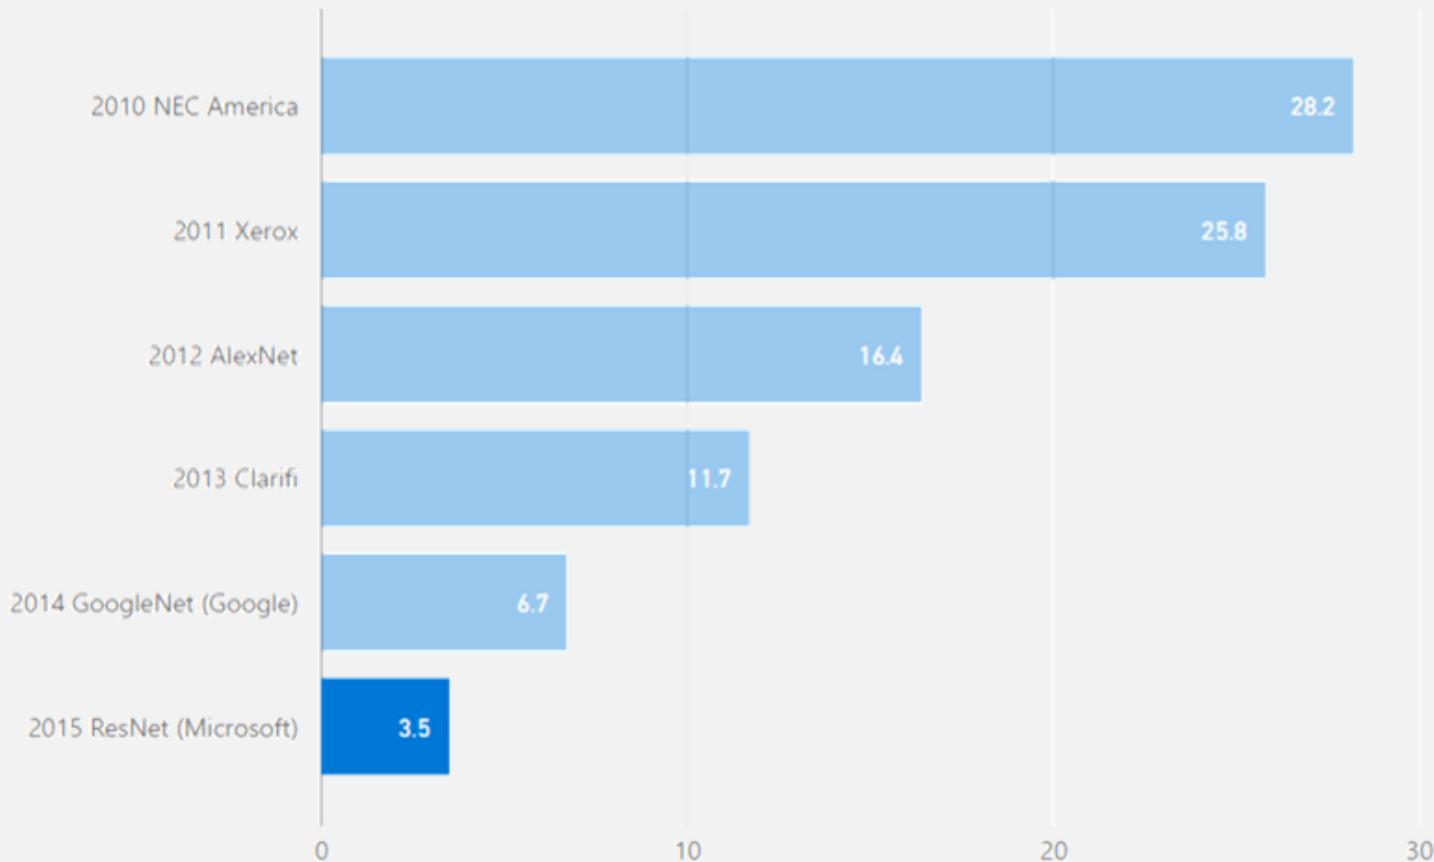
# Convolutional Networks

# REMEMBER THIS THING?

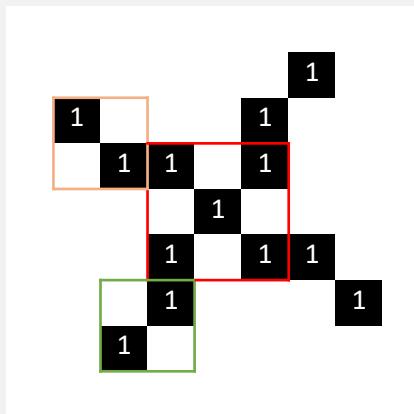
---



# IMAGENET CHALLENGE



# FILTERING



# FILTERING

...looking for common characteristics

Filter 1

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 2

-1	-1	1
-1	1	-1
1	-1	-1

Filter 3

1	-1	1
-1	1	-1
1	-1	1

1	1	1
1	1	1
1	1	1

1.0

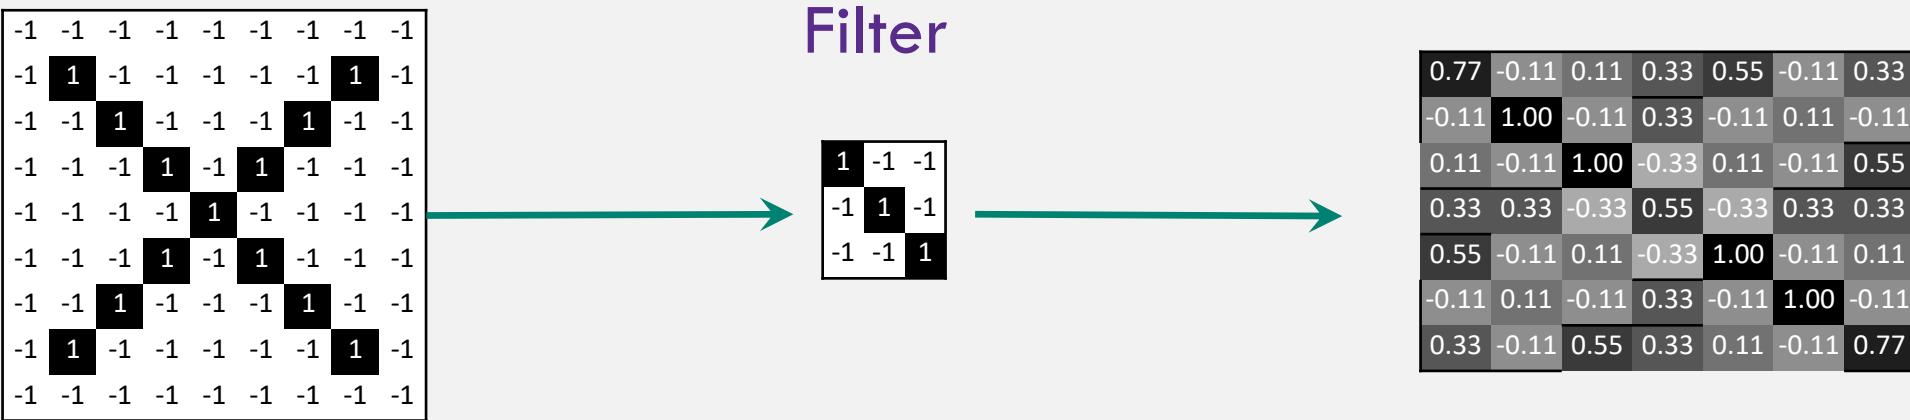
-1	1	-1
1	1	1
-1	1	-1

0.11

1	1	-1
1	1	1
-1	1	1

0.55

# CONVOLUTION



# CONVOLUTION

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# CONVOLUTION

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

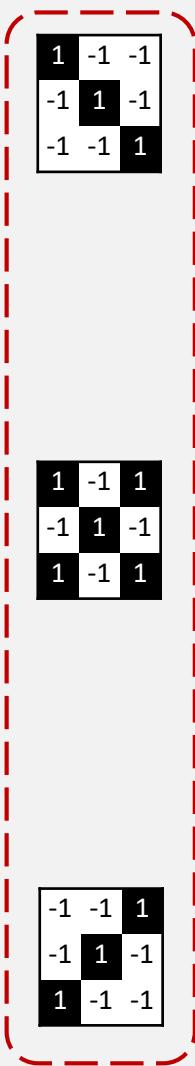


-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

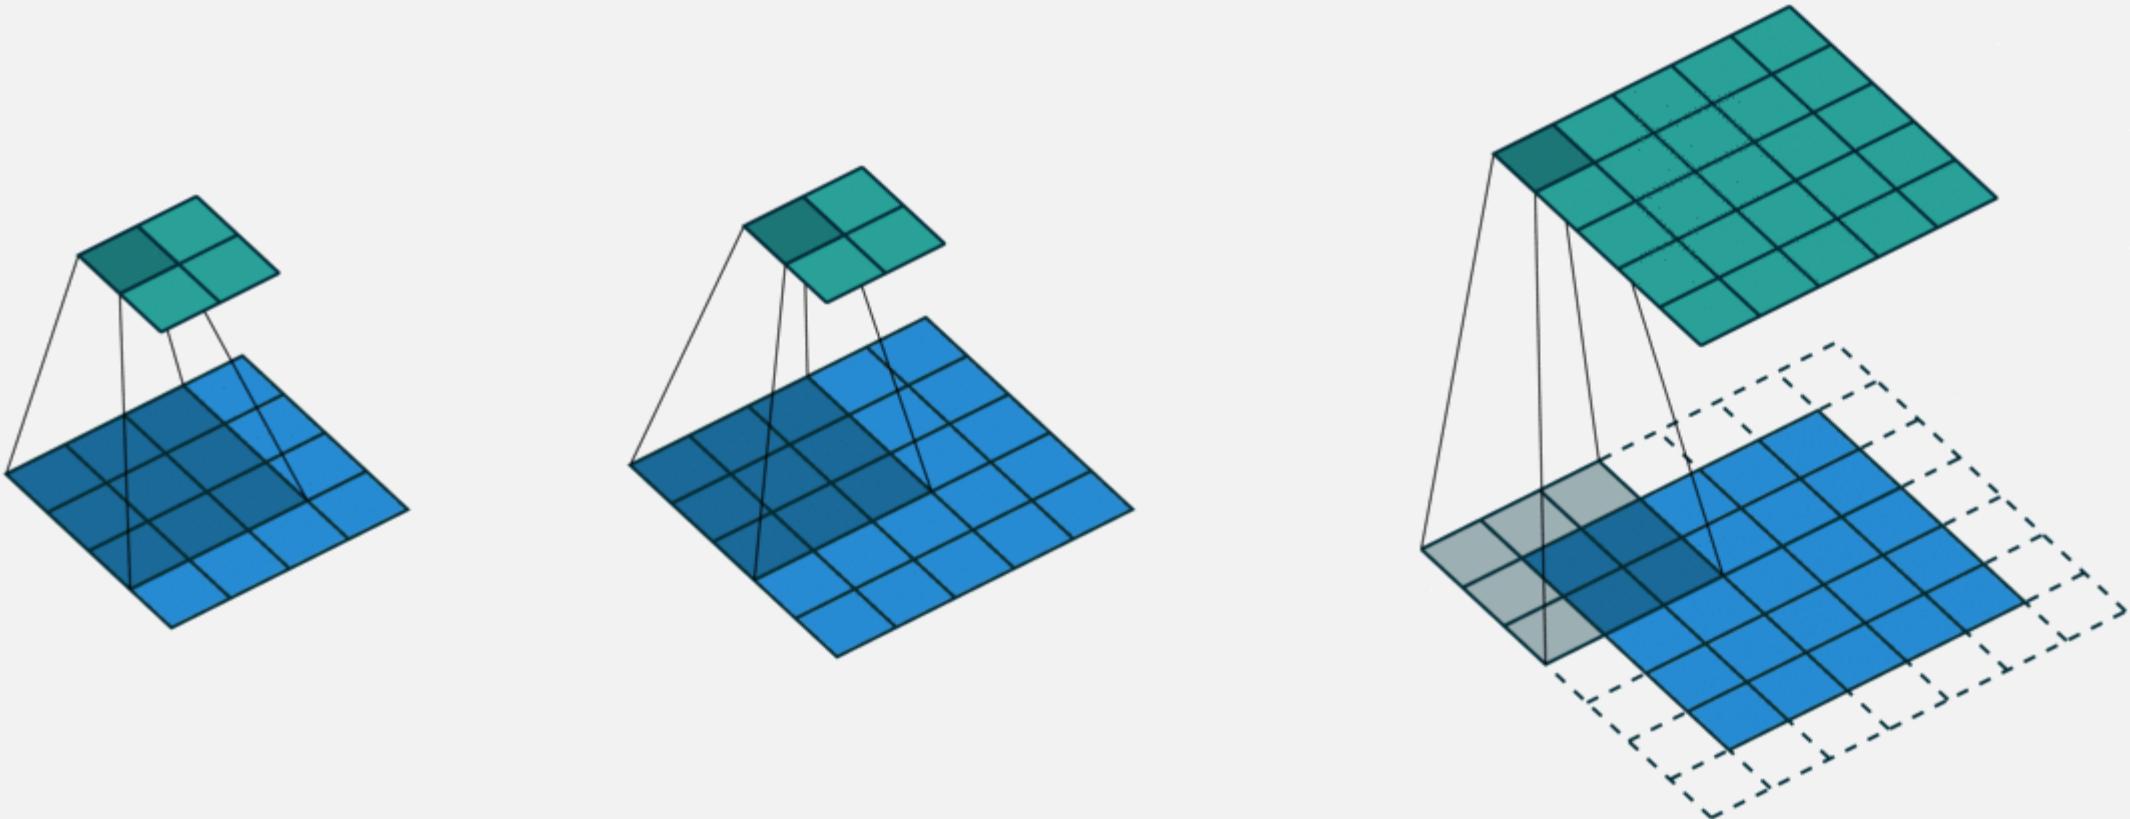


0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

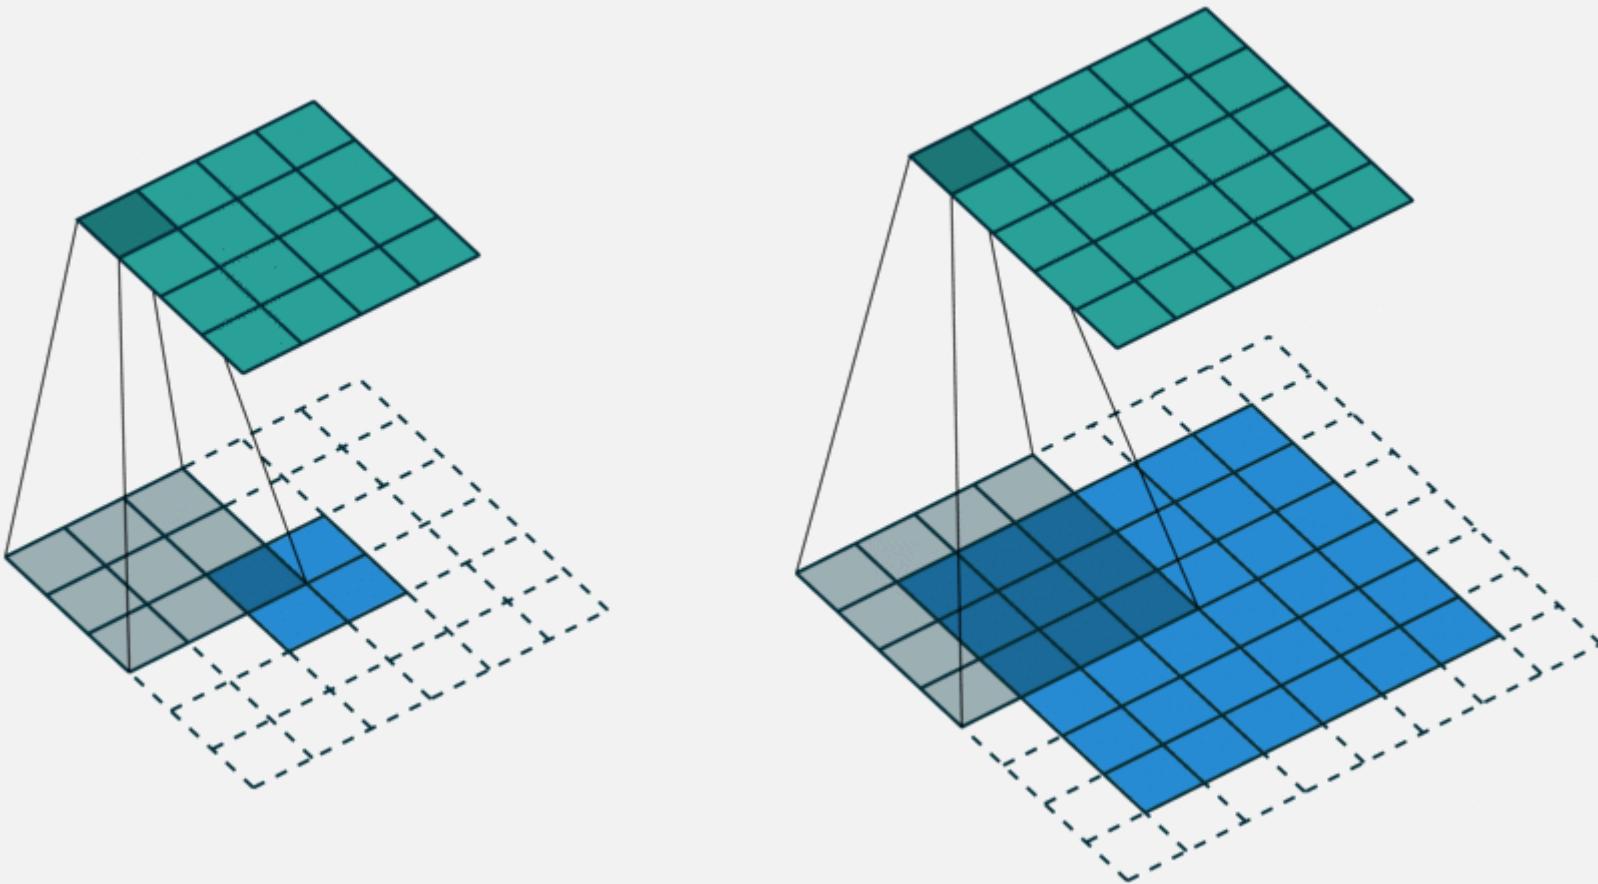
0.33	-0.55	-0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# STRIDES AND PADDING



# DECONVOLUTION – TRANSPOSED CONVOLUTION



# POOLING

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# POOLING

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# AVERAGE POOLING – STRIDE 1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.39	0.22	0.17	0.28	0.11	0.06
0.22	0.45	0.22	0.00	0.00	0.11
0.17	0.22	0.22	0.00	0.00	0.28
0.28	0.00	0.00	0.22	0.22	0.17
0.11	0.00	0.00	0.22	0.45	0.22
0.06	0.11	0.28	0.17	0.22	0.39

# MAX POOLING – STRIDE 1

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	1.00	0.33	0.55	0.55	0.33
1.00	1.00	1.00	0.33	0.11	0.55
0.33	1.00	1.00	0.55	0.33	0.55
0.55	0.33	0.55	1.00	1.00	0.33
0.55	0.11	0.33	1.00	1.00	1.00
0.33	0.55	0.55	0.33	1.00	1.00

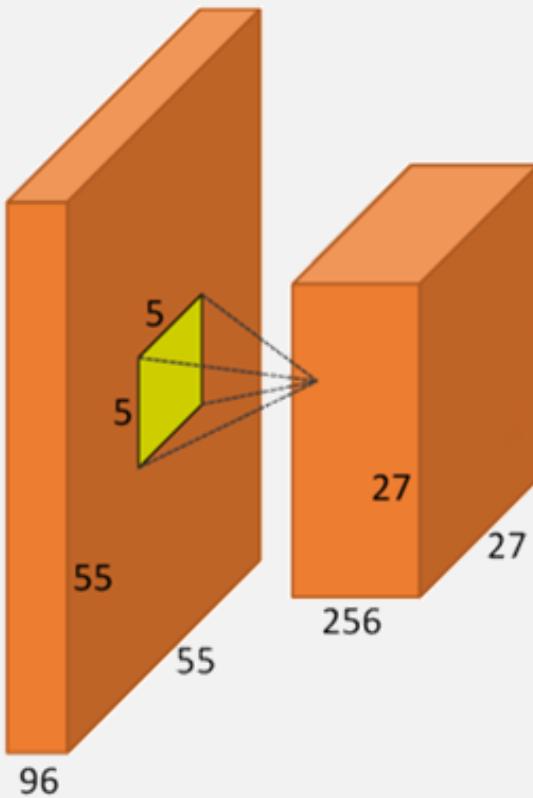
# AVERAGE POOLING – STRIDE 2

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

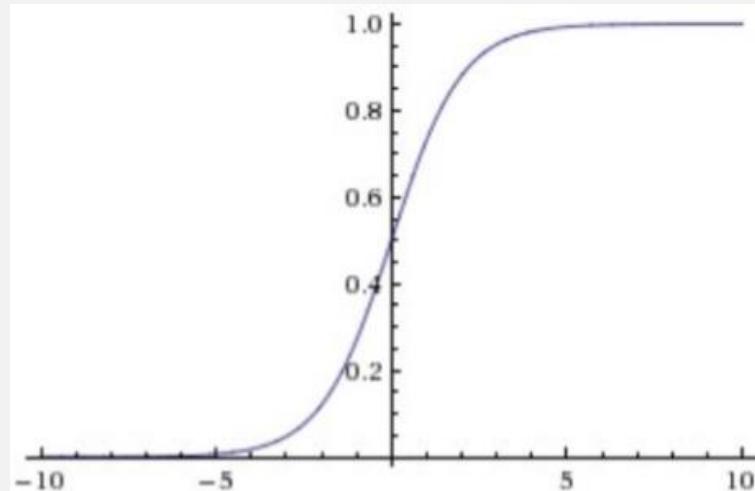
# WHY POOLING?

# ACTIVATION FUNCTIONS?



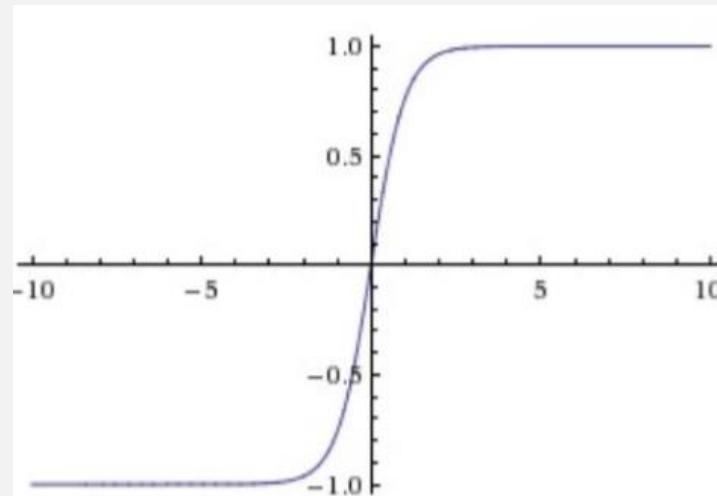
# ACTIVATION FUNCTIONS?

## LOGISTIC FUNCTION (“sigmoid”):



- Simple
- Drawbacks:
  - Saturates
  - Non-centered in zero
- We don't use it anymore, but we thank her for the services provided to the cause :)

# ACTIVATION FUNCTIONS?

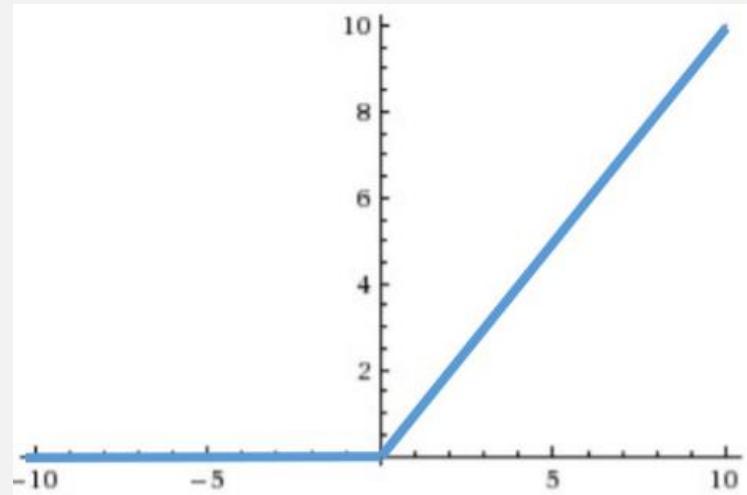


## HYPERBOLIC TANGENT ( $\tanh$ ):

- Solves the issue with not being centered in zero
- Still ‘killing gradients’ due to saturating behaviour
- It is always better than the logistic function

# ACTIVATION FUNCTIONS?

## RECTIFIED LINEAR UNITS (ReLU):

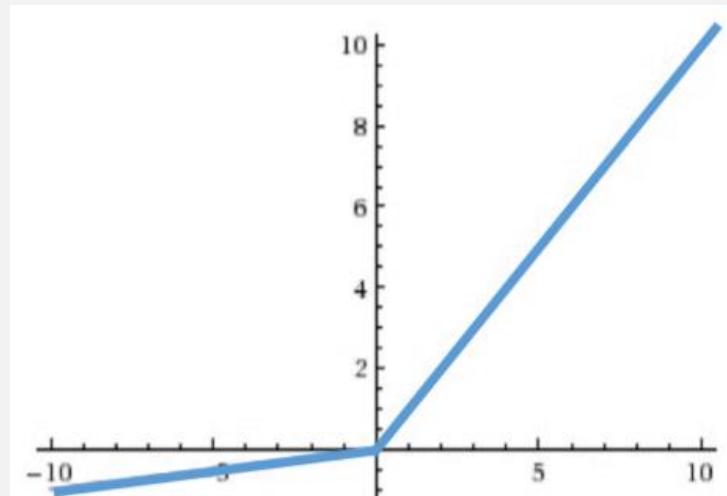


- Optimization for the SGD convergence
- More simple from a computational point of view
- Fragile

# ACTIVATION FUNCTIONS?

## LEAKY ReLU:

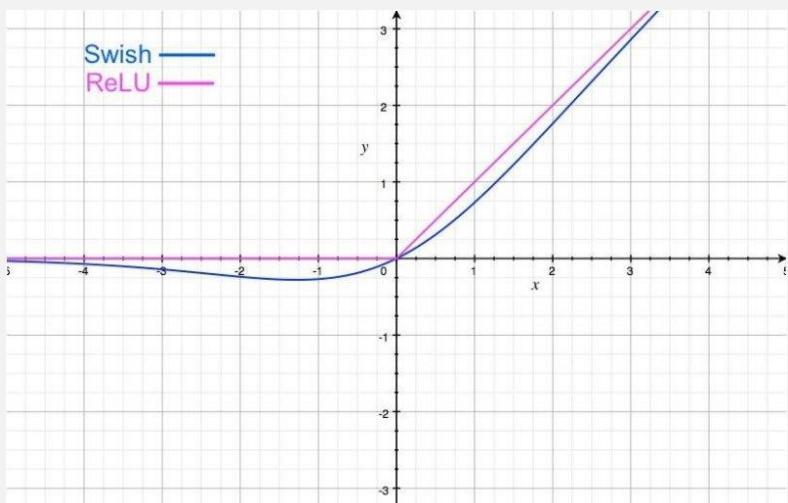
- More stable than traditional ReLU
- Particular case of MaxOut



# ACTIVATION FUNCTIONS?

## SWISH:

- Self-Gating activation function
- Better performance than ReLU



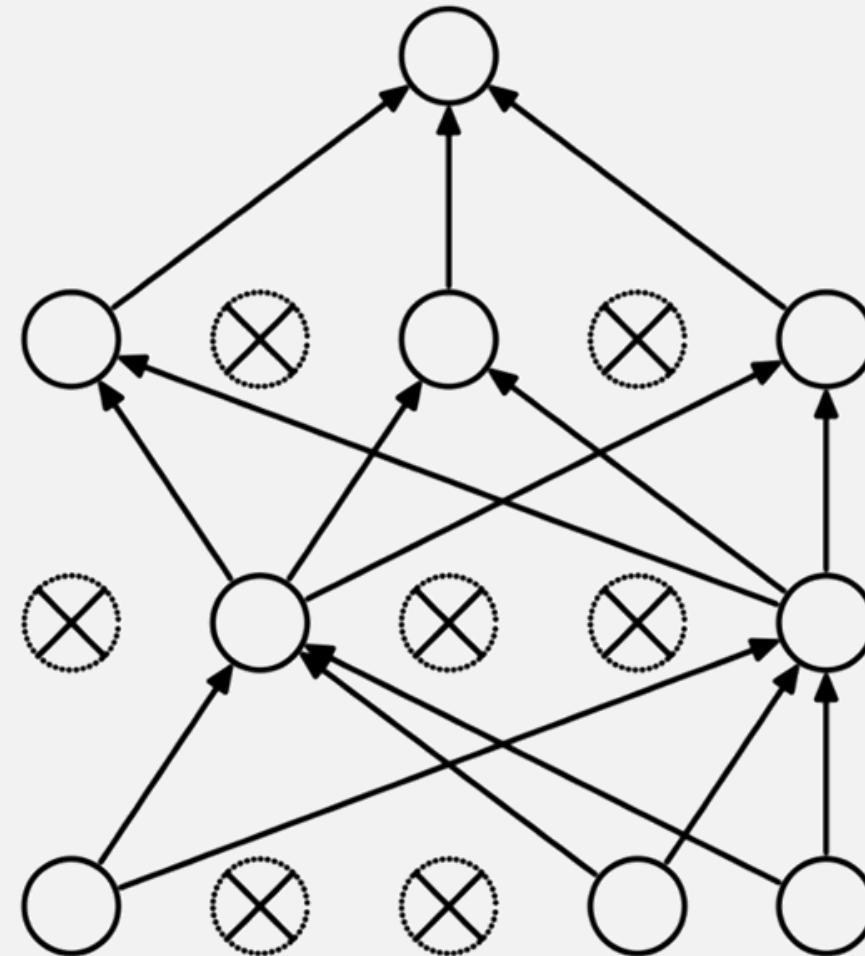
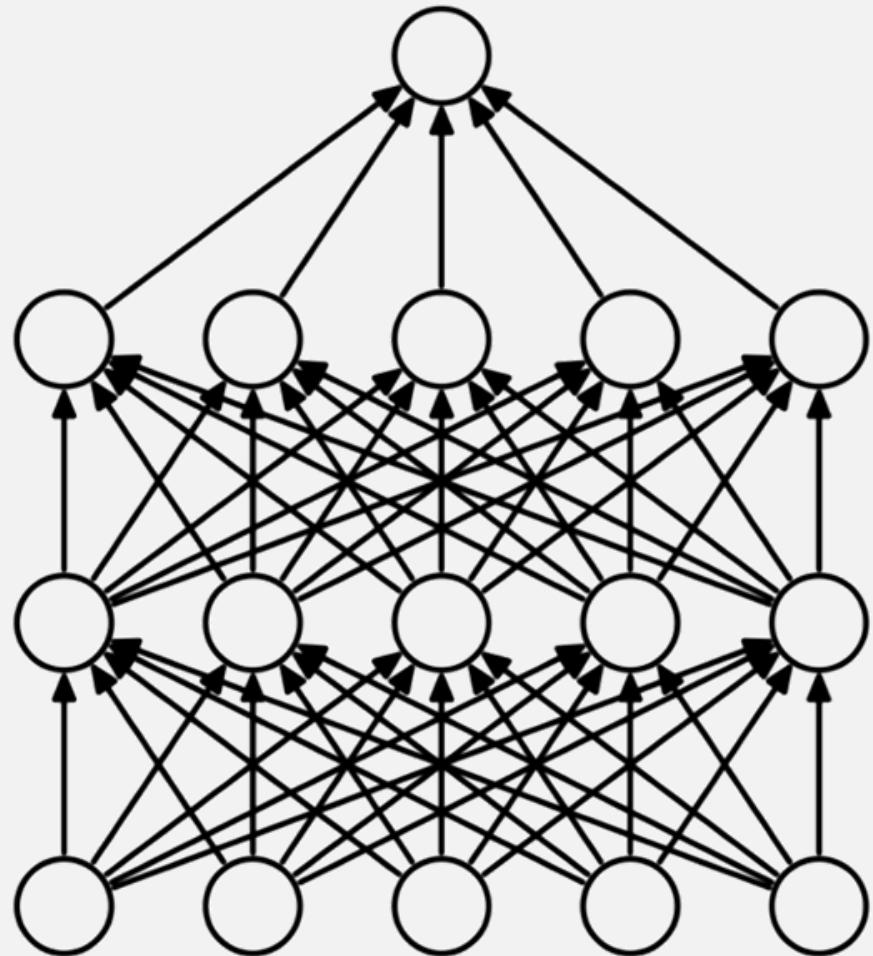
# RECTIFIED LINEAR UNITS (ReLU)

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

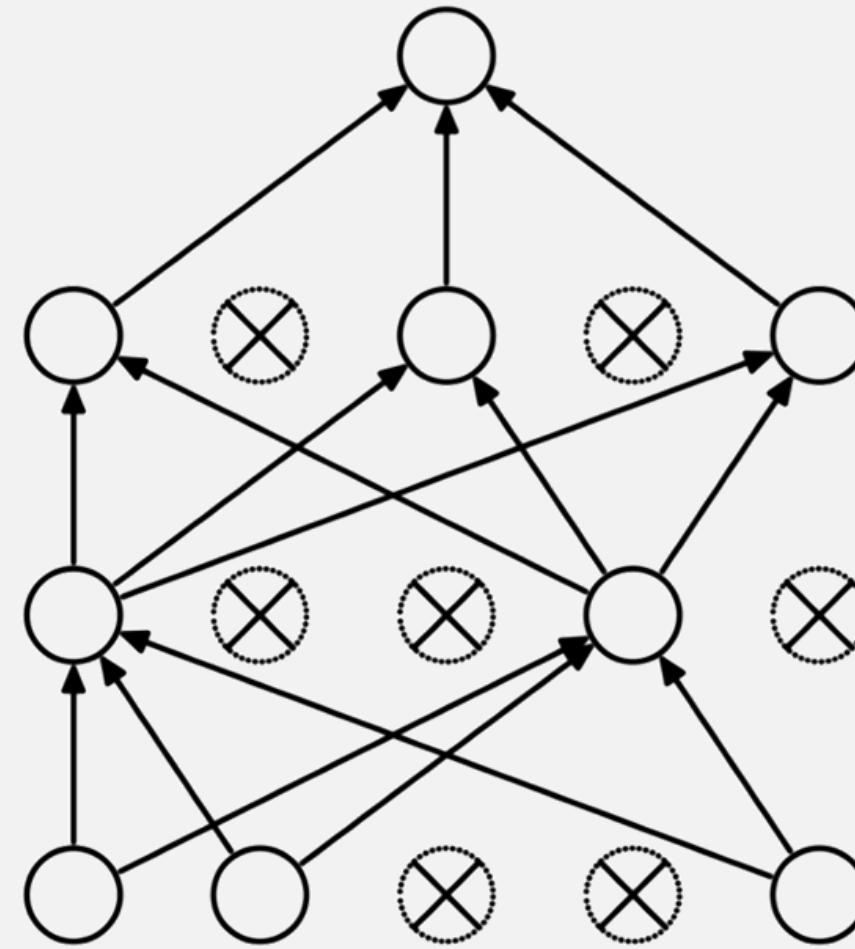
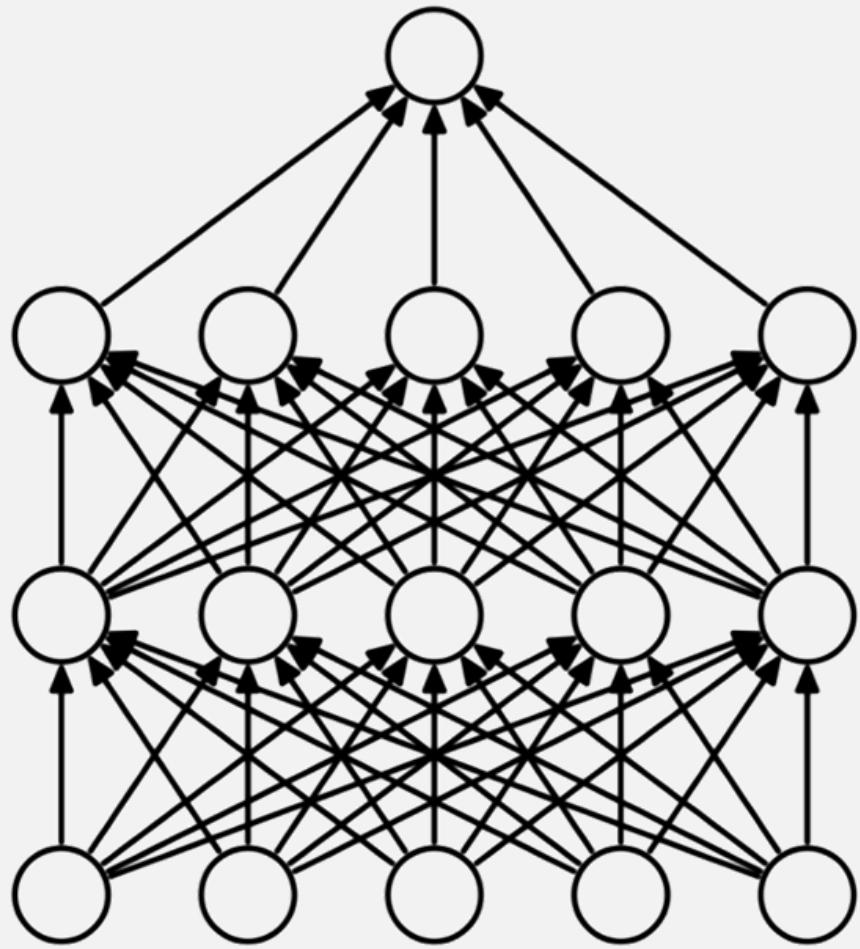


0.77	0.00	0.11	0.33	0.55	0.00	0.33
0.00	1.00	0.00	0.33	0.00	0.11	0.00
0.11	0.00	1.00	0.00	0.11	0.00	0.55
0.33	0.33	0.00	0.55	0.00	0.33	0.33
0.55	0.00	0.11	0.00	1.00	0.00	0.11
0.00	0.11	0.00	0.33	0.00	1.00	0.00
0.33	0.00	0.55	0.33	0.11	0.00	0.77

# DROPOUT



# DROPOUT



# WHAT IS MISSING?

Filter 1      Filter 2      Filter 3

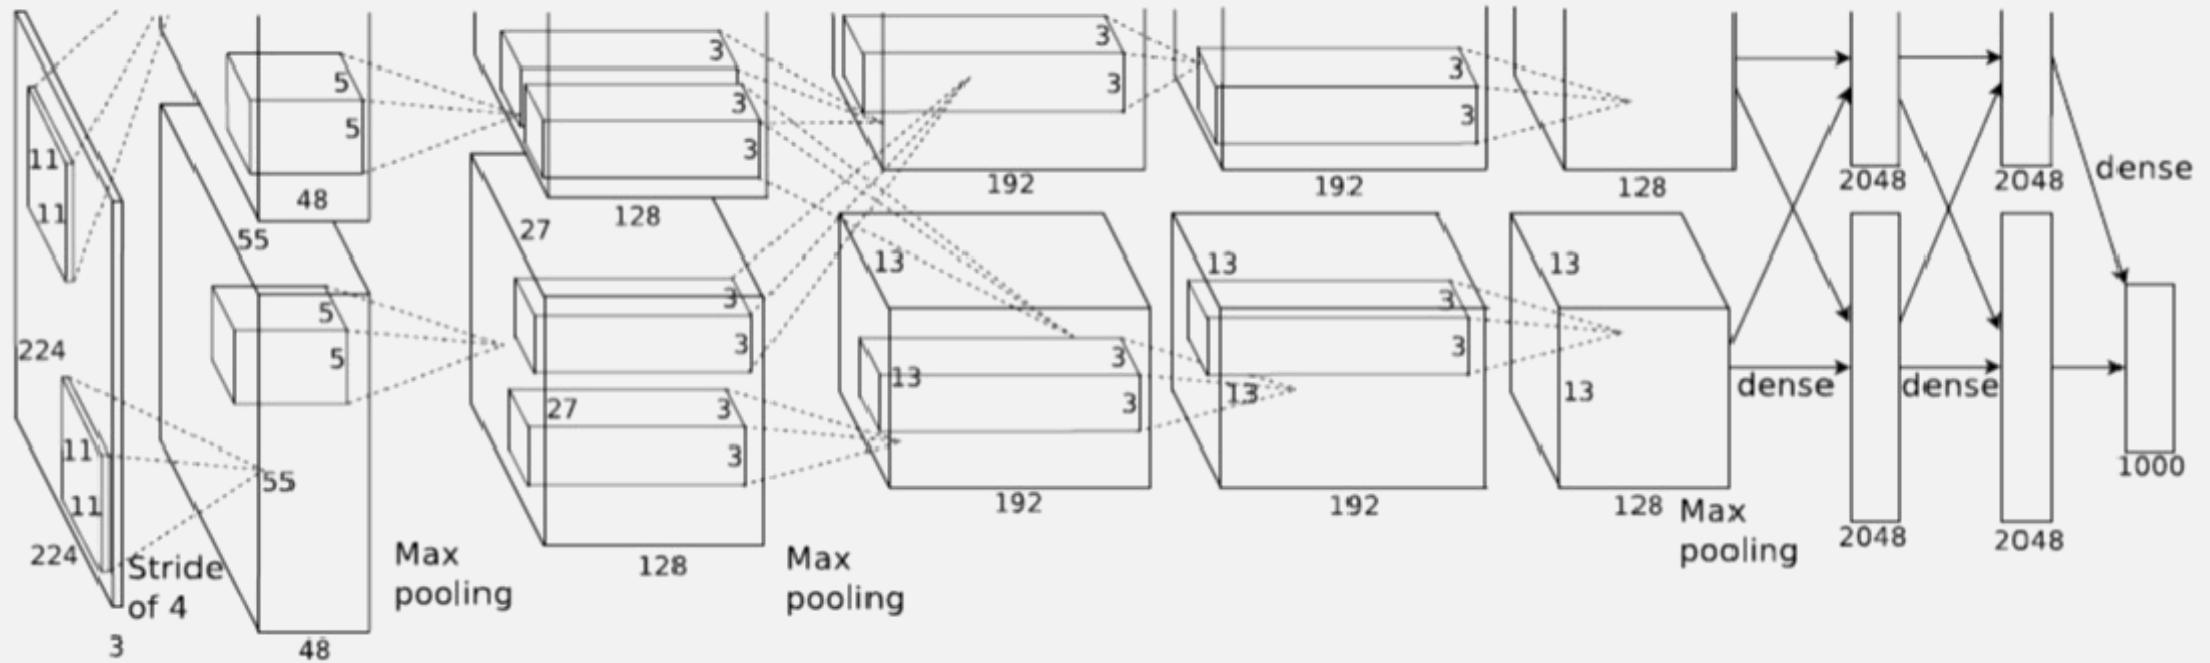
$$\begin{matrix} 1 & -1 & -1 \\ -1 & \boxed{1} & -1 \\ -1 & -1 & 1 \end{matrix}$$

$$\begin{matrix} -1 & -1 & 1 \\ -1 & \boxed{1} & -1 \\ 1 & -1 & -1 \end{matrix}$$

$$\begin{matrix} 1 & -1 & 1 \\ -1 & \boxed{1} & -1 \\ 1 & -1 & 1 \end{matrix}$$

How have these filters been  
selected/created?

# AlexNet



# HOW A CNN SEES

---

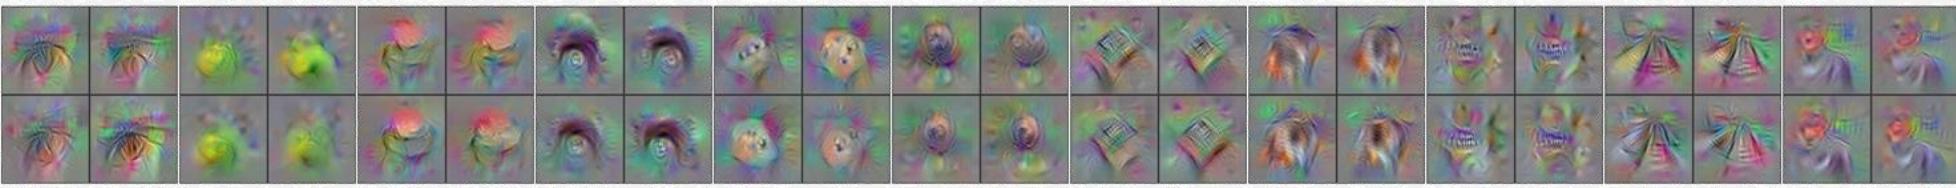
Layer 2



# HOW A CNN SEES

---

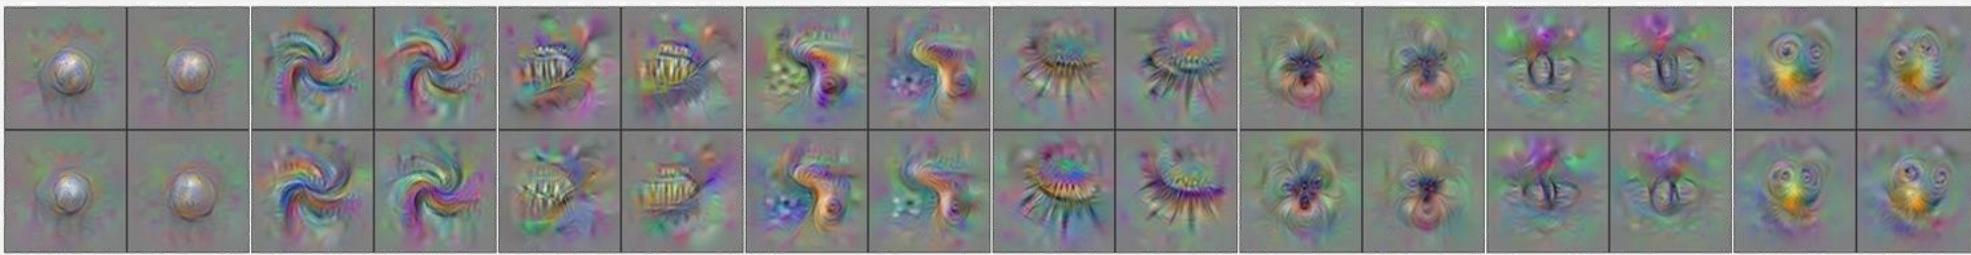
Layer 3



# HOW A CNN SEES

---

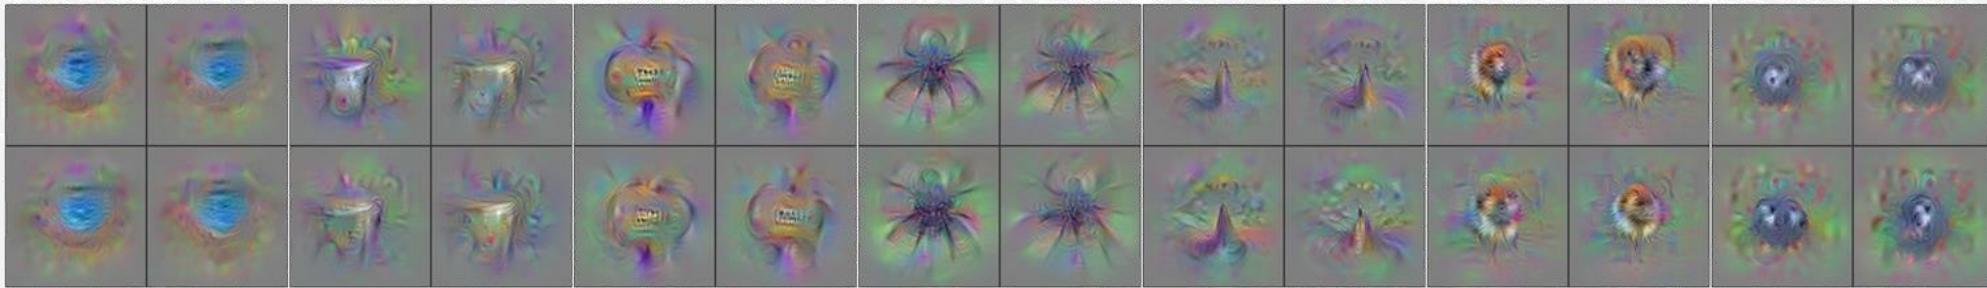
Layer 4



# HOW A CNN SEES

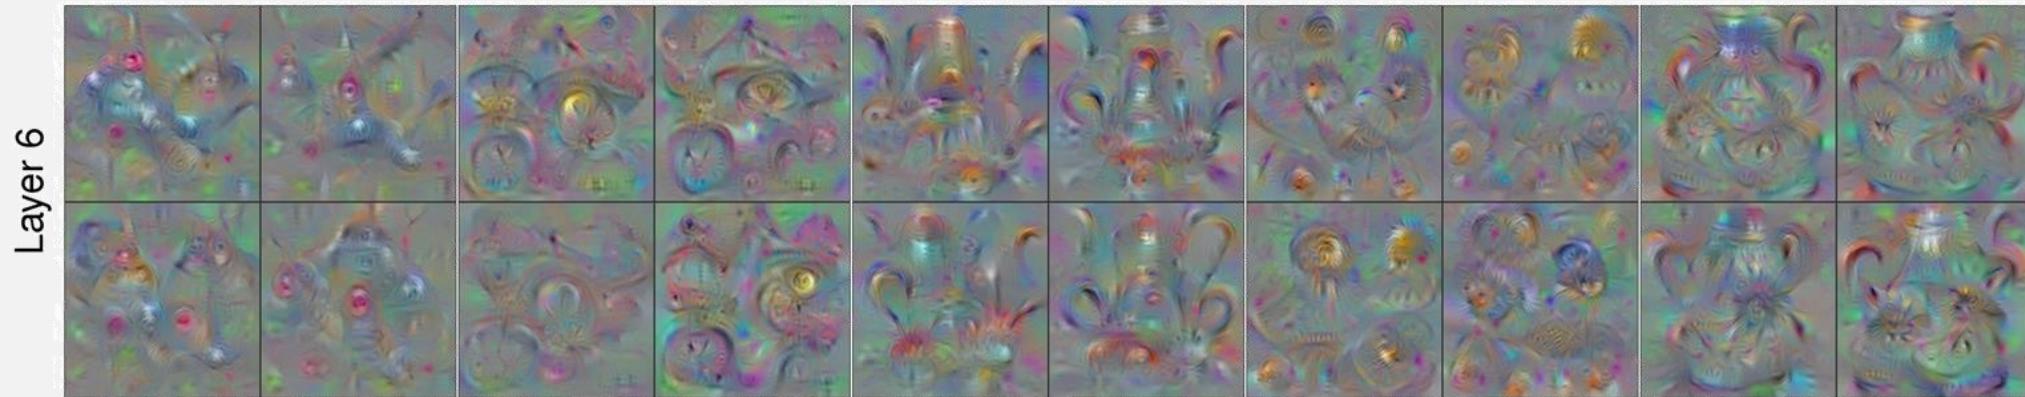
---

Layer 5



# HOW A CNN SEES

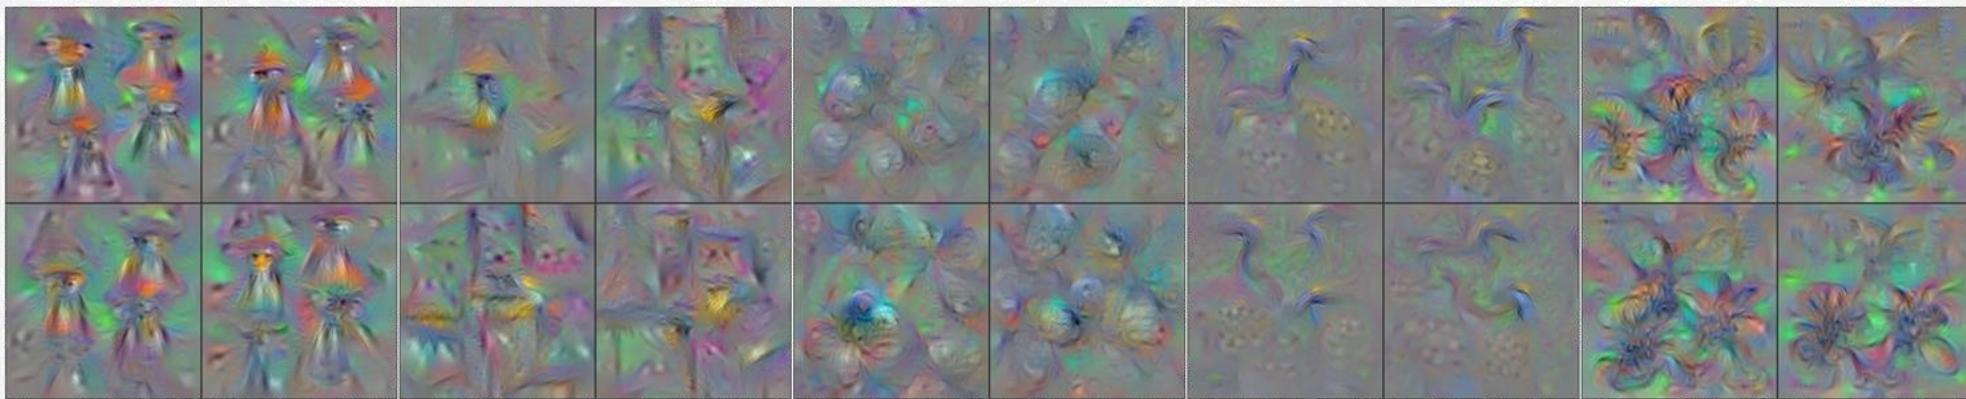
---



# HOW A CNN SEES

---

Layer 7



# HOW A CNN SEES

---

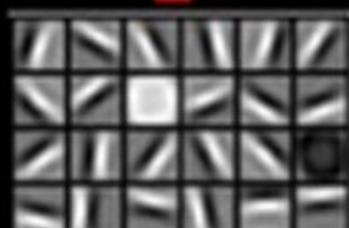
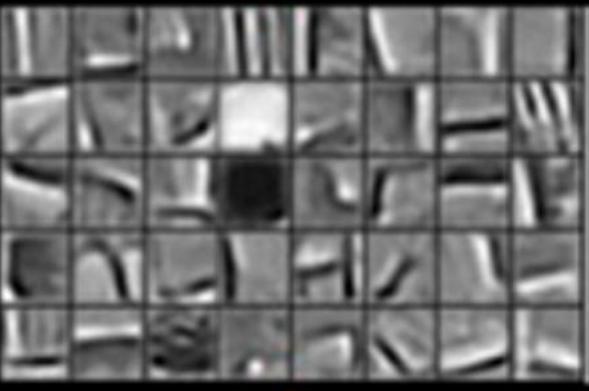
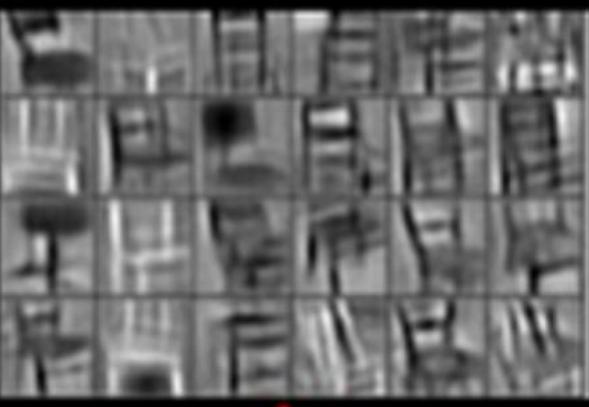
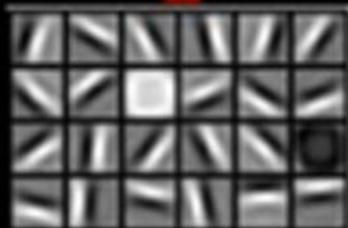
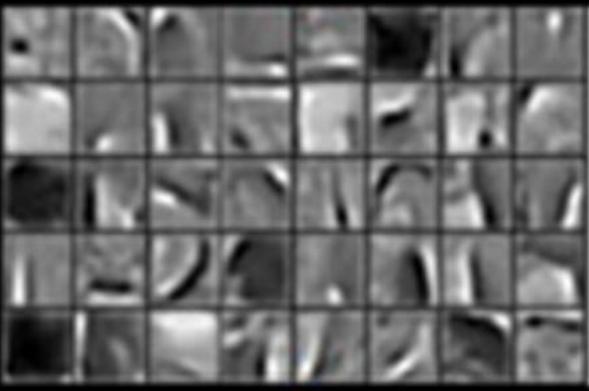
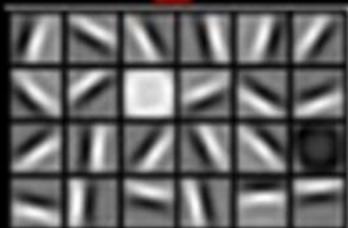
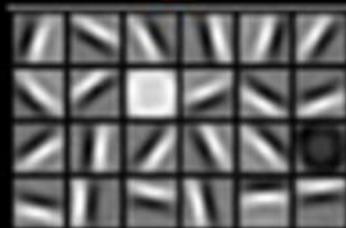


Faces

Cars

Elephants

Chairs



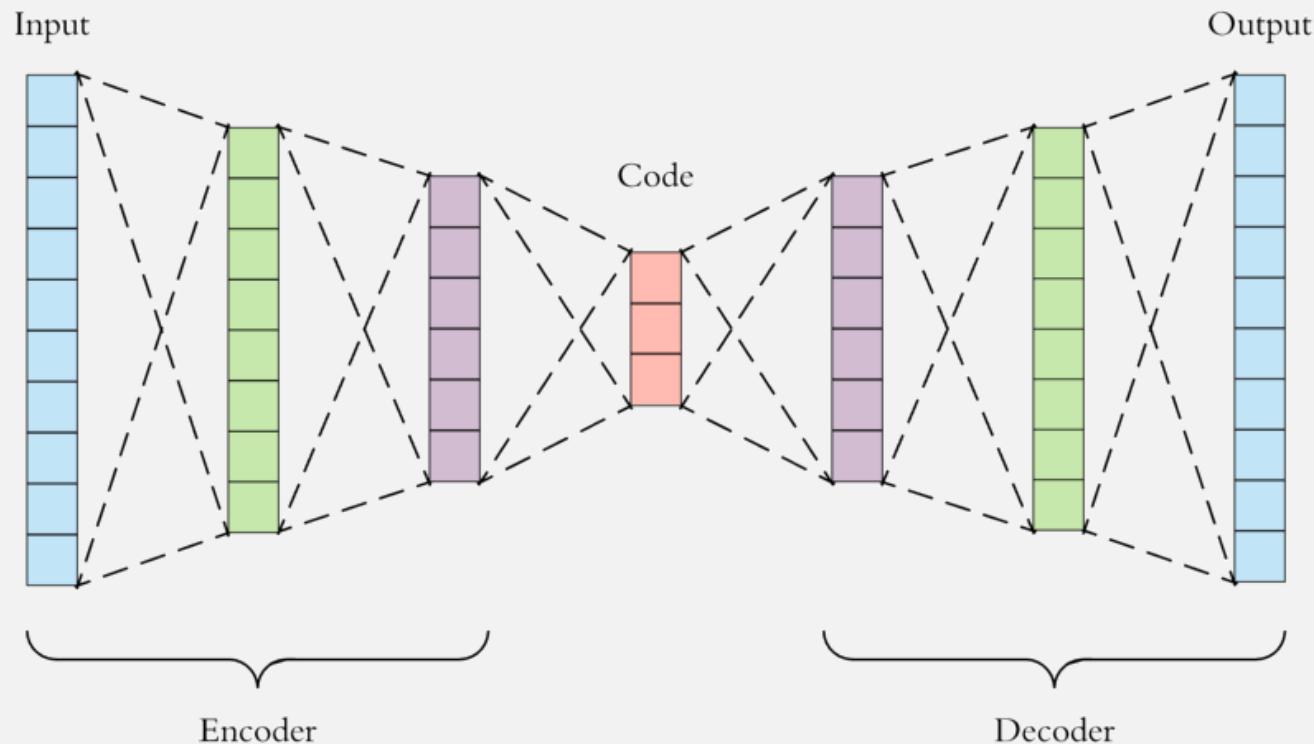
# LAB 3: Convolutional Networks



# DL: Autoencoders

# BASIC AUTOENCODER

$$\hat{x} = h_{W,b} \approx x$$



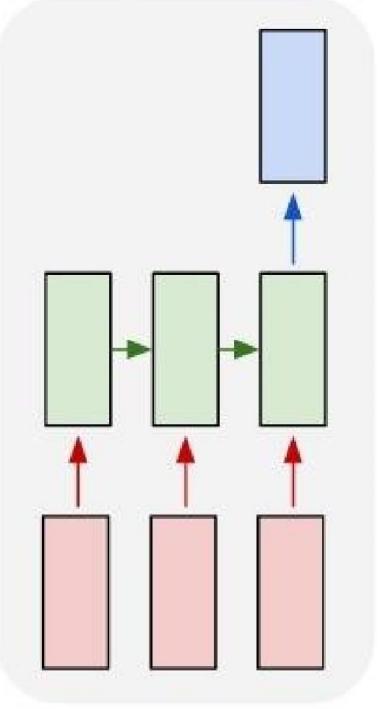
# LAB 5: Autoencoder

# DL: Recurrent Neural Networks

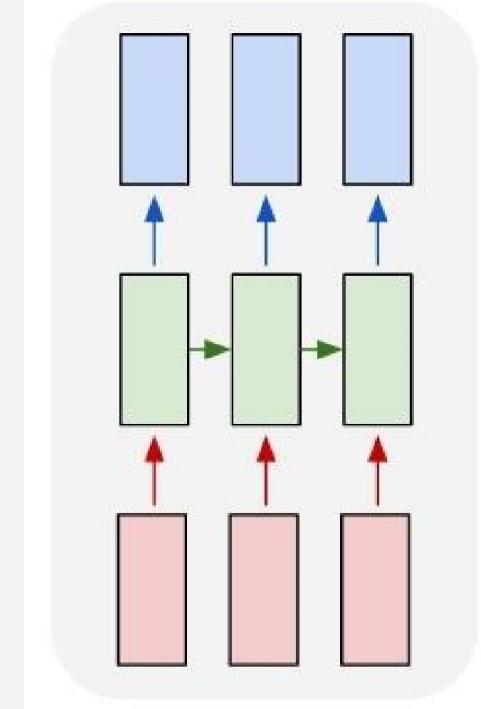
# HANDLING SEQUENCES

---

many to one

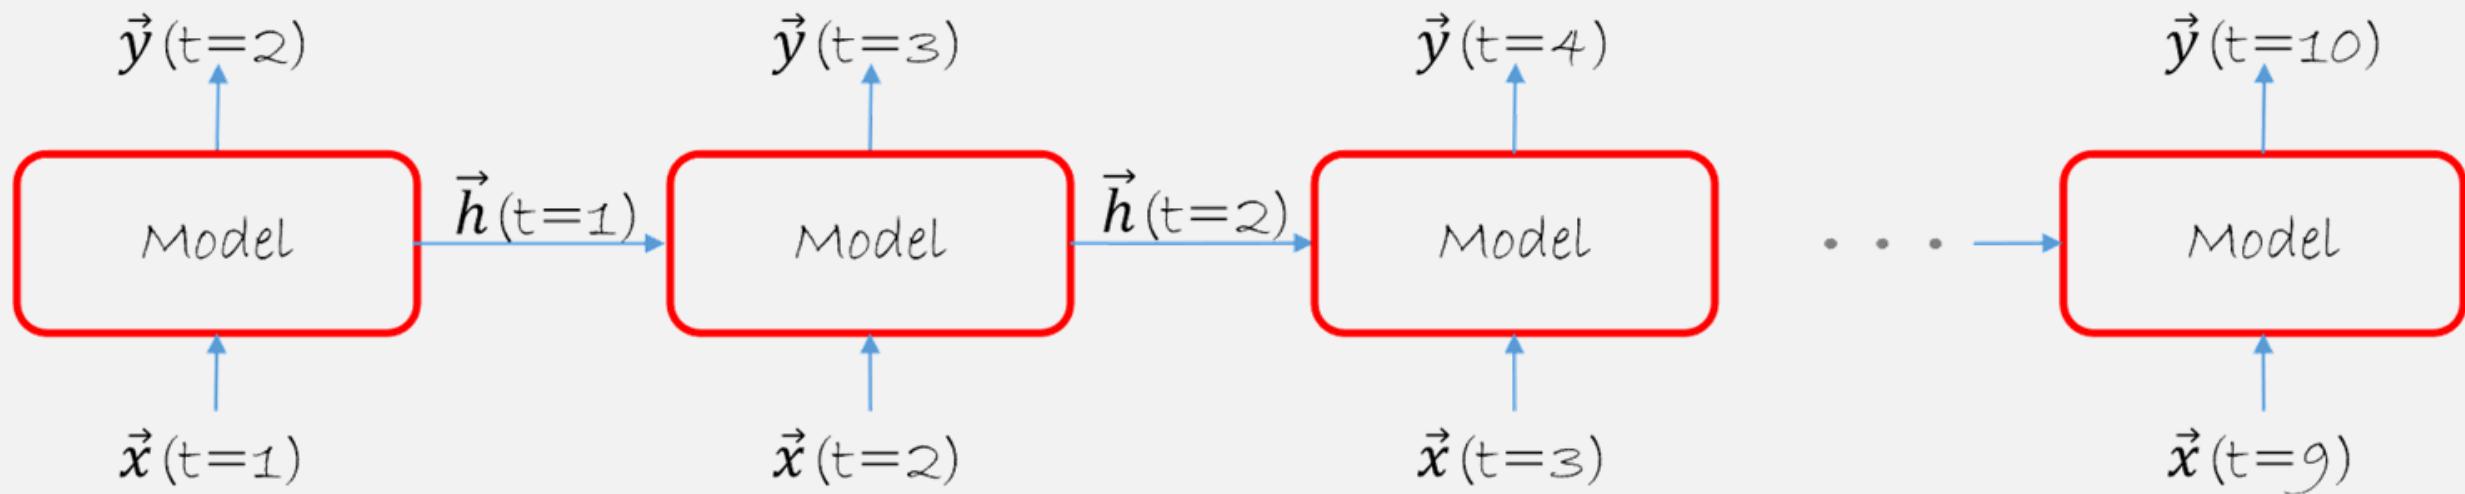


many to many



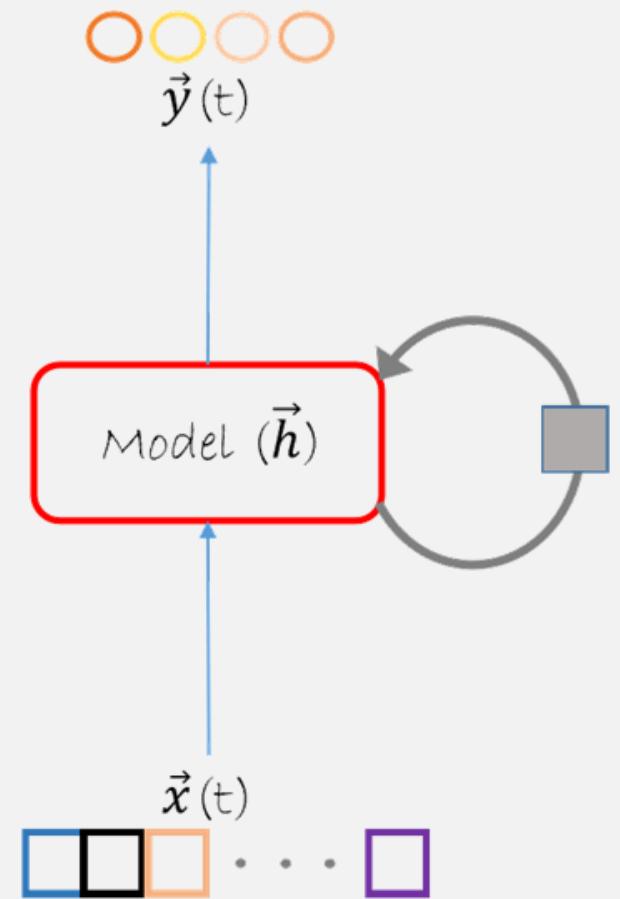
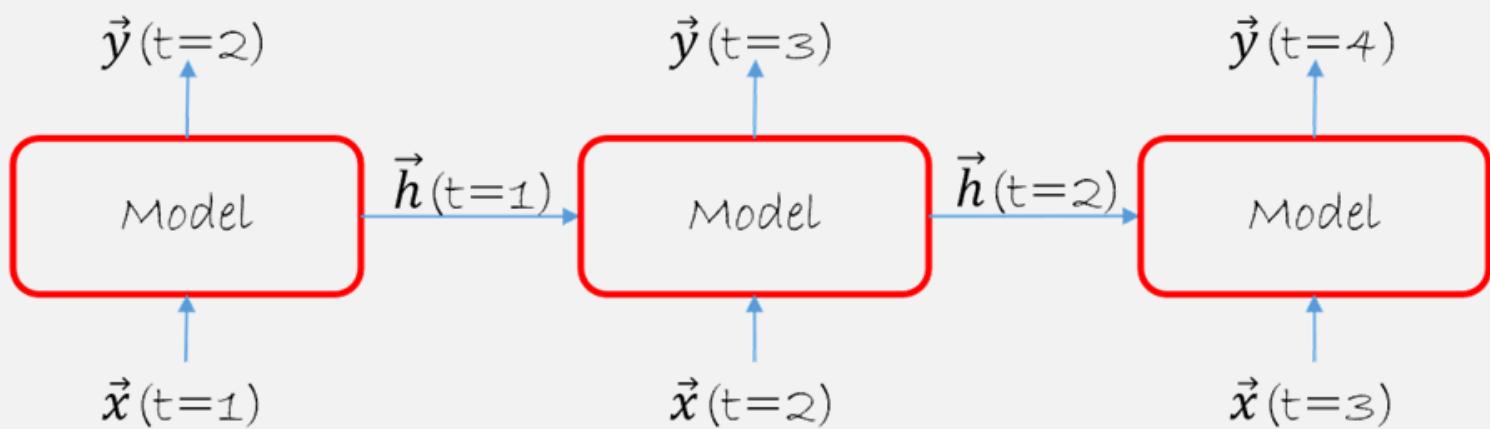
# TRACKING THE HISTORY

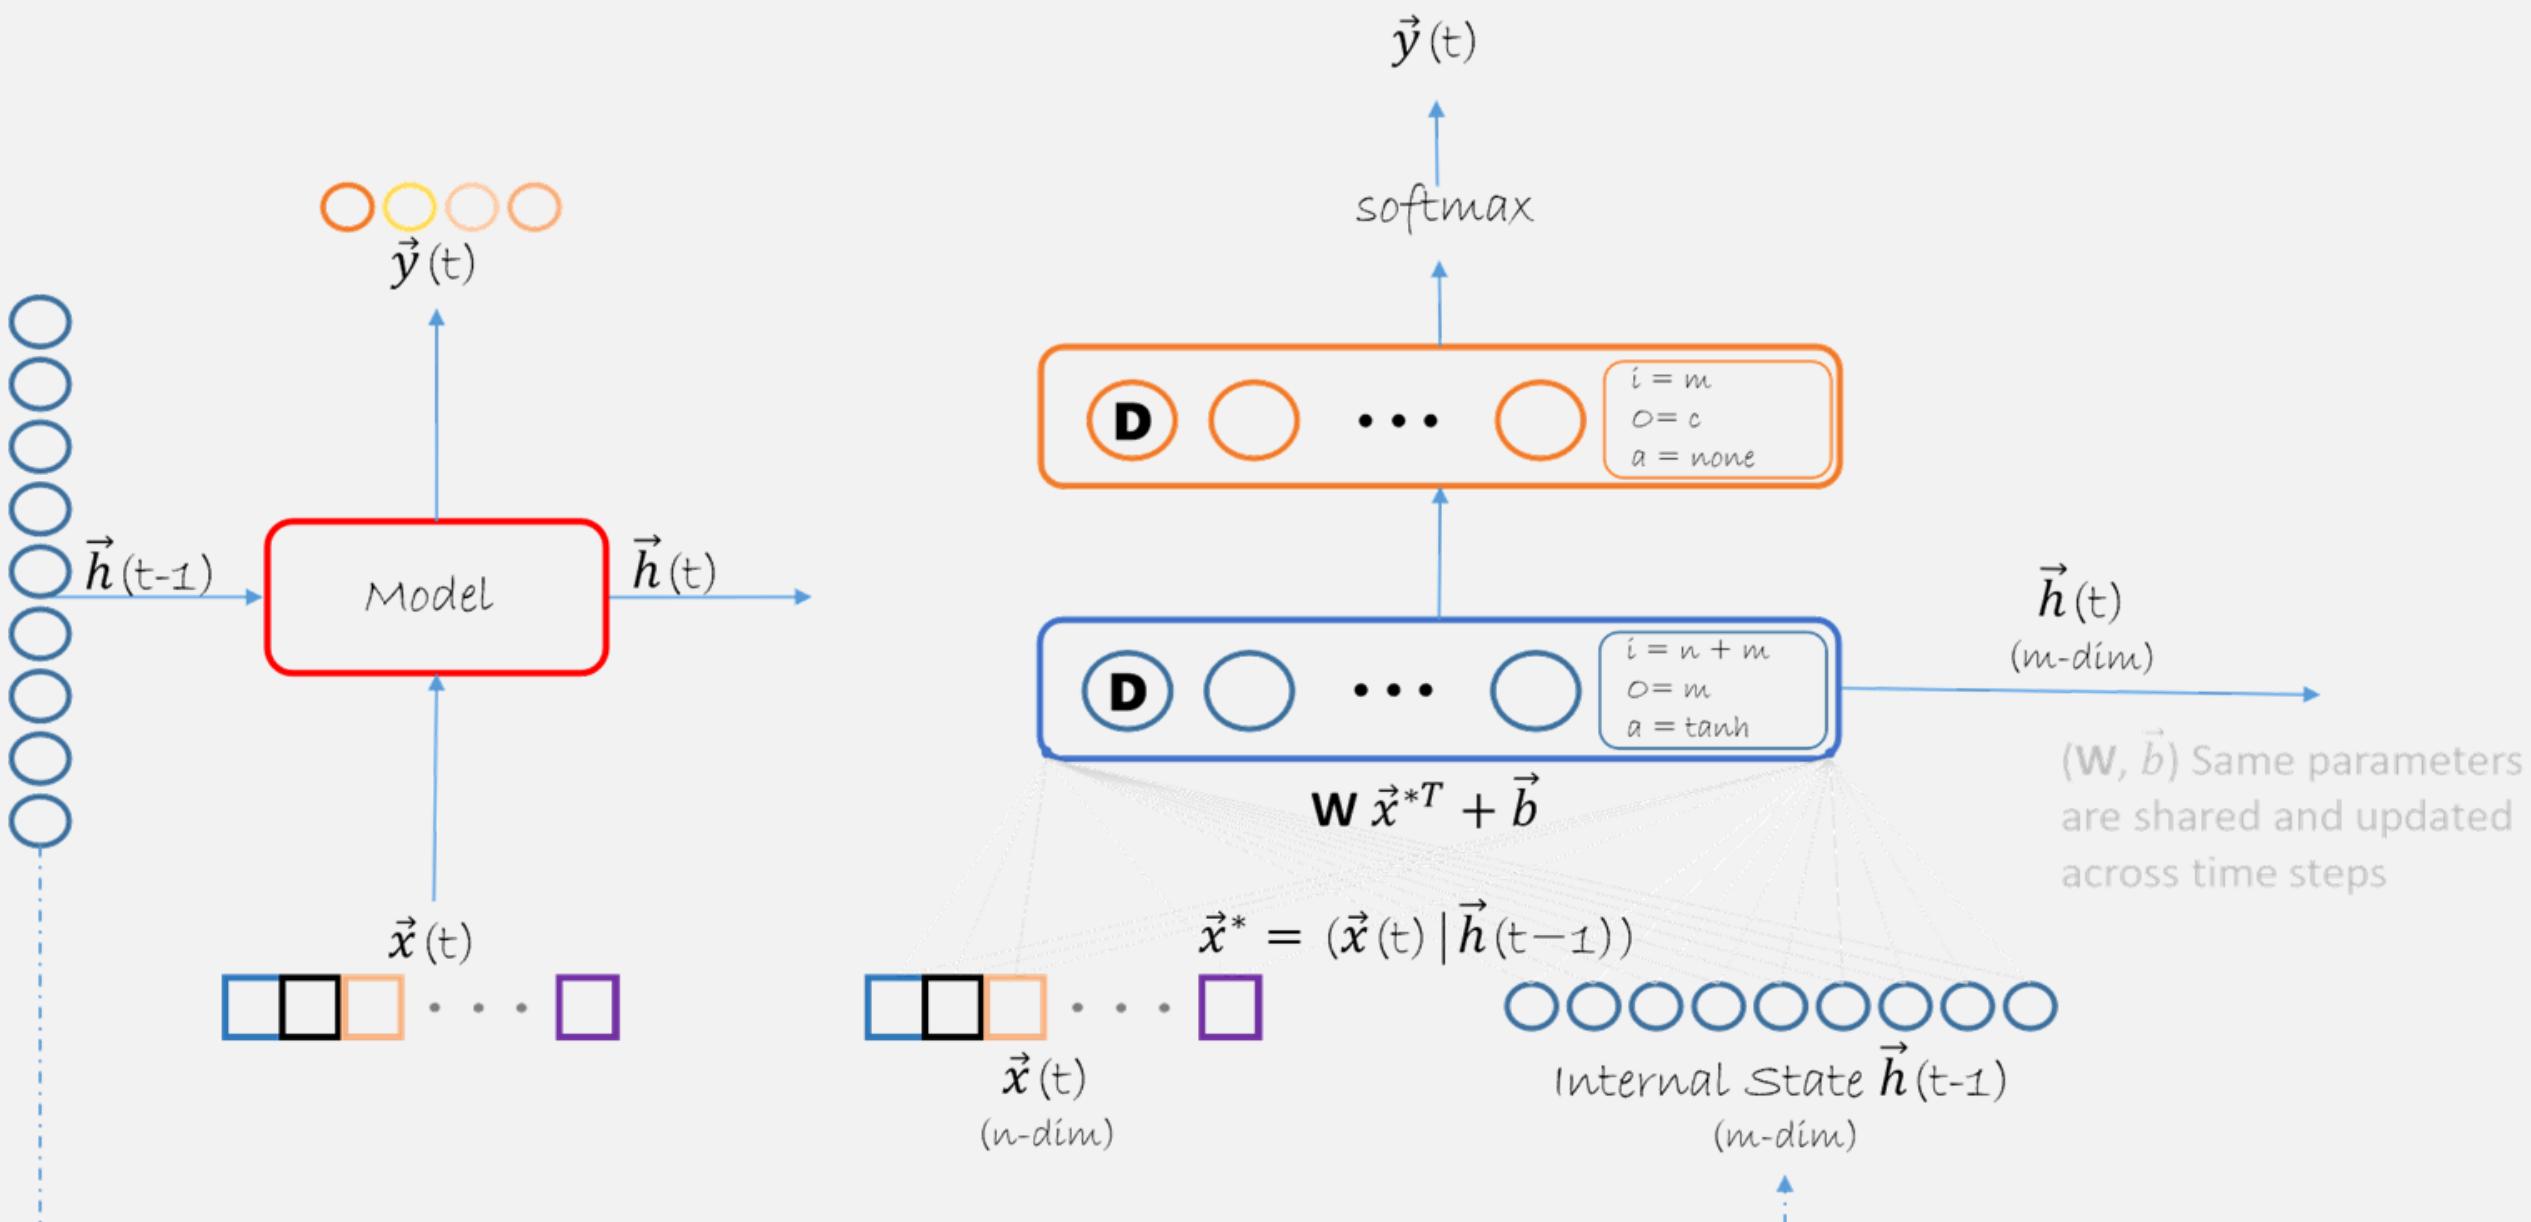
---



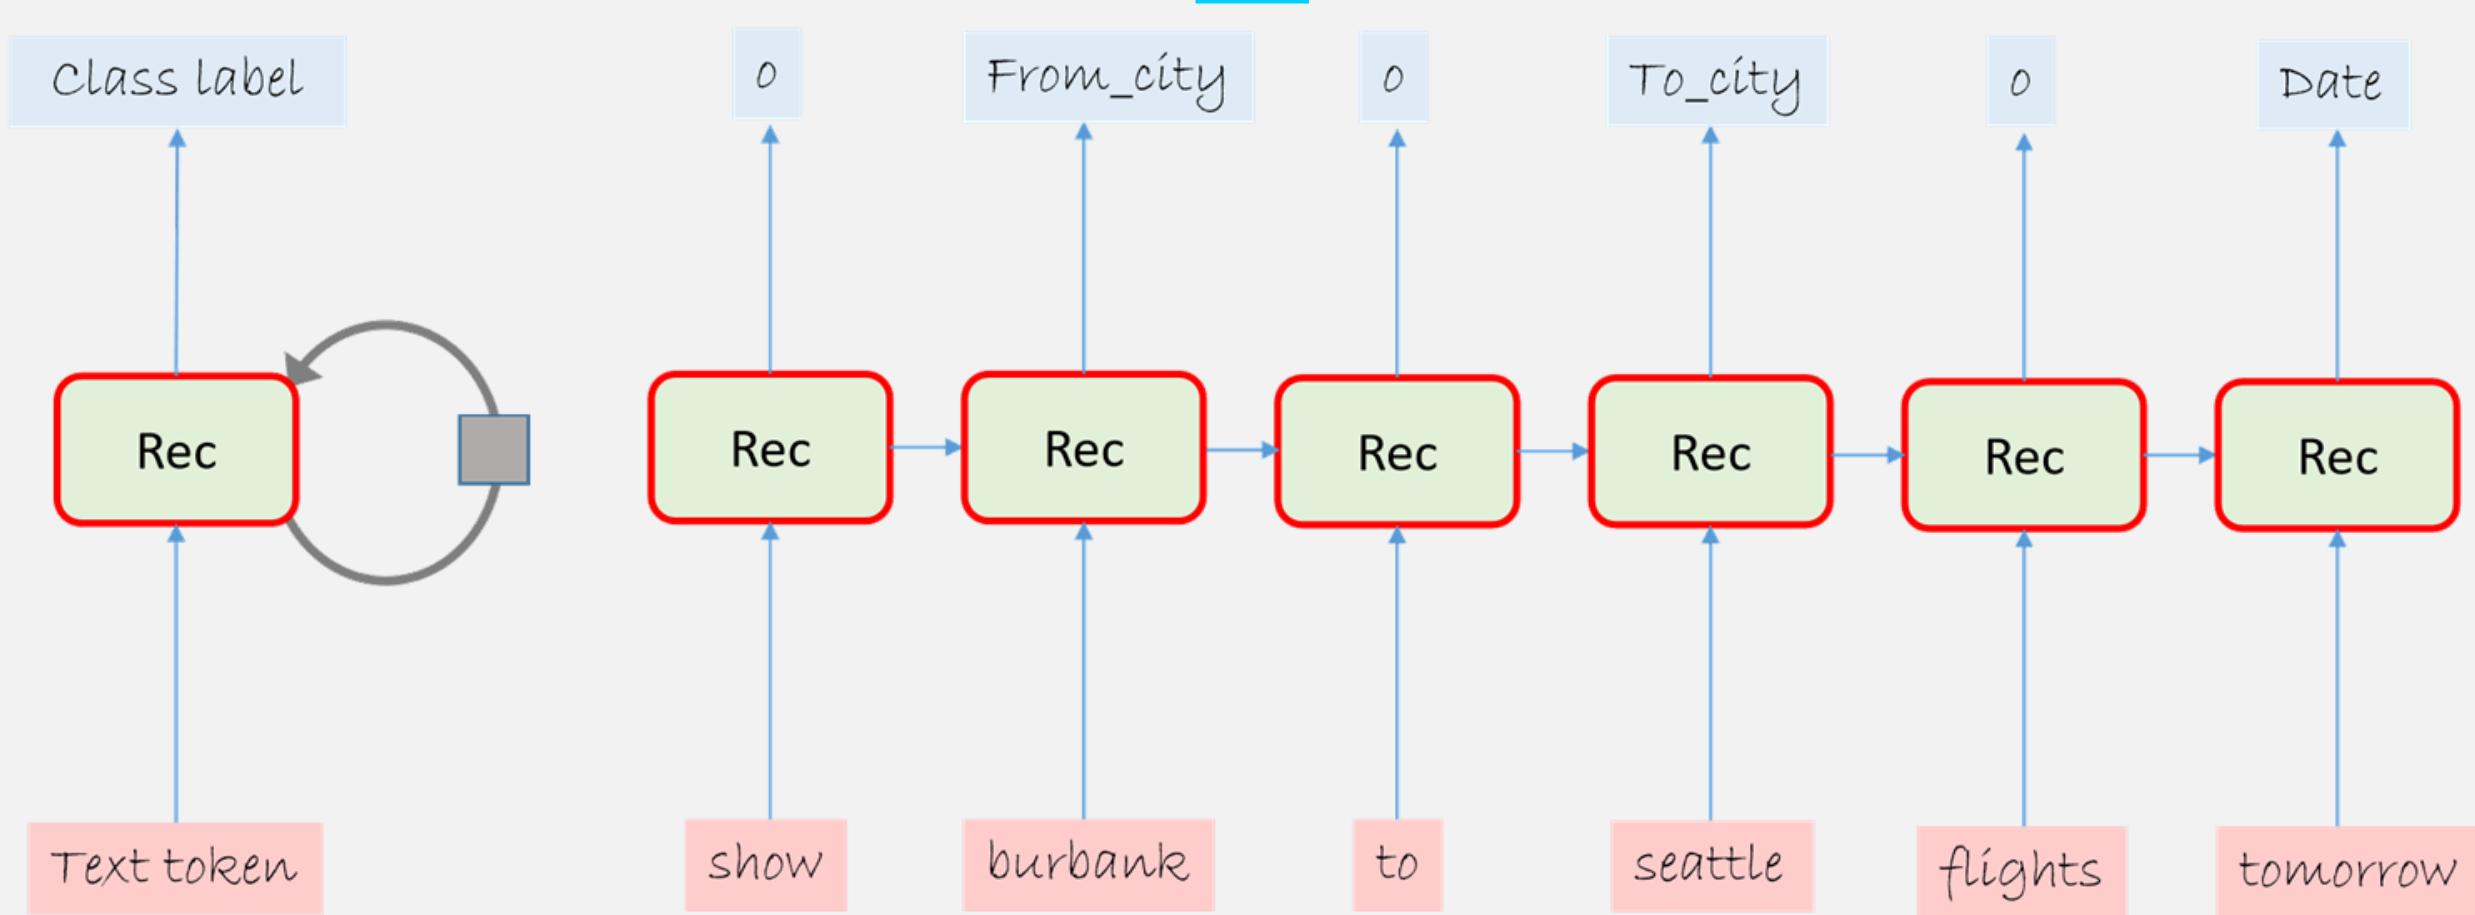
# RECURRENCE

---

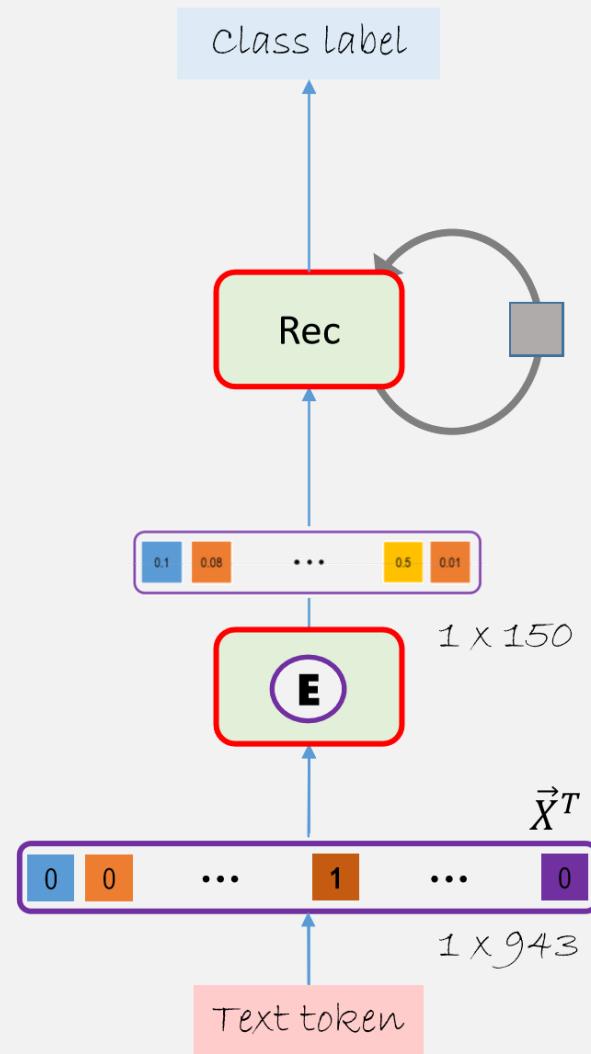


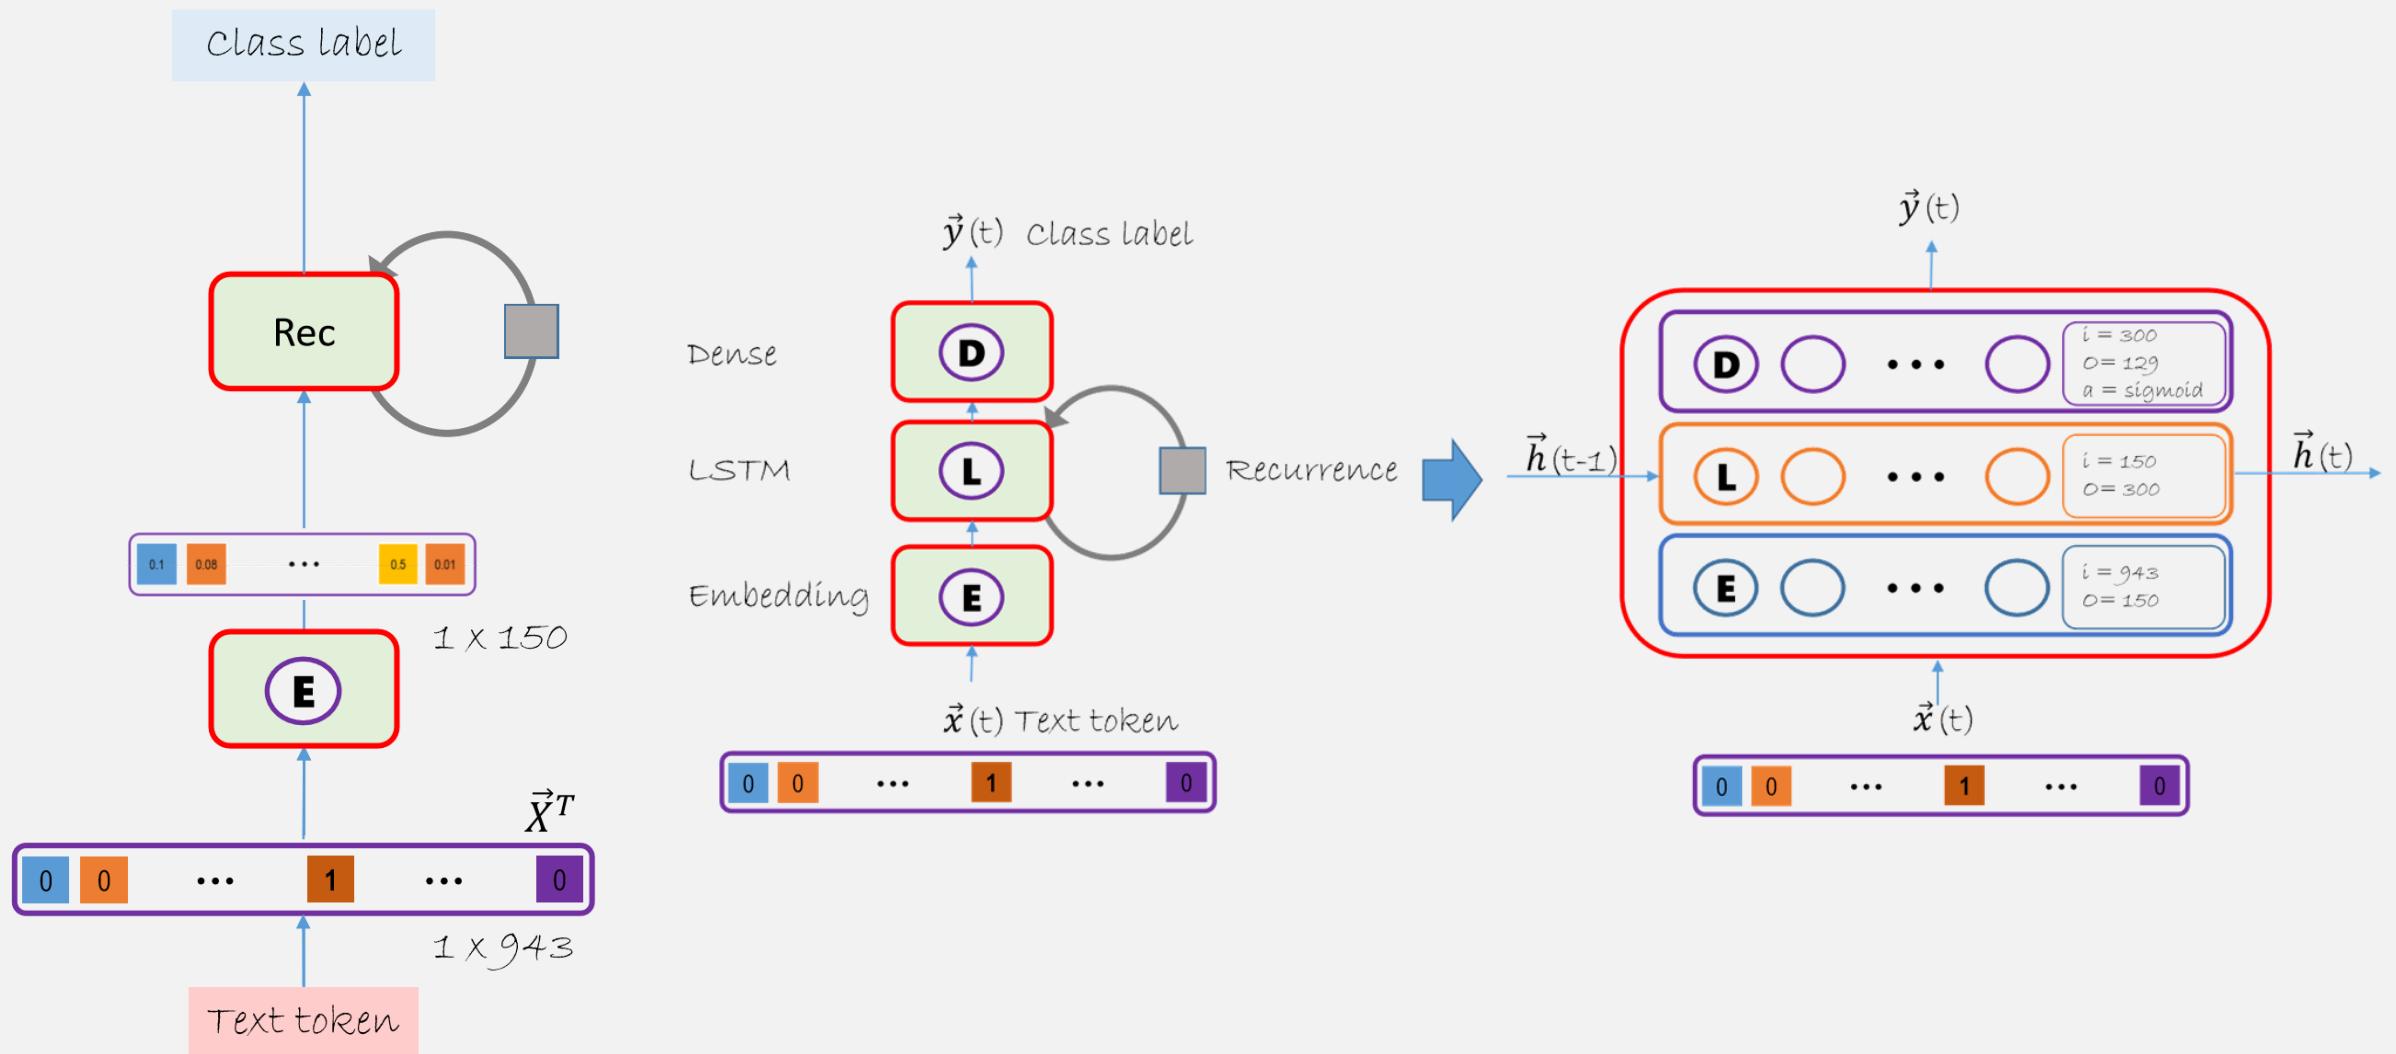


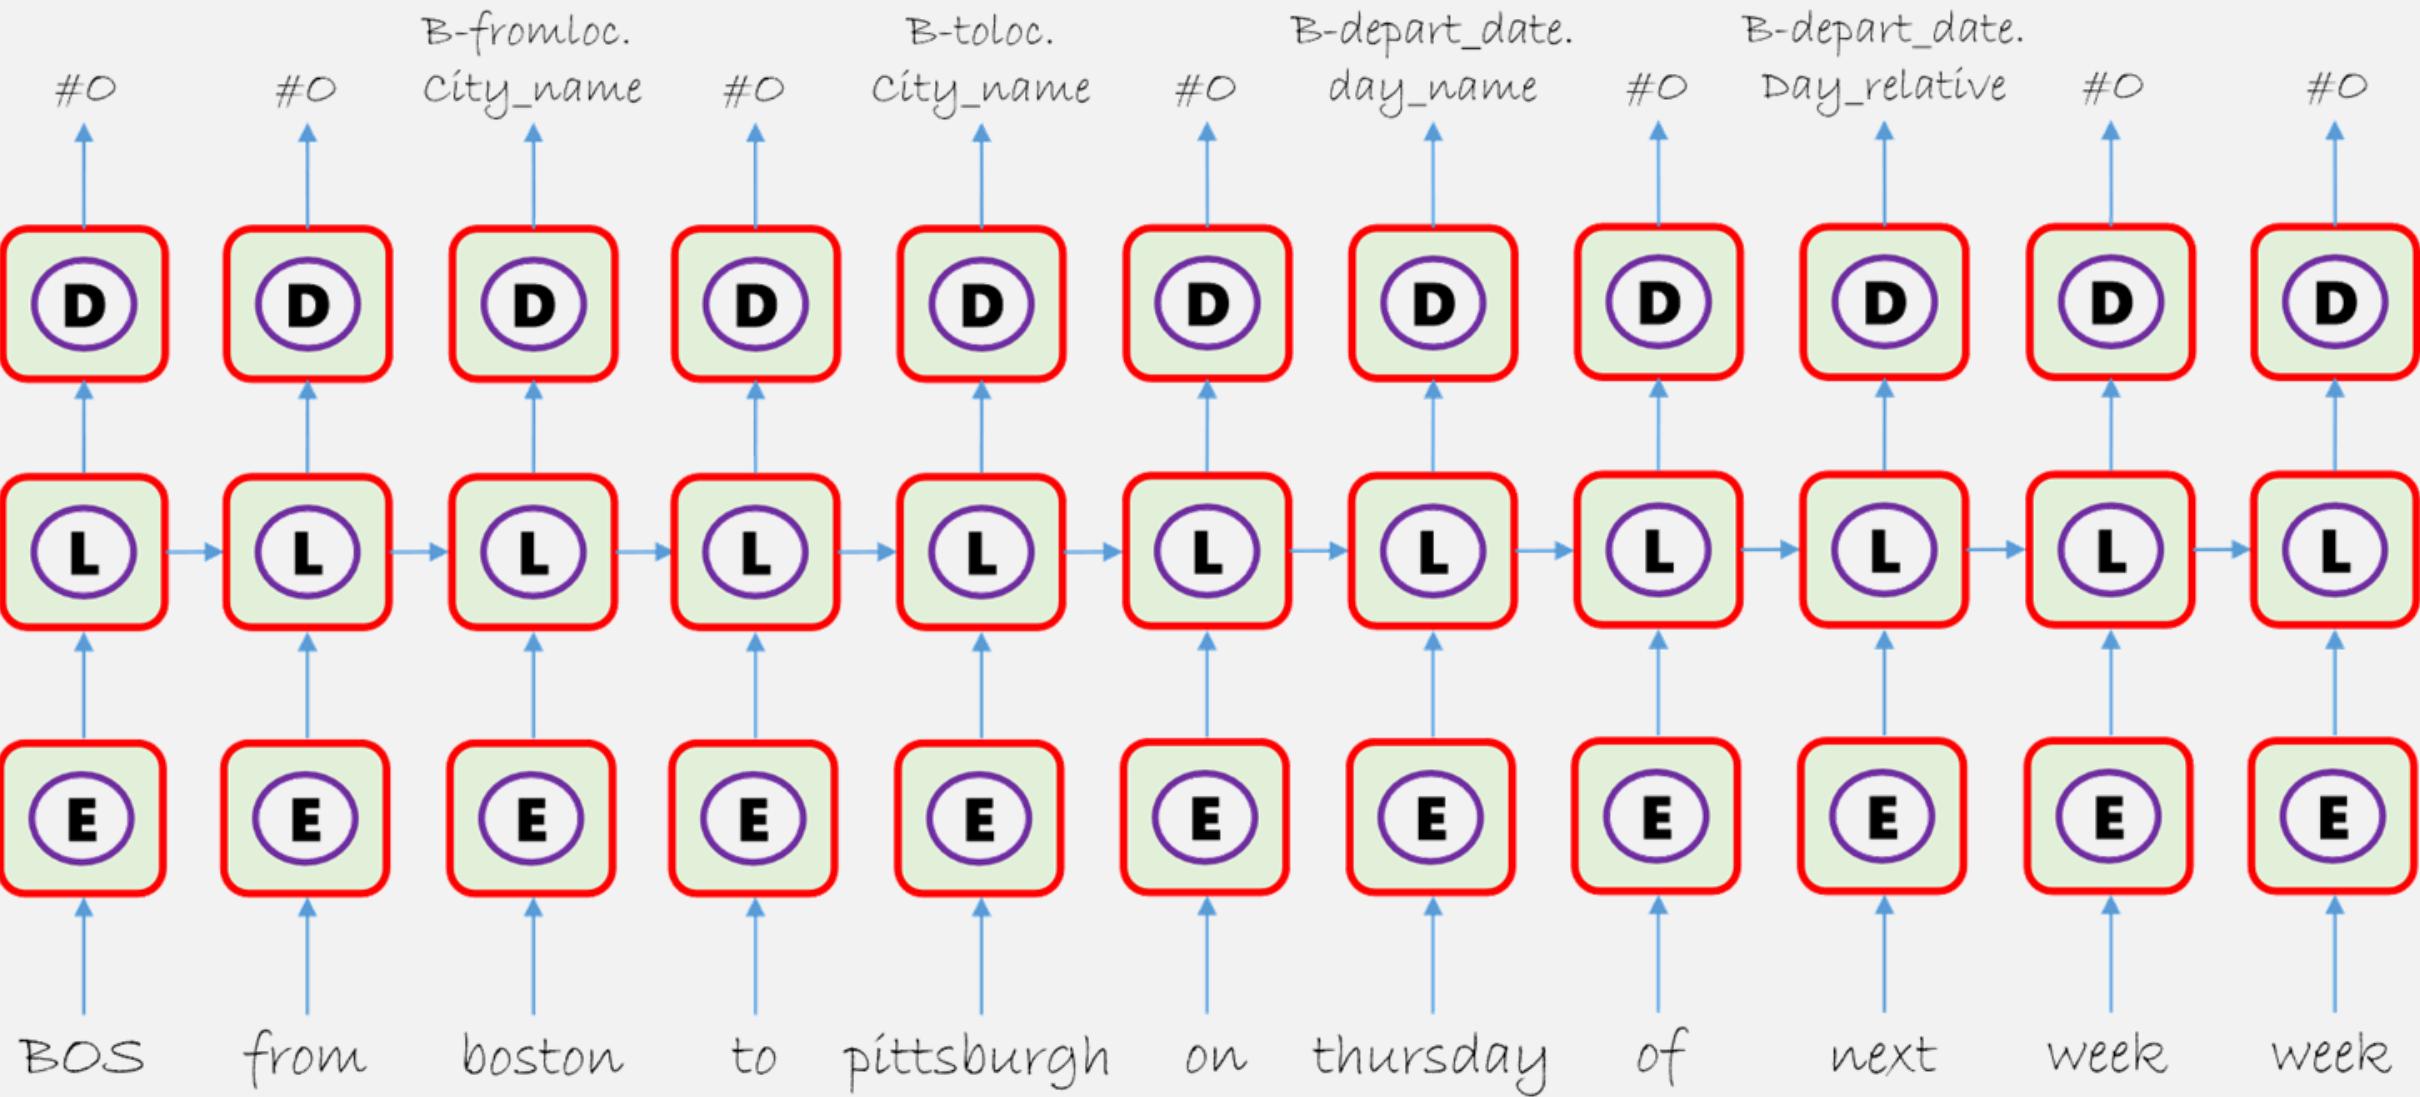
# LONG SHORT TERM MEMORY



Sequence Id	Input Word (sample)	Word Index (in vocabulary) S0	Word Label	Label Index (S2)
19	# BOS	178:1	# O	128:1
19	# please	688:1	# O	128:1
19	# give	449:1	# O	128:1
19	# me	581:1	# O	128:1
19	# the	827:1	# O	128:1
19	# flights	429:1	# O	128:1
19	# from	444:1	# O	128:1
19	# boston	266:1	# B-fromloc.city_name	48:1
19	# to	851:1	# O	128:1
19	# pittsburgh	682:1	# B-toloc.city_name	78:1
19	# on	654:1	# O	128:1
19	# thursday	845:1	# B-depart_date.day_name	26:1
19	# of	646:1	# O	128:1
19	# next	621:1	# B-depart_date.date_relative	25:1
19	# week	910:1	# O	128:1
19	# EOS	179:1	# O	128:1







# Some Advanced Ideas

# Counting people in a multitude

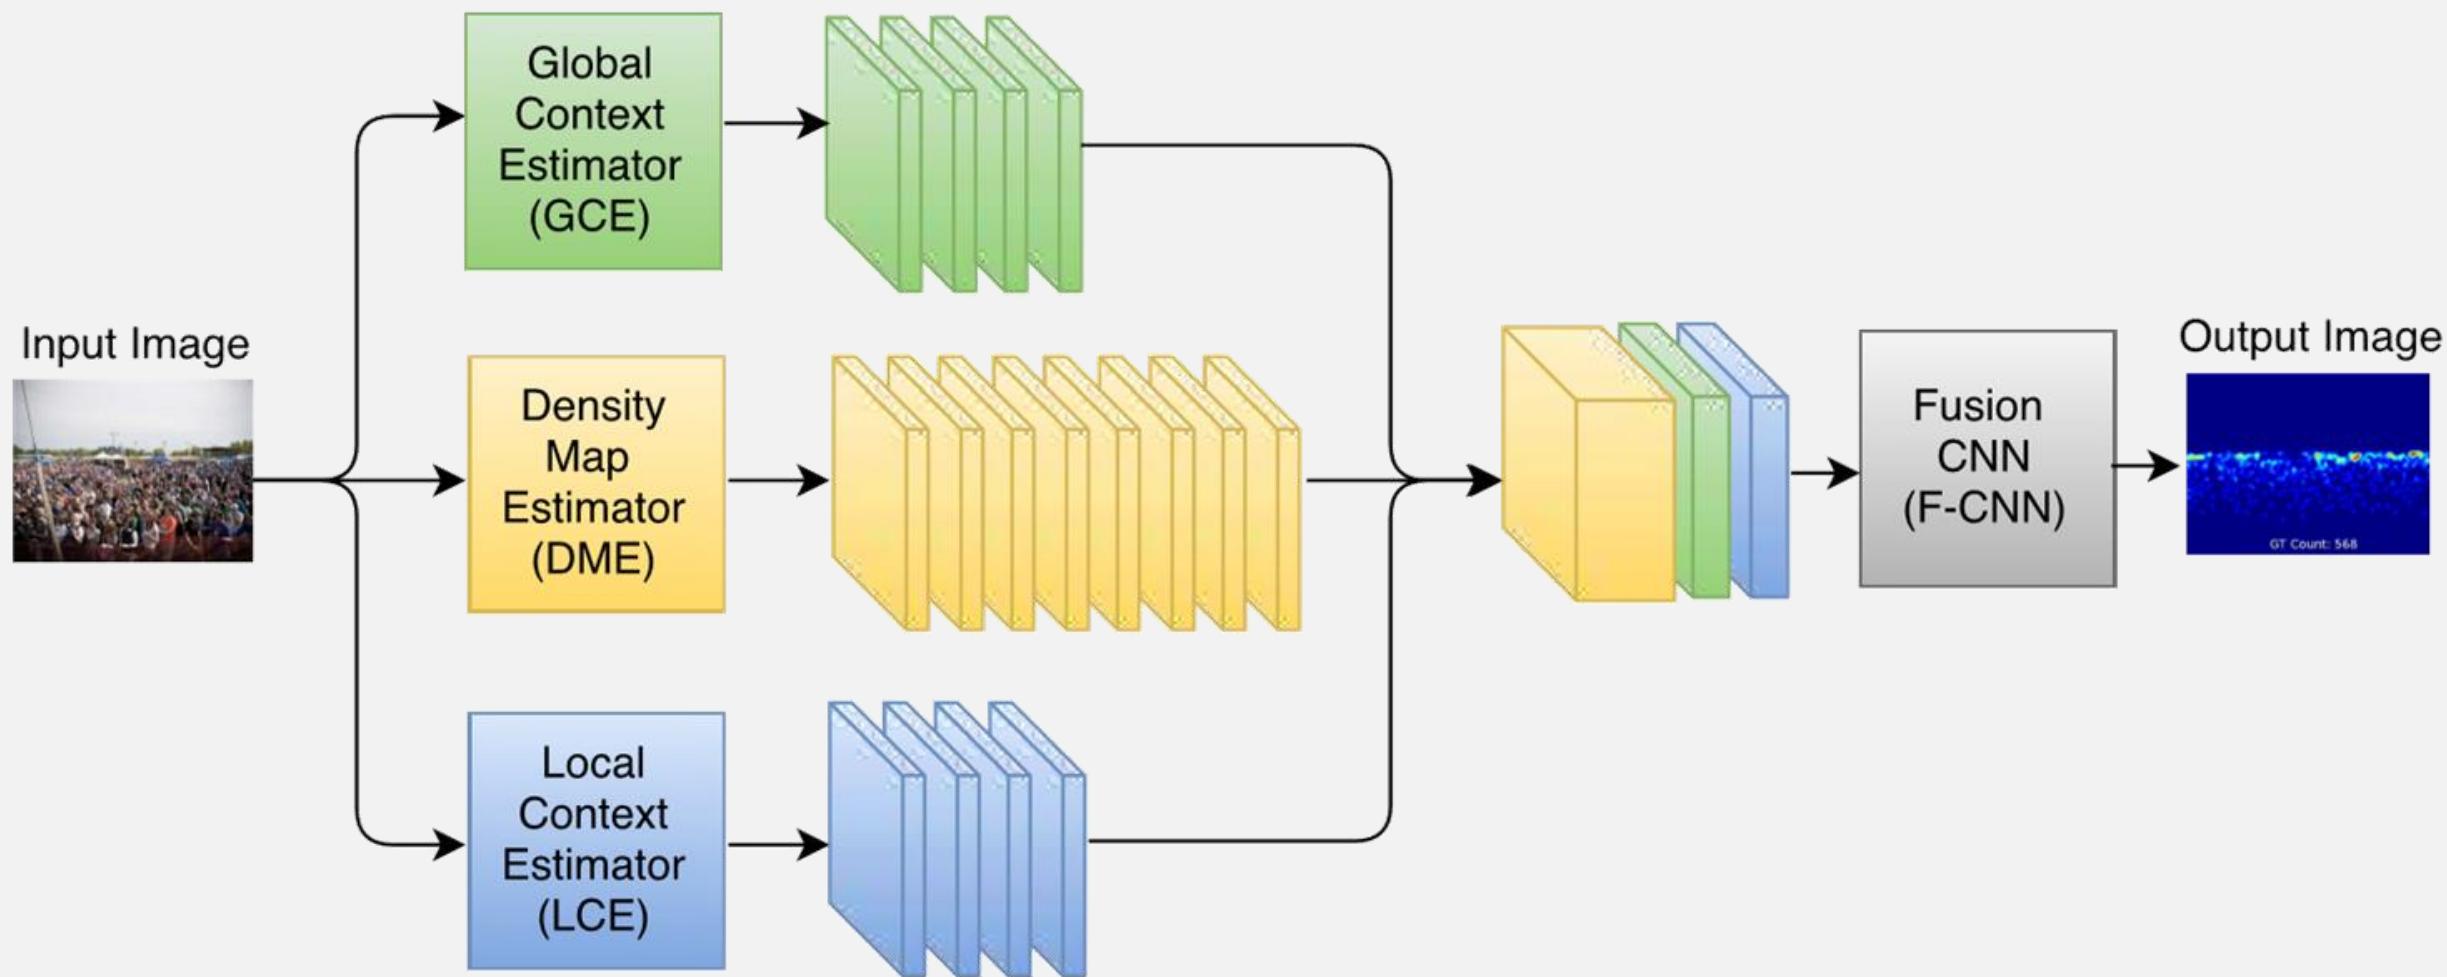
---

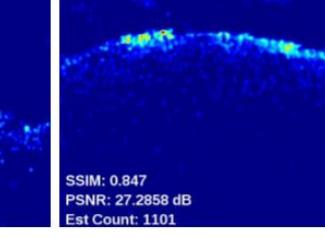
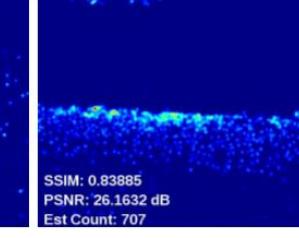
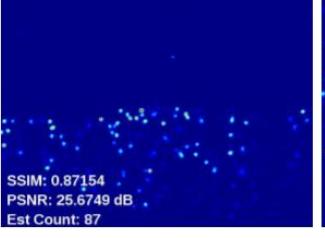
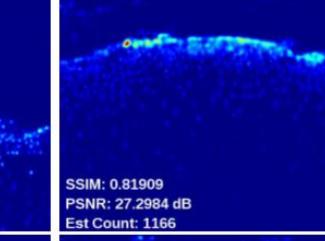
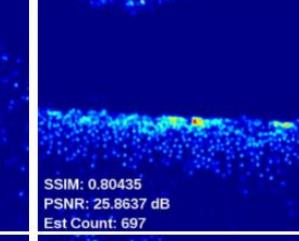
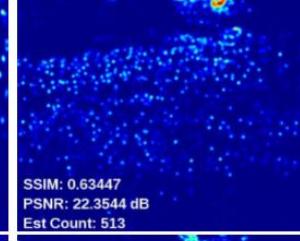
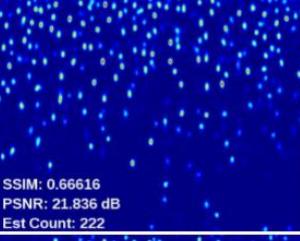
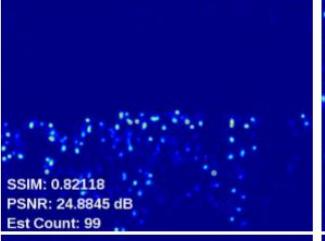
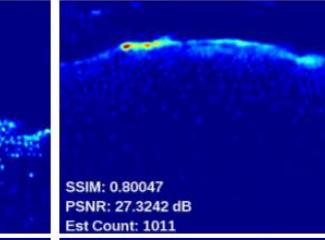
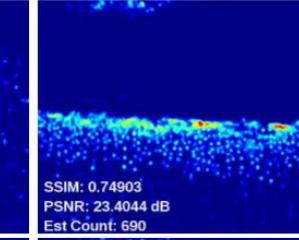
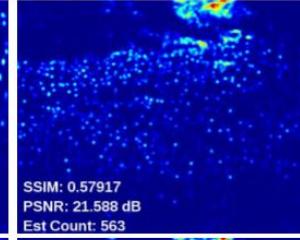
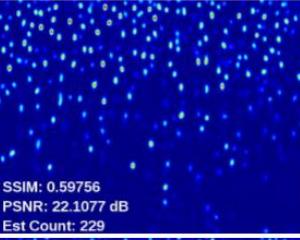
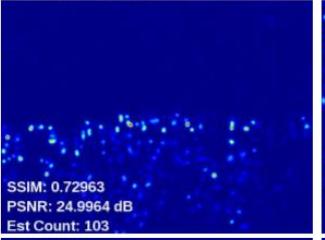
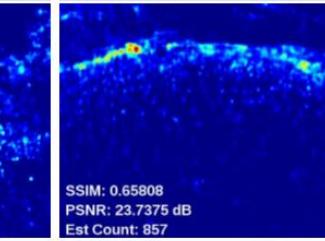
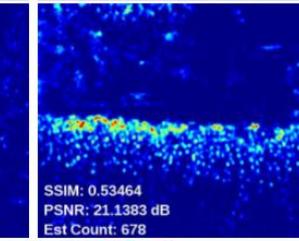
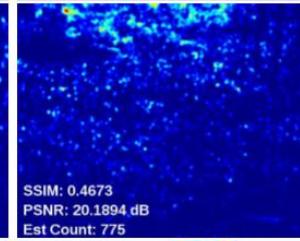
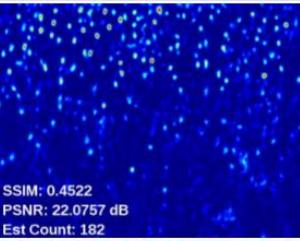
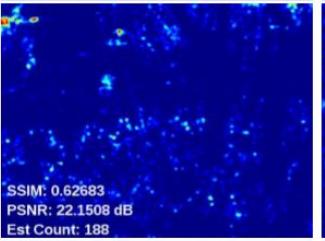
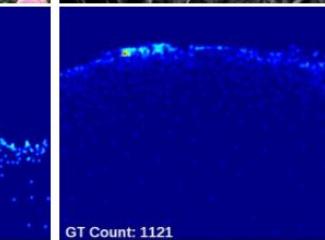
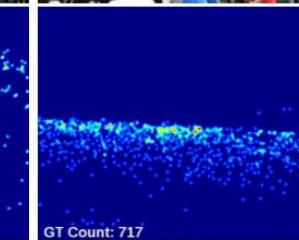
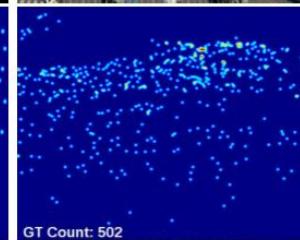
## CONTEXTUAL PYRAMID CNN (CP-CNN)

A new method named Contextual Pyramid CNN (CP-CNN) is proposed here to generate density maps and influx estimations, by explicitly incorporating global and local context information. Composed of four modules: Global Context Estimator (GCE), Local Context Estimator (LCE), Density Map Estimator (DME) and a Fusion-CNN (F-CNN) convolutional network.

**Vishwanath A. Sindagi, Vishal M. Patel;** The IEEE International Conference on Computer Vision (ICCV), 2017, pp. 1861-1870









# Transferring style between images

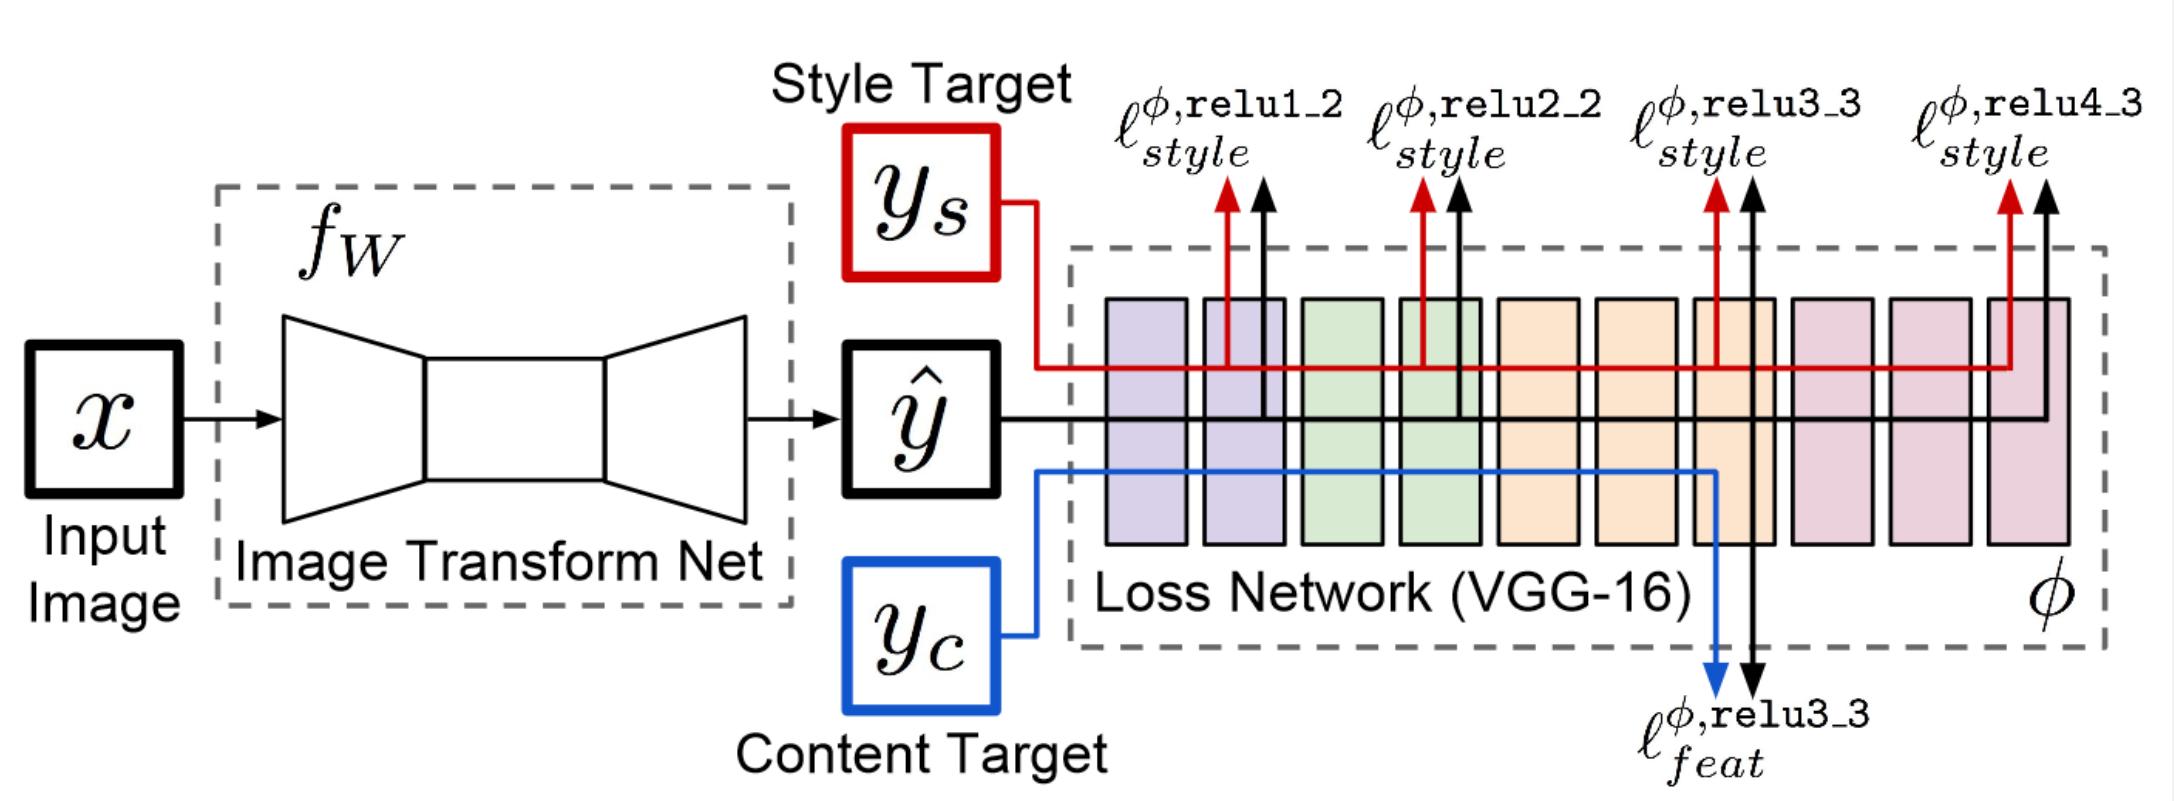
---

## CONVOLUTIONAL NEURAL NETWORKS

By using a perceptual loss functions based on high-level features extracted from pretrained networks, networks for image transformation tasks can be trained, and by fine tuning the loss function different features can be kept for the source image and the style image.

**Justin Johnson, Alexandre Alahi, Li Fei-Fei;** Perceptual Losses for Real-Time Style Transfer and Super-Resolution, 2016









# Using GANs to drive design decisions

## GENERATIVE ADVERSARIAL NETWORKS

By taking advantage of Generational Adversarial Networks, synthetic images based on the training data can be generated. Including an external array of features, the generated images can be tailored to a specific set of requirements.

***Jaime Deverall, Jiwoo Lee, Miguel Ayala; Using Generative Adversarial Networks to Design Shoes***

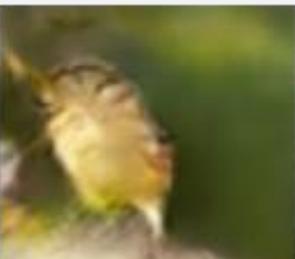


Text  
description

This bird is blue with white and has a very short beak



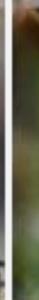
This bird has wings that are brown and has a yellow belly



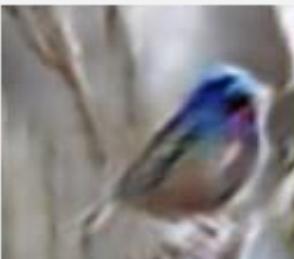
A white bird with a black crown and yellow beak



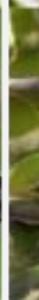
This bird is white, black, and brown in color, with a brown beak



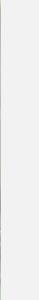
The bird has small beak, with reddish brown crown and gray belly



This is a small, black bird with a white breast and white on the wingbars.



This bird is white black and yellow in color, with a short black beak



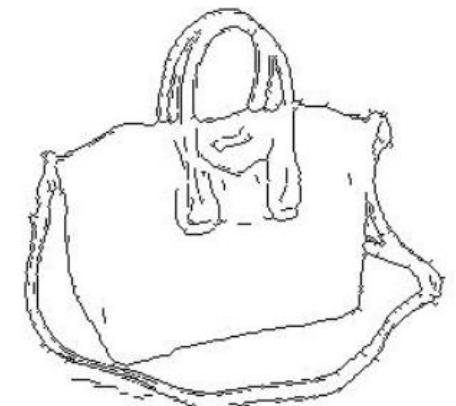
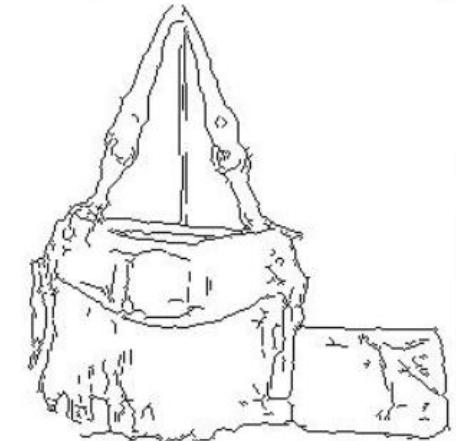
Input



Ground truth



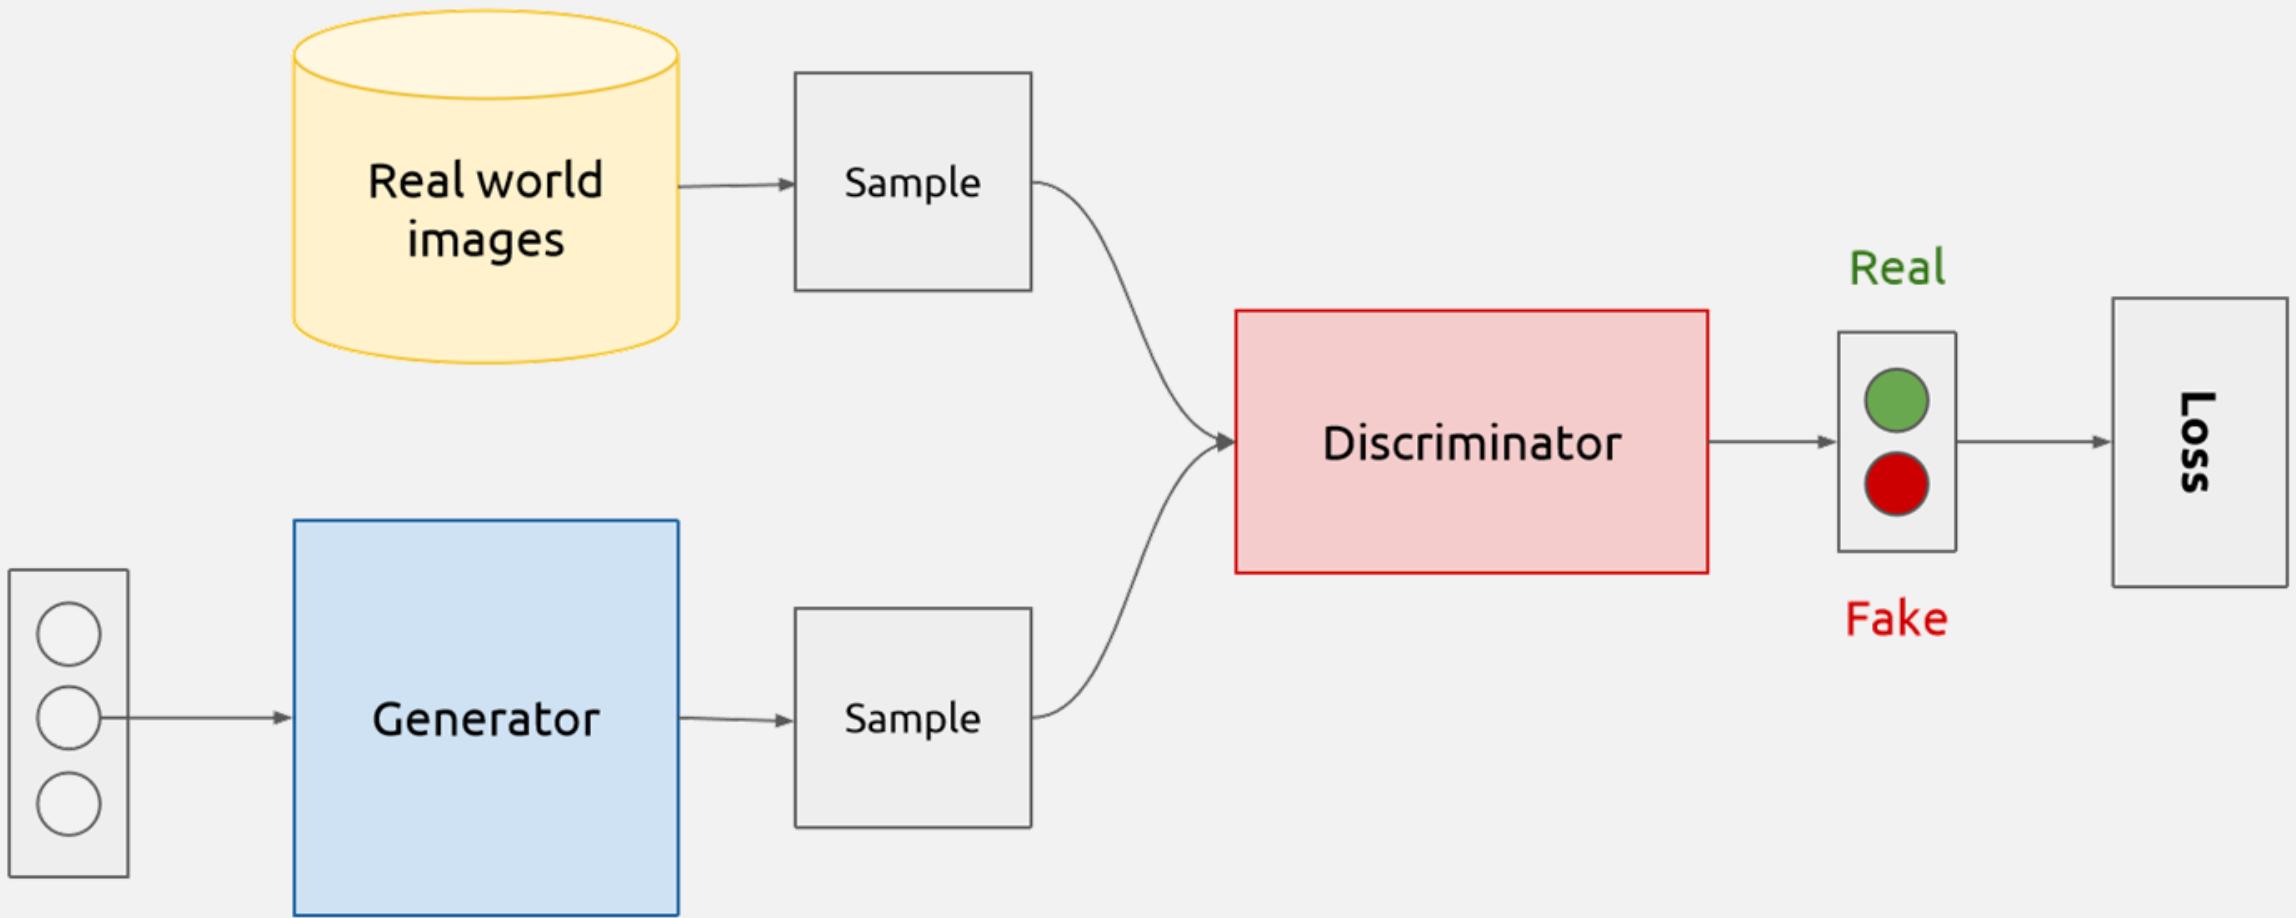
Output





plain concepts

Latent random variable



# DEEP LEARNING - MOVING FAST!

---

- Swish Activation Function
- Capsule Networks
- ...

# SOME INTERESTING PAPERS

“ImageNet Classification with Deep Convolutional Neural Networks”

[Krizhevsky, Sutskever, Hinton]

<http://www.cs.toronto.edu/~fritz/absps/imagenet.pdf>

“Visualizing and Understanding Convolutional Networks”

[Zeiler, Fergus]

arXiv: 1311.2901v3

“Optimization Methods for Large-Scale Machine Learning”

[Bottou, Curtis, Nocedal]

arXiv: 1606.04838v2

# SOME INTERESTING PAPERS

“Swish: A Self-Gated Activation Function”

[Ramachandran, Zoph, Quoc V.]

arXiv: 1710.05941v2

“Dynamic Routing Between Capsules”

[Sabour, Frosst, Hinton]

arXiv: 1710.09829